



US008561044B2

(12) **United States Patent**
Chen et al.

(10) **Patent No.:** **US 8,561,044 B2**
(45) **Date of Patent:** **Oct. 15, 2013**

(54) **OPTIMIZED CODE GENERATION
TARGETING A HIGH LOCALITY
SOFTWARE CACHE**

(75) Inventors: **Tong Chen**, Yorktown Heights, NY
(US); **Alexandre E. Eichenberger**,
Chappaqua, NY (US); **Marc Gonzalez**
Tallada, Barcelona (ES); **John K.**
O'Brien, South Salem, NY (US);
Kathryn M. O'Brien, South Salem, NY
(US); **Zehra N. Sura**, Yorktown Heights,
NY (US); **Tao Zhang**, Yorktown
Heights, NY (US)

(73) Assignee: **International Business Machines
Corporation**, Armonk, NY (US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 704 days.

(21) Appl. No.: **12/246,602**

(22) Filed: **Oct. 7, 2008**

(65) **Prior Publication Data**

US 2010/0088673 A1 Apr. 8, 2010

(51) **Int. Cl.**
G06F 9/44 (2006.01)

(52) **U.S. Cl.**
USPC **717/151**; 717/110; 717/168; 714/38;
714/216

(58) **Field of Classification Search**
USPC 717/110, 151, 168; 714/38, 216;
711/216
See application file for complete search history.

(56) **References Cited**

U.S. PATENT DOCUMENTS

5,553,254 A * 9/1996 Berstis et al. 712/207
6,094,702 A 7/2000 Williams et al.
6,549,983 B1 4/2003 Han et al.

6,564,297 B1 5/2003 Kosche
6,622,208 B2 9/2003 North
6,950,902 B2 9/2005 Sakai et al.
7,124,249 B1 10/2006 Darcy
7,243,195 B2 7/2007 O'Brien et al.
2004/0133747 A1 * 7/2004 Coldewey 711/137
2004/0163080 A1 8/2004 Menon et al.
2004/0205307 A1 * 10/2004 Pullen et al. 711/158
2005/0081107 A1 * 4/2005 DeWitt et al. 714/38
2005/0086653 A1 * 4/2005 Heishi et al. 717/151
2006/0080372 A1 4/2006 Barua et al.

(Continued)

OTHER PUBLICATIONS

Speeding up Irregular Applications in Shared-memory Multiproces-
sor, Zheng Zhang, ISCA 1995, Santa Margherita Ligure Italy.*

(Continued)

Primary Examiner — Chat Do

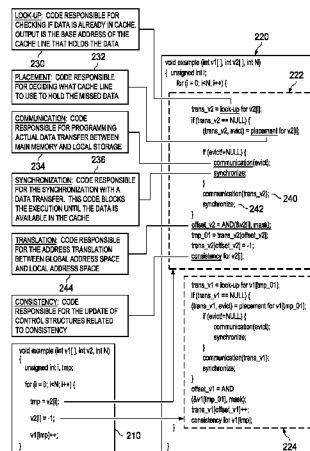
Assistant Examiner — Lynda Dinh

(74) Attorney, Agent, or Firm — Stephen J. Walder, Jr.;
William Stock

(57) **ABSTRACT**

Mechanisms for optimized code generation targeting a high
locality software cache are provided. Original computer code
is parsed to identify memory references in the original com-
puter code. Memory references are classified as either regular
memory references or irregular memory references. Regular
memory references are controlled by a high locality cache
mechanism. Original computer code is transformed, by a
compiler, to generate transformed computer code in which
the regular memory references are grouped into one or more
memory reference streams, each memory reference stream
having a leading memory reference, a trailing memory refer-
ence, and one or more middle memory references. Trans-
forming of the original computer code comprises inserting,
into the original computer code, instructions to execute ini-
tialization, lookup, and cleanup operations associated with
the leading memory reference and trailing memory reference
in a different manner from initialization, lookup, and cleanup
operations for the one or more middle memory references.

19 Claims, 13 Drawing Sheets



(56)

References Cited

U.S. PATENT DOCUMENTS

2006/0112237 A1 5/2006 Chen et al.
 2006/0212440 A1 9/2006 Heishi et al.
 2006/0265568 A1 11/2006 Burton
 2007/0260846 A1 11/2007 Burton
 2007/0261042 A1 11/2007 Chen et al.
 2007/0283098 A1 12/2007 O'Brien et al.
 2009/0158248 A1 6/2009 Linderman et al.
 2009/0248985 A1 10/2009 Ayguade et al.
 2009/0249318 A1 10/2009 Ayguade et al.

OTHER PUBLICATIONS

Antonio González, A Data Cache with Multiple Caching Strategies Tuned to Different Types of Locality, ACM 1995 International Conference on Supercomputing, Barcelona, Spain, ISBN: 0-89791-728-6, pp. 338-347.*

Eichenberger, Alexandre E. et al., "Using advanced compiler technology to exploit the performance of the Cell Broadband Engine architecture", IBM Systems Journal, vol. 45, No. 1, 2006, pp. 59-84.
 Fryman, Joshua B. et al., "SoftCache: A Technique for Power and Area Reduction in Embedded Systems", GTECH Tech Report GIT-CERCS-03-06, 12 pages.

Gonzalez, Marc et al., "Hybrid Access-Specific Software Cache Techniques for the Cell BE Architecture", Proceedings of the Seventeenth International Conference on Parallel Architectures and Compilation Techniques (PACT'08), Toronto, Canada, Oct. 25-29, 2008, 11 pages.

Jacob, Bruce et al., "Software-Managed Address Translation", IEEE, Proceedings of the Third International Symposium on High Performance Computer Architecture, Feb. 1-5, 1997, 12 pages.

U.S. Appl. No. 12/057,447, filed Mar. 28, 2008, Ayguade et al.

U.S. Appl. No. 12/057,456, filed Mar. 28, 2008, Ayguade et al.

Miller, Jason E. et al., "Software-based Instruction Caching for Embedded Processors", ACM, ASPLOS'06, Oct. 21-25, 2006, pp. 293-302.

Moritz, Csaba A. et al., "Hot Pages: Software Caching for Raw Microprocessors", MIT-LCS Technical Memo LCS-TM-599, Aug. 1999, 12 pages.

Witchel, Emmett et al., "Direct Addressed Caches for Reduced Power Consumption", 34th Annual International Symposium on Microarchitecture, MICRO-34, Austin, Texas, Dec. 2001, 11 pages.
 Balart, Jairo et al., "A Novel Asynchronous Software Cache Implementation for the Cell-BE Processor", Proceedings of the Languages and Compilers for Parallel Computing, 20th International Workshop, LCPC 2007, Oct. 11-13, 2007, pp. 125-140.

U.S. Appl. No. 12/057,447.

U.S. Appl. No. 12/057,456.

Final Office Action mailed Jan. 3, 2012 for U.S. Appl. No. 12/057,447, 32 pages.

Final Office Action mailed Dec. 29, 2011 for U.S. Appl. No. 12/057,456, 30 pages.

Response to Office Action filed with the USPTO on Oct. 17, 2011 for U.S. Appl. No. 12/057,456, 19 pages.

Response to Office Action filed with the USPTO on Oct. 20, 2011 for U.S. Appl. No. 12/057,447, 25 pages.

Advisory Action mailed Mar. 19, 2012 for U.S. Appl. No. 12/057,447, 5 pages.

Advisory Action mailed Mar. 19, 2012 for U.S. Appl. No. 12/057,456, 4 pages.

Interview Summary mailed Mar. 5, 2012 for U.S. Appl. No. 12/057,447, 3 pages.

Interview Summary mailed Mar. 5, 2012 for U.S. Appl. No. 12/057,456, 3 pages.

Response to Final Office Action filed with the USPTO on Feb. 29, 2012 for U.S. Appl. No. 12/057,456, 23 pages.

Response to Final Office Action filed with the USPTO on Mar. 2, 2012 for U.S. Appl. No. 12/057,447, 27 pages.

Al-Sukhni, Hassan F., "Identifying and Exploiting Memory Access Characteristics for Prefetching Linked Data Structures", Thesis submitted to the Faculty of the Graduate School of the University of Colorado, 2006, 9 pages.

Office Action mailed Jun. 7, 2012 for U.S. Appl. No. 12/057,447, 31 pages.

Office Action mailed Jun. 22, 2012 for U.S. Appl. No. 12/057,456, 29 pages.

Final Office Action dated Nov. 20, 2012 for U.S. Appl. No. 12/057,447, 28 pages.

Final Office Action mailed Dec. 6, 2012, for U.S. Appl. No. 12/057,456, 25 pages.

Response to Office Action filed with the USPTO on Sep. 7, 2012 for U.S. Appl. No. 12/057,447, 31 pages.

Response to Office Action filed with the USPTO on Sep. 21, 2012 for U.S. Appl. No. 12/057,456, 30 pages.

Appeal Brief filed with the USPTO on Apr. 1, 2013 for U.S. Appl. No. 12/057,447, 41 pages.


Appeal Brief filed with the USPTO on Apr. 30, 2013 for U.S. Appl. No. 12/057,456, 35 pages.

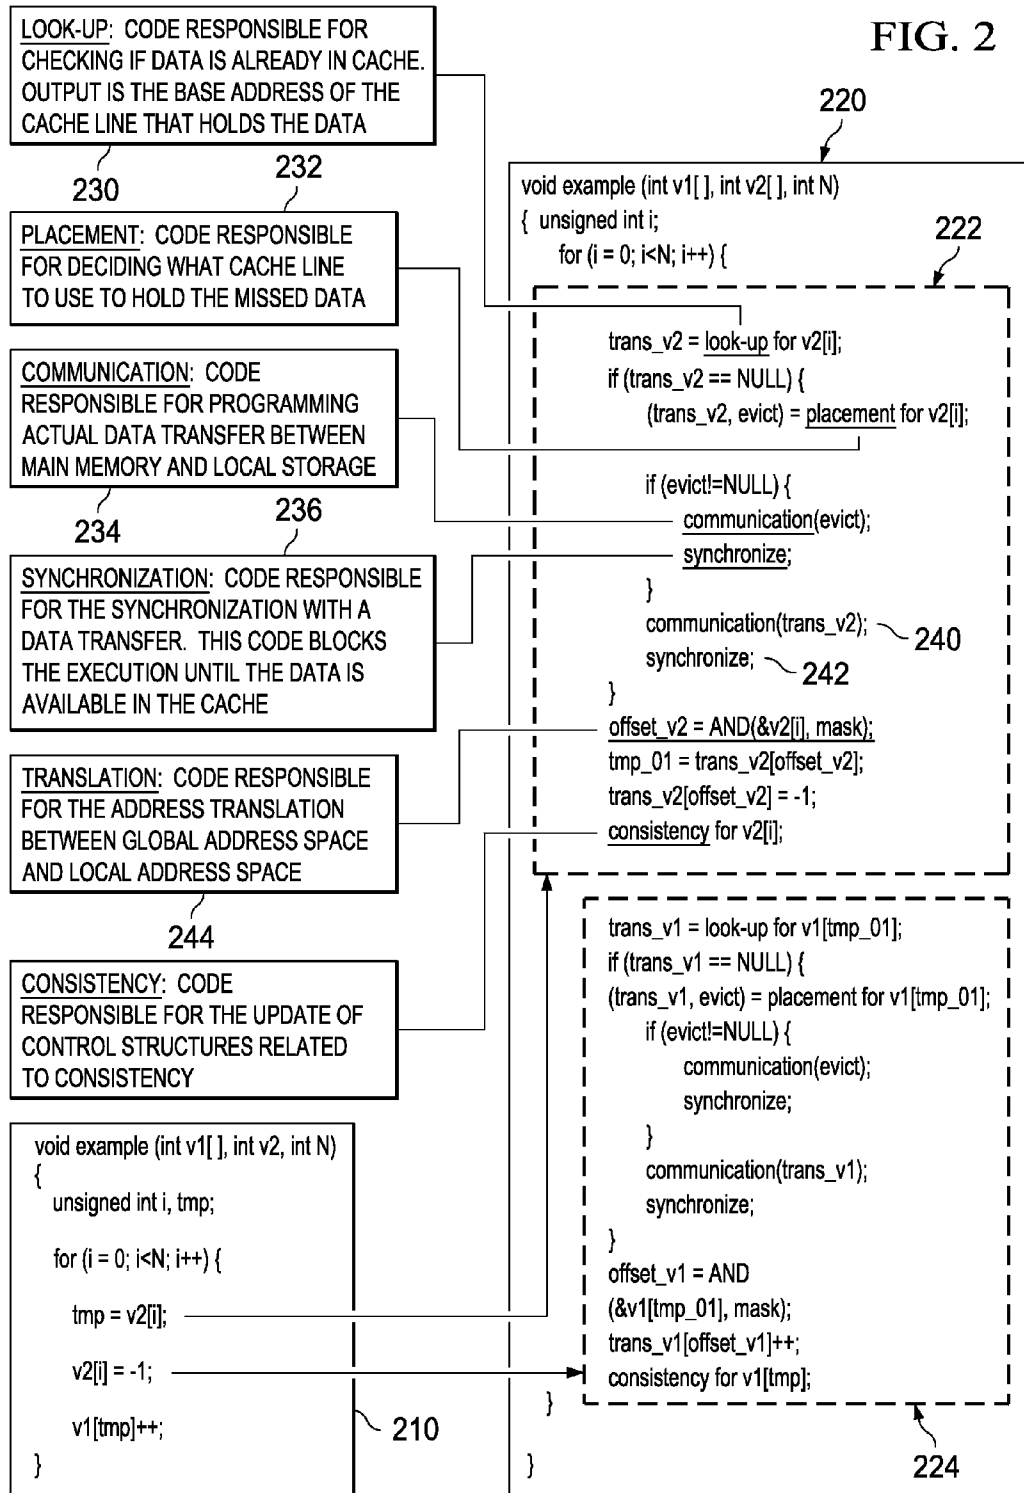
Notice of Allowance mailed Apr. 29, 2013 for U.S. Appl. No. 12/057,447, 13 pages.

Notice of Allowance mailed Jun. 7, 2013 for U.S. Appl. No. 12/057,456, 13 pages.

* cited by examiner

FIG. 1








FIG. 3

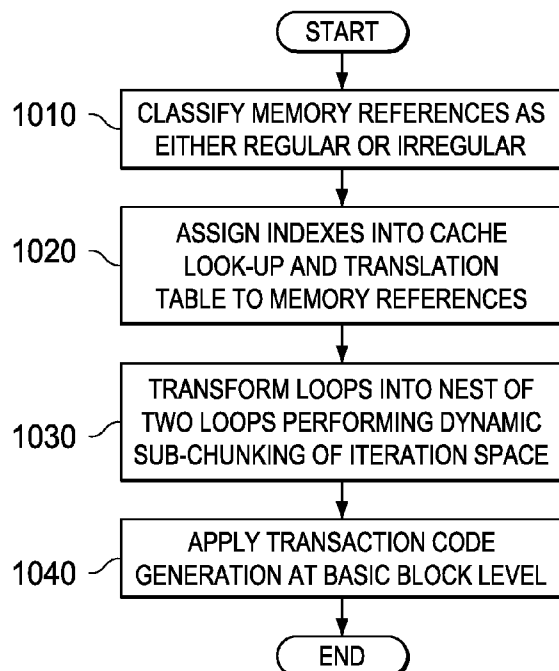


FIG. 10


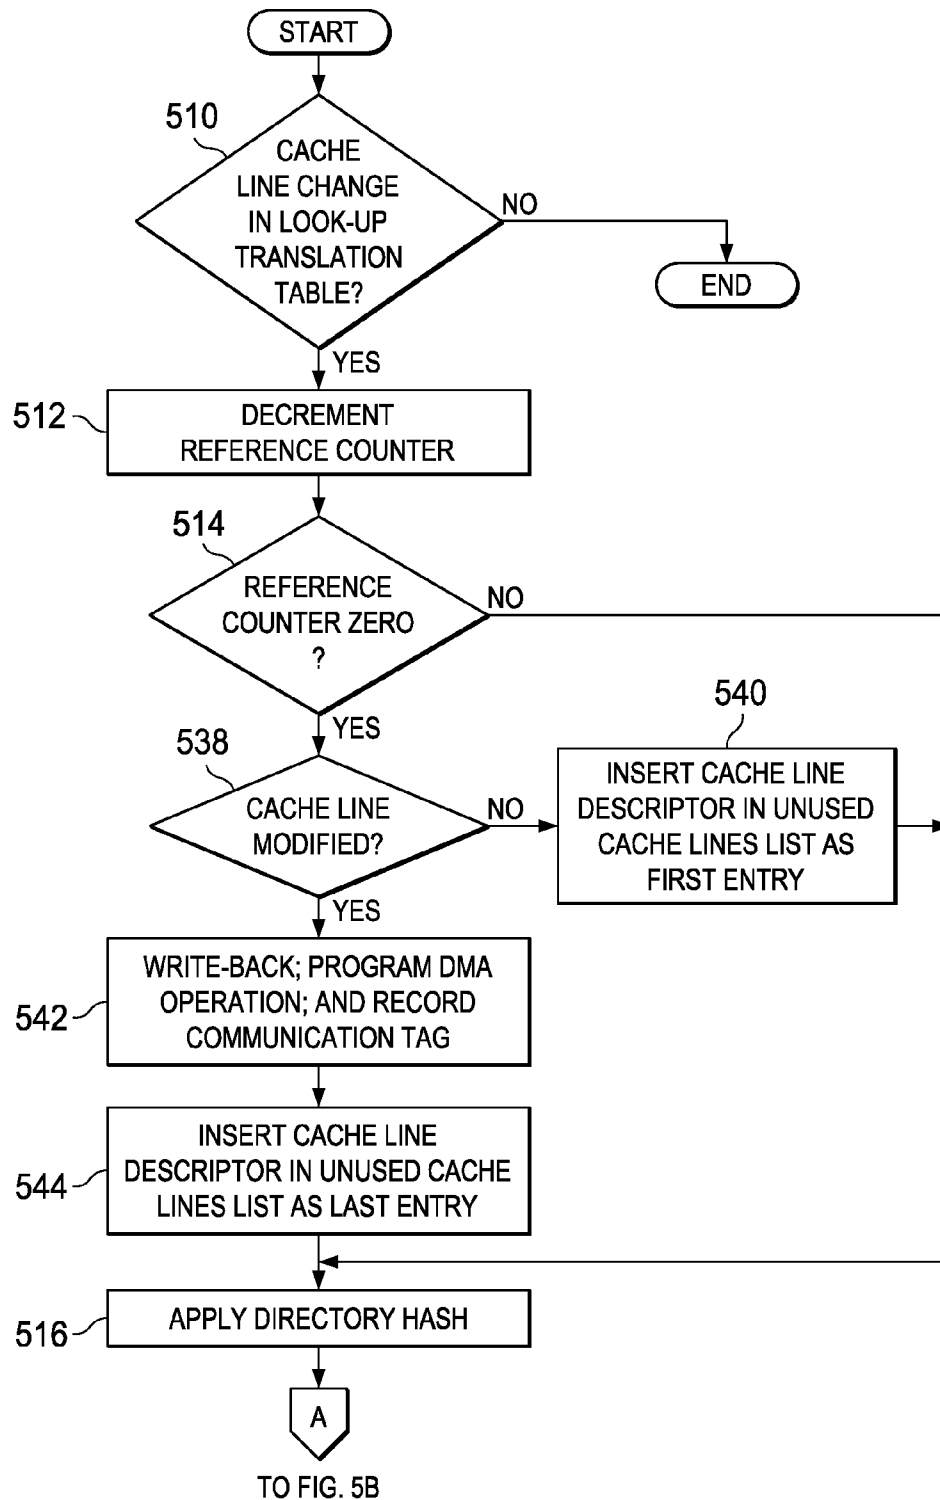


FIG. 4

FIG. 5A



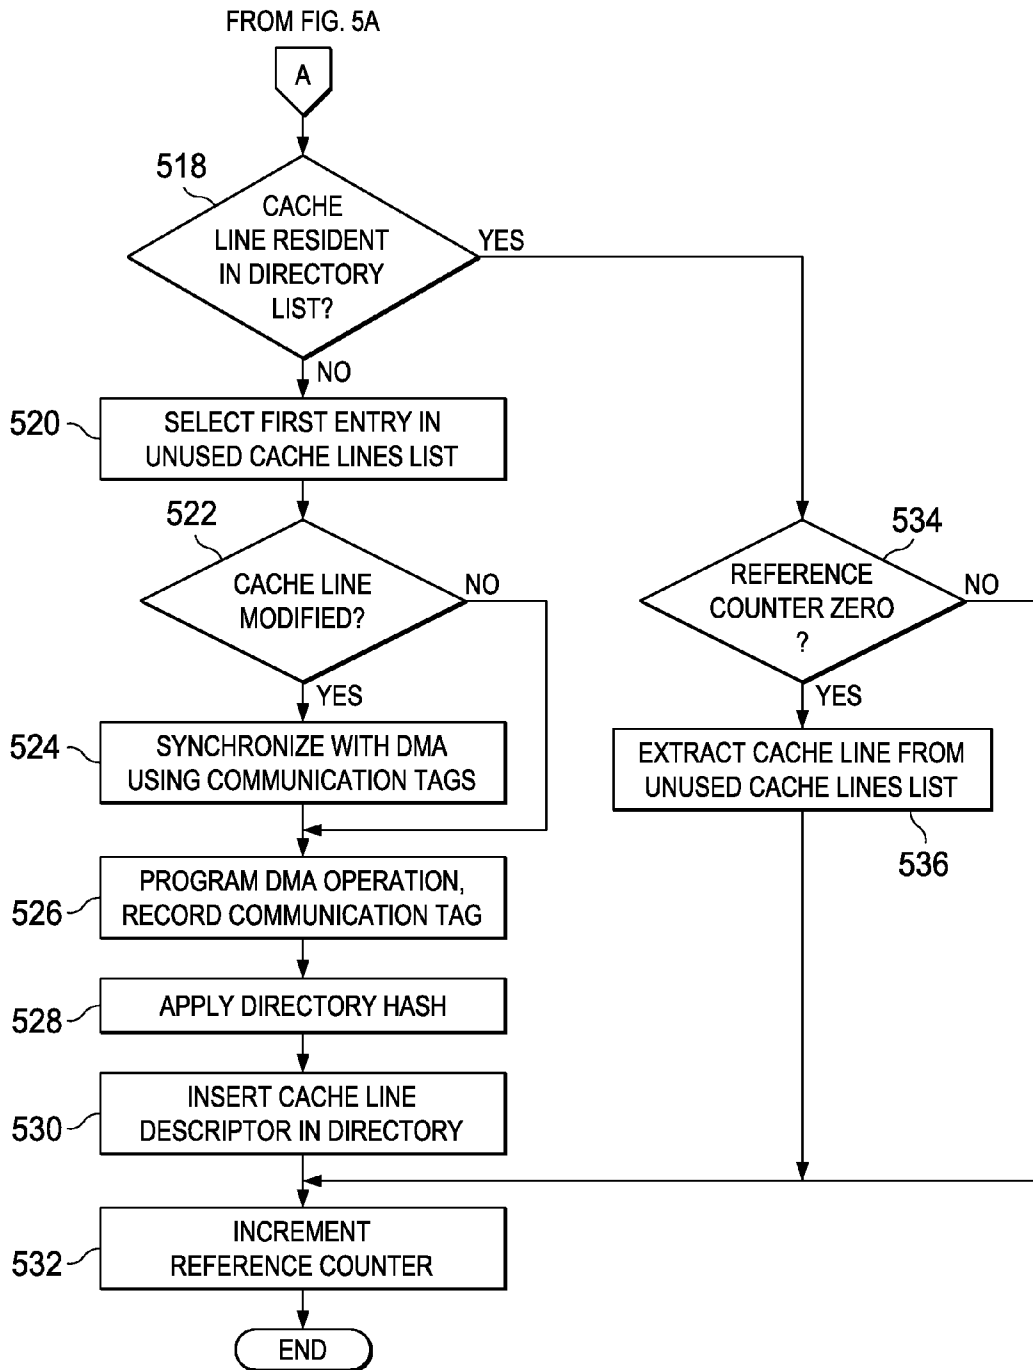


FIG. 5B




FIG. 6

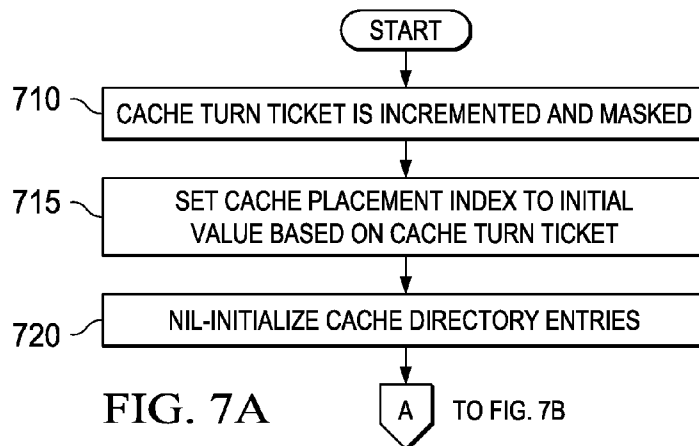


FIG. 7A




FIG. 7B

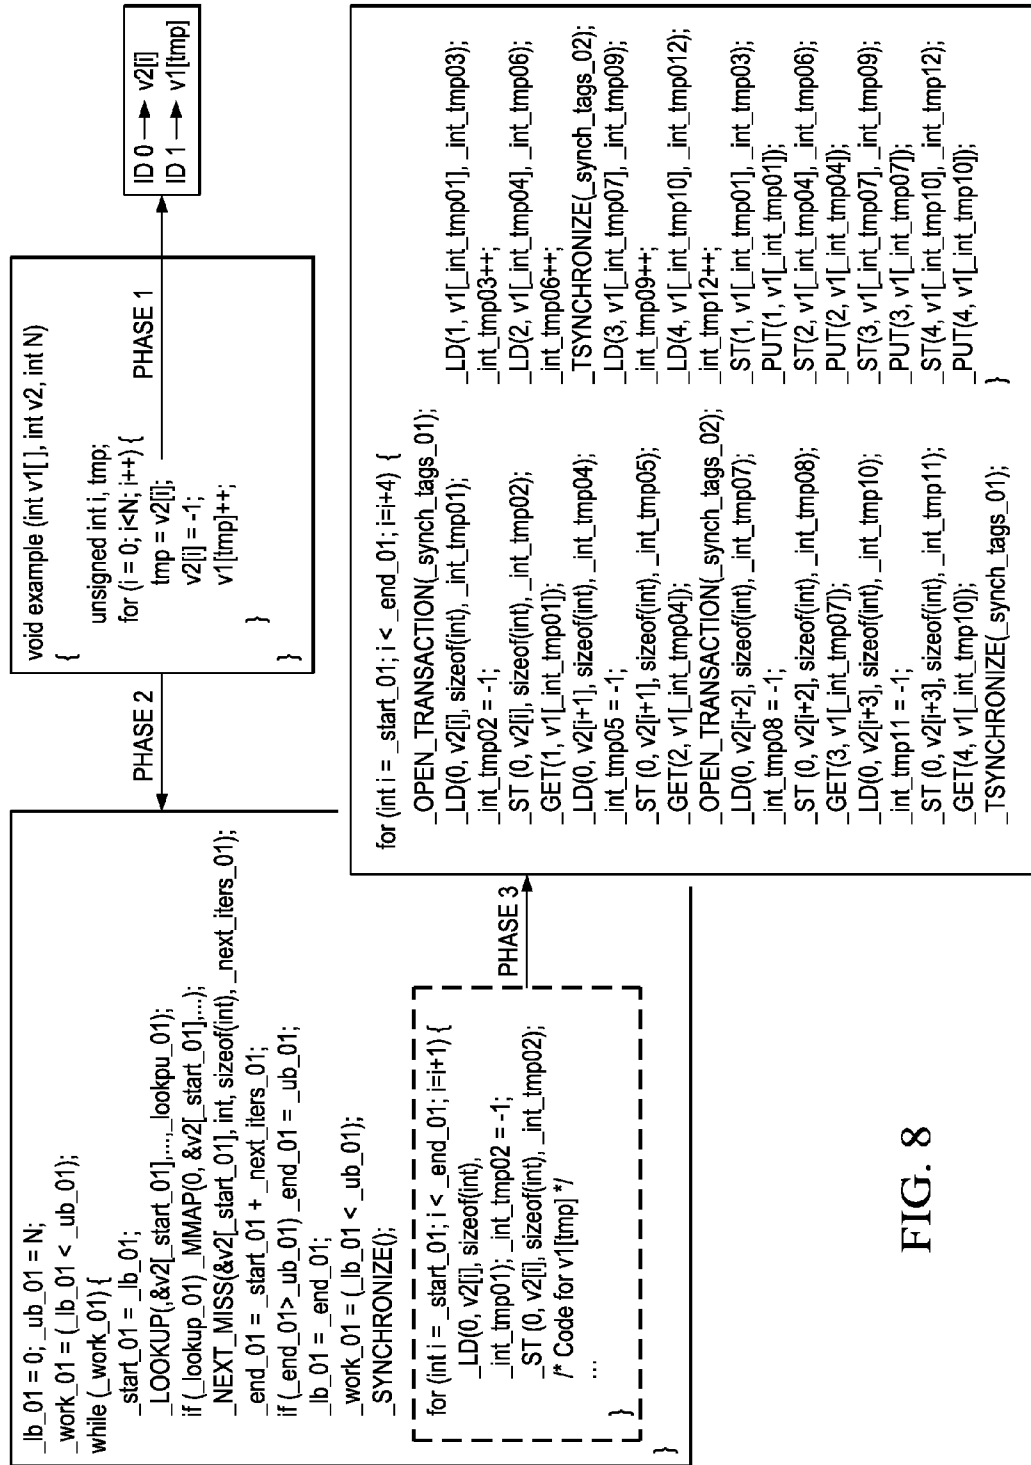


FIG. 8

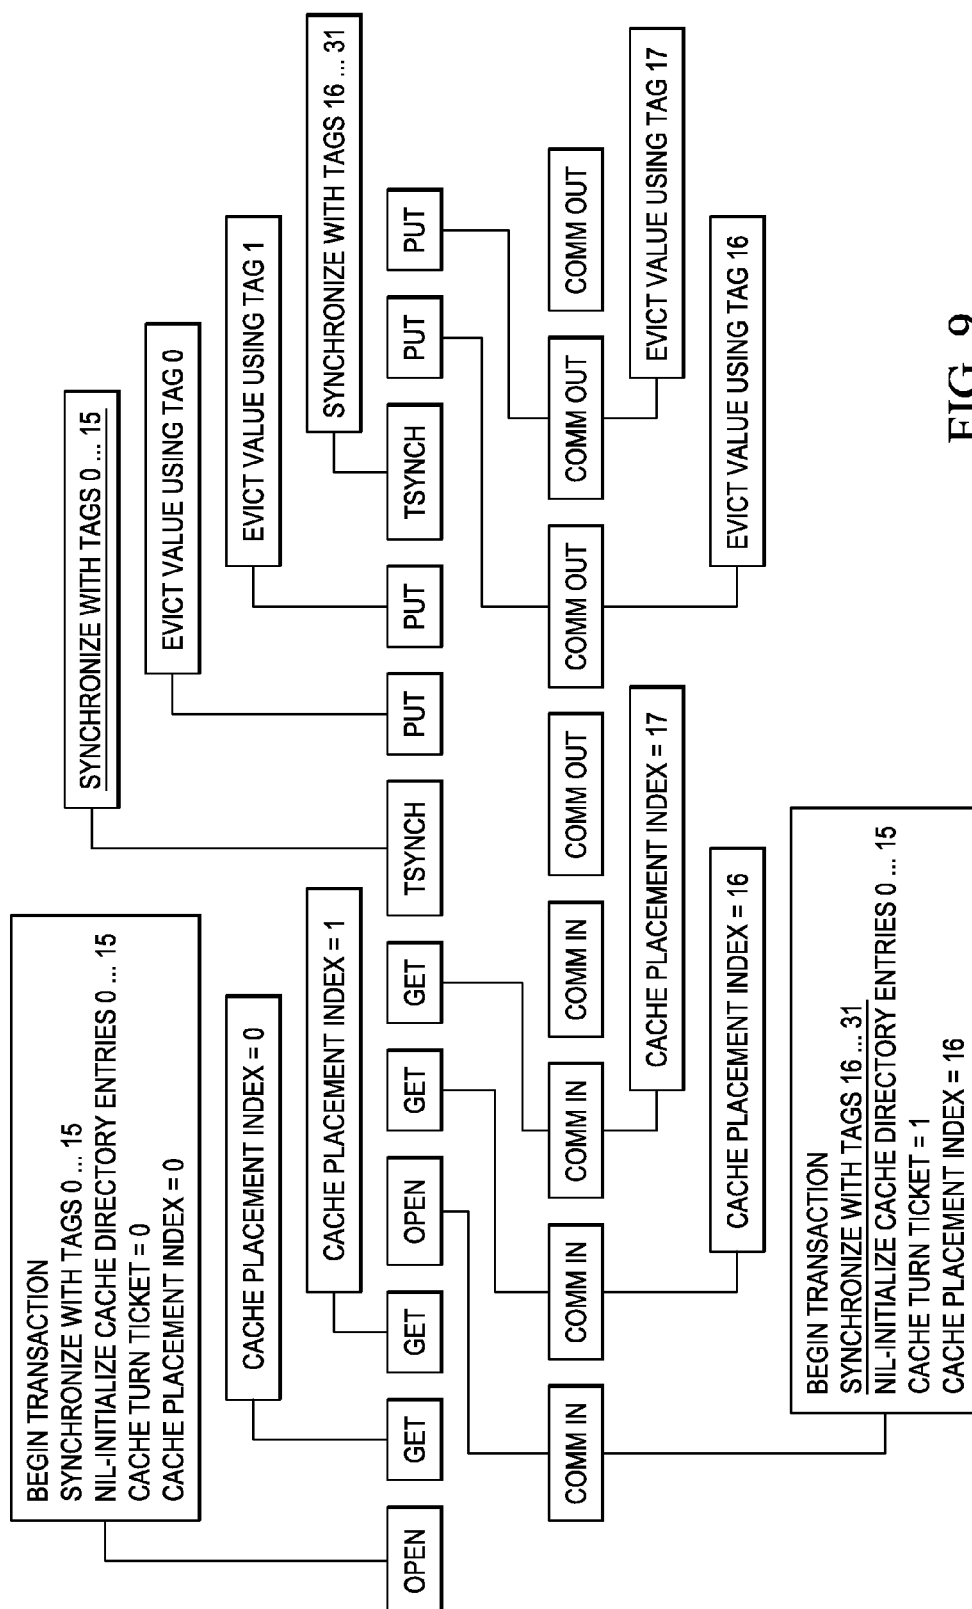


FIG. 11A

```
for(i=0; i<N; i++) {  
    .... a[x*i+y] ....  
}
```

FIG. 11B





FIG. 12



		1310	1320	1330	
		ref	a[i+3] (leading)	a[i-1] (trailing)	a[i+1] (middle)
1340	init		h1 = Map()	h3 =Lookup() if (!h3) h3 = Map();	h2 = Lookup()
1350	lookup (steady state)		if (!NumInLine()) h1 = Map() trip = Min(trip,...)	if (!NumInLine()) Free(h3) h3 = h1 trip = Min(trip,...)	if (!NumInLine()) h2 = h1 trip = Min(trip,...)
1360	clean		Free(h1)	if (h1!=h3) Free(h3)	

FIG. 13




FIG. 14

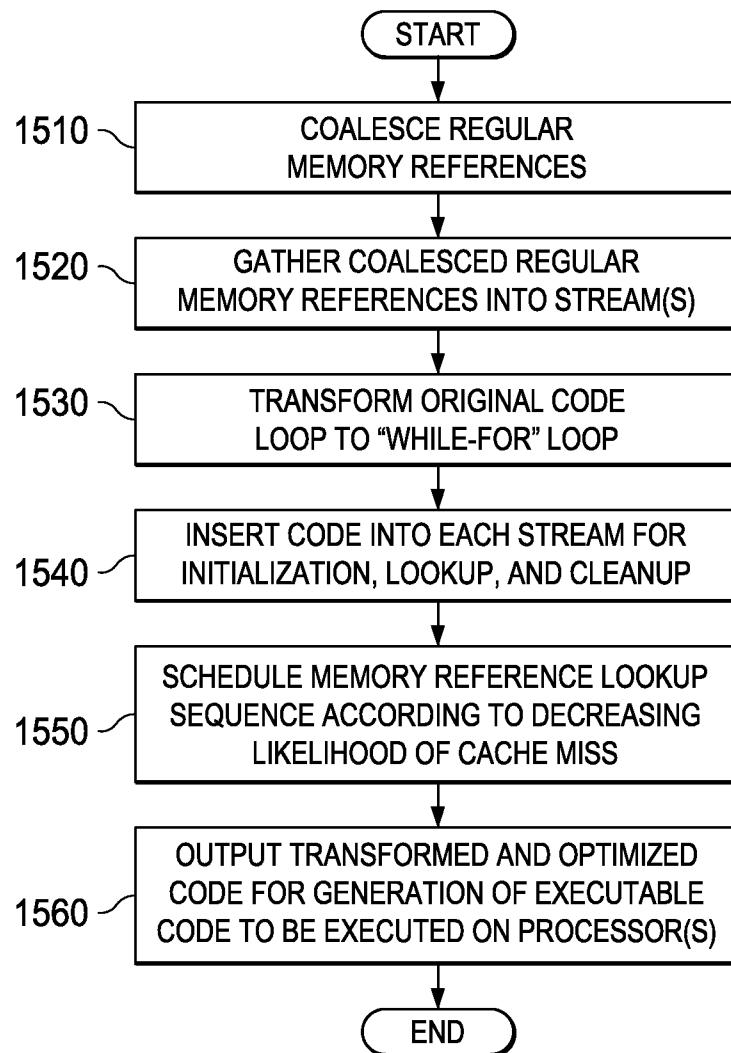


FIG. 15

1

OPTIMIZED CODE GENERATION TARGETING A HIGH LOCALITY SOFTWARE CACHE

BACKGROUND OF THE INVENTION

1. Field of the Invention

The present application relates generally to an improved data processing apparatus and method and more specifically to mechanisms for generating optimized code that targets a high locality software cache.

2. Background of the Invention

Multi-core systems are becoming more prevalent in today's computing environments. A multi-core system combines two or more independent cores, or processors, into a single package composed of either a single integrated circuit (IC) or multiple ICs packaged together. For example, a dual-core system contains two cores while a quad-core processor contains four cores. Cores in such a multi-core system may share a single coherent cache at the highest on-device cache level (e.g., a L2 cache for the Intel Core 2) or may have separate caches (e.g., the current AMD dual-core processors). The processors, or cores, also share the same interconnect to the rest of the system. Each core independently implements optimizations such as superscalar execution, pipelining, and multithreading. The most commercially significant multi-core processors are those used in personal computers and gaming consoles, e.g., the Cell Broadband Engine (CBE) available from International Business Machines Corporation of Armonk, N.Y., which is presently used in the Playstation 3 gaming console available from Sony Corporation.

The amount of performance gained by the use of a multi-core system depends on the problems being solved and the algorithms used, as well as their implementation in software. For example, for some parallel problems, a dual-core processor with two cores running at 2 GHz may perform very nearly as fast as a single core of 4 GHz. However, other problems may not yield as much of a speed-up from the use of multiple cores. Even if such a speed-up is not achieved, the system will typically perform multitasking more efficiently since it can run two or more programs at once, one on each core.

Ease of programming is one of the main impediments for the broad acceptance of multi-core systems. This is because present multi-core systems do not have hardware support for transparent data transfer between local and global memories. To address this issue, software caches have been used as a robust approach to provide the user with a transparent view of the memory architecture. A software cache is a hardware cache that is managed by software. While software caches allow local and global memories to be viewed together as a single memory device, software cache approaches can suffer from poor performance for a variety of reasons.

BRIEF SUMMARY OF THE INVENTION

In one illustrative embodiment, a method, in a data processing system, is provided for optimizing regular memory references in original computer code. The method may comprise parsing the computer code to identify memory references in the computer code. The method may further comprise classifying the memory references in the computer code as either regular memory references or irregular memory references. Accesses to a software cache by regular memory references may be controlled by a high locality cache mechanism. Moreover, the method may comprise transforming the original computer code to generate transformed computer code in which the regular memory references are grouped into

2

one or more memory reference streams, each memory reference stream having a leading memory reference, a trailing memory reference, and one or more middle memory references. The transforming of the original computer code may comprise inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations associated with the leading memory reference and trailing memory reference in a different manner from initialization, lookup, and cleanup operations for the one or more middle memory references.

In other illustrative embodiments, a computer program product comprising a computer useable or readable medium having a computer readable program is provided. The computer readable program, when executed on a computing device, causes the computing device to perform various ones, and combinations of, the operations outlined above with regard to the method illustrative embodiment.

In yet another illustrative embodiment, a system/apparatus is provided. The system/apparatus may comprise one or more processors and a memory coupled to the one or more processors. The memory may comprise instructions which, when executed by the one or more processors, cause the one or more processors to perform various ones, and combinations of, the operations outlined above with regard to the method illustrative embodiment.

These and other features and advantages of the present invention will be described in, or will become apparent to those of ordinary skill in the art in view of, the following detailed description of the example embodiments of the present invention.

BRIEF DESCRIPTION OF THE SEVERAL VIEWS OF THE DRAWINGS

The invention, as well as a preferred mode of use and further objectives and advantages thereof, will best be understood by reference to the following detailed description of illustrative embodiments when read in conjunction with the accompanying drawings, wherein:

FIG. 1 is an exemplary block diagram of a Cell Broadband Engine architecture in accordance with one illustrative embodiment;

FIG. 2 is an example of a portion of code and its transformation for ensuring proper access of a distributed memory architecture;

FIG. 3 is an exemplary block diagram of an overall software cache design architecture in accordance with one illustrative embodiment;

FIG. 4 is an exemplary block diagram of the high locality cache of the software cache design in accordance with one illustrative embodiment;

FIGS. 5A and 5B are a flowchart outlining an exemplary operation of a high locality cache in accordance with one illustrative embodiment;

FIG. 6 is an exemplary block diagram of a transactional cache in accordance with one illustrative embodiment;

FIGS. 7A and 7B are a flowchart outlining an exemplary operation of a transactional cache in accordance with one illustrative embodiment;

FIG. 8 is an exemplary diagram of an executable code generation by a compiler in accordance with one illustrative embodiment;

FIG. 9 is an exemplary diagram illustrating how the transactional cache structures are modified along one iteration of the unrolled loop of FIG. 8;

FIG. 10 is a flowchart outlining an exemplary operation of a compiler in accordance with one illustrative embodiment;

FIG. 11A illustrates a very simple portion of a code loop;
FIG. 11B illustrates a basic template for transformed code
corresponding to the code loop in FIG. 11A in accordance
with one illustrative embodiment;

FIG. 12 illustrates a coalescing of memory references to
generate pairs of memory references in which, if one memory
reference is present in a cache line, the other memory refer-
ence is guaranteed to be within the cache line, in accordance
with one illustrative embodiment;

FIG. 13 is a table of code to be inserted for leading, trailing,
and middle memory references in accordance with one illus-
trative embodiment;

FIG. 14 is an example diagram that illustrates how the
order in which a memory reference is checked may influence
computation/communication overlap; and

FIG. 15 illustrates an example operation for optimizing
memory reference checks for regular memory references in a
code loop of original code in accordance with one illustrative
embodiment.

DETAILED DESCRIPTION OF THE INVENTION

As discussed above, the Cell Broadband Engine (CBE)
processor is one of the recent chip multiprocessor architec-
tures with application specific processor features. Moreover,
as mentioned above, in the CBE processor architecture, soft-
ware is responsible for initiating DMA transfers. In so doing,
the software is further responsible for ensuring memory con-
sistency and proper address translation.

In an architecture where software is responsible for data
transfers, memory consistency, and address translation, the
performance of the system is very sensitive to the compiler
technology. The compiler has to address performance in three
specific areas. First, each memory reference has to be moni-
tored to ensure correct translation from a global address space
to a local store address space. This monitoring code repre-
sents an overhead and thus, it is important to minimize the
impact of this monitoring code on the performance of the
system. Second, communication must be scheduled to mini-
mize the cost of data transfers. Thus, the ability to overlap
computation with communication becomes an attractive opti-
mization. Third, memory consistency must be maintained by
software which again is an important source of overhead that
needs to be addressed.

In order to address these various areas of potential perfor-
mance degradation, it is not possible to rely entirely on the
compiler for optimization. The runtime system must also ease
the compiler's job by including the necessary capabilities to
enhance the chances for optimization. The illustrative
embodiments set forth herein provide software caching
mechanisms that provide a number of contributions for
addressing these issues with regard to programmability of
processors having distributed memory architectures, such as
the CBE processor architecture.

As will be appreciated by one skilled in the art, the present
invention may be embodied as a system, method or computer
program product. Accordingly, the present invention may
take the form of an entirely hardware embodiment, an entirely
software embodiment (including firmware, resident software,
micro-code, etc.) or an embodiment combining software and
hardware aspects that may all generally be referred to herein
as a "circuit," "module" or "system." Furthermore, the
present invention may take the form of a computer program
product embodied in any tangible medium of expression hav-
ing computer usable program code embodied in the medium.

Any combination of one or more computer usable or com-
puter readable medium(s) may be utilized. The computer-

usable or computer-readable medium may be, for example
but not limited to, an electronic, magnetic, optical, electro-
magnetic, infrared, or semiconductor system, apparatus,
device, or propagation medium. More specific examples (a
non-exhaustive list) of the computer-readable medium would
include the following: an electrical connection having one or
more wires, a portable computer diskette, a hard disk, a ran-
dom access memory (RAM), a read-only memory (ROM), an
erasable programmable read-only memory (EPROM or Flash
memory), an optical fiber, a portable compact disc read-only
memory (CDROM), an optical storage device, a transmission
media such as those supporting the Internet or an intranet, or
a magnetic storage device. Note that the computer-usable or
computer-readable medium could even be paper or another
suitable medium upon which the program is printed, as the
program can be electronically captured, via, for instance,
optical scanning of the paper or other medium, then com-
piled, interpreted, or otherwise processed in a suitable man-
ner, if necessary, and then stored in a computer memory. In the
context of this document, a computer-usable or computer-
readable medium may be any medium that can contain, store,
communicate, propagate, or transport the program for use by
or in connection with the instruction execution system, appa-
ratus, or device. The computer-usable medium may include a
propagated data signal with the computer-usable program
code embodied therewith, either in baseband or as part of a
carrier wave. The computer usable program code may be
transmitted using any appropriate medium, including but not
limited to wireless, wireline, optical fiber cable, RF, etc.

Computer program code for carrying out operations of the
present invention may be written in any combination of one or
more programming languages, including an object oriented
programming language such as Java, Smalltalk, C++ or the
like and conventional procedural programming languages,
such as the "C" programming language or similar program-
ming languages. The program code may execute entirely on
the user's computer, partly on the user's computer, as a stand-
alone software package, partly on the user's computer and
partly on a remote computer or entirely on the remote com-
puter or server. In the latter scenario, the remote computer
may be connected to the user's computer through any type of
network, including a local area network (LAN) or a wide area
network (WAN), or the connection may be made to an external
computer (for example, through the Internet using an
Internet Service Provider).

The illustrative embodiments are described below with
reference to flowchart illustrations and/or block diagrams of
methods, apparatus (systems) and computer program prod-
ucts according to the illustrative embodiments of the inven-
tion. It will be understood that each block of the flowchart
illustrations and/or block diagrams, and combinations of
blocks in the flowchart illustrations and/or block diagrams,
can be implemented by computer program instructions. These
computer program instructions may be provided to a
processor of a general purpose computer, special purpose
computer, or other programmable data processing apparatus
to produce a machine, such that the instructions, which
execute via the processor of the computer or other program-
mable data processing apparatus, create means for imple-
menting the functions/acts specified in the flowchart and/or
block diagram block or blocks.

These computer program instructions may also be stored in
a computer-readable medium that can direct a computer or
other programmable data processing apparatus to function in
a particular manner, such that the instructions stored in the
computer-readable medium produce an article of manufac-

5

ture including instruction means which implement the function/act specified in the flowchart and/or block diagram block or blocks.

The computer program instructions may also be loaded onto a computer or other programmable data processing apparatus to cause a series of operational steps to be performed on the computer or other programmable apparatus to produce a computer implemented process such that the instructions which execute on the computer or other programmable apparatus provide processes for implementing the functions/acts specified in the flowchart and/or block diagram block or blocks.

The flowchart and block diagrams in the Figures illustrate the architecture, functionality, and operation of possible implementations of systems, methods and computer program products according to various embodiments of the present invention. In this regard, each block in the flowchart or block diagrams may represent a module, segment, or portion of code, which comprises one or more executable instructions for implementing the specified logical function(s). It should also be noted that, in some alternative implementations, the functions noted in the block may occur out of the order noted in the figures. For example, two blocks shown in succession may, in fact, be executed substantially concurrently, or the blocks may sometimes be executed in the reverse order, depending upon the functionality involved. It will also be noted that each block of the block diagrams and/or flowchart illustration, and combinations of blocks in the block diagrams and/or flowchart illustration, can be implemented by special purpose hardware-based systems that perform the specified functions or acts, or combinations of special purpose hardware and computer instructions.

In order to understand the advances provided by the mechanisms of the illustrative embodiments, it is first important to understand a representative distributed memory architecture in which the mechanisms of the illustrative embodiments may be implemented. Thus, FIG. 1 is provided as an exemplary block diagram of a data processing system in which aspects of the illustrative embodiments may be implemented. The exemplary data processing system shown in FIG. 1 is an example of the Cell Broadband Engine (CBE) data processing system. While the CBE will be used in the description of the illustrative embodiments of the present invention, the present invention is not limited to such, as will be readily apparent to those of ordinary skill in the art upon reading the following description. The mechanisms of the illustrative embodiments may be applied to any distributed memory architecture and are not limited to use with the CBE processor architecture.

As shown in FIG. 1, the CBE 100 includes a power processor element (PPE) 110 having a processor (PPU) 116 and its L1 and L2 caches 112 and 114, and multiple synergistic processor elements (SPEs) 120-134 that each has its own synergistic processor unit (SPU) 140-154, memory flow control 155-162, local memory or store (LS) 163-170, and bus interface unit (BIU unit) 180-194 which may be, for example, a combination direct memory access (DMA), memory management unit (MMU), and bus interface unit. A high bandwidth internal element interconnect bus (EIB) 196, a bus interface controller (BIC) 197, and a memory interface controller (MIC) 198 are also provided.

The local memory or local store (LS) 163-170 is a non-coherent addressable portion of a large memory map which, physically, may be provided as small memories coupled to the SPU 140-154. The local stores 163-170 may be mapped to different address spaces. These address regions are continuous in a non-aliased configuration. A local store 163-170 is

6

associated with its corresponding SPU 140-154 and SPE 120-134 by its address location, such as via the SPU Identification Register, described in greater detail hereafter. Any resource in the system has the ability to read/write from/to the local store 163-170 as long as the local store is not placed in a secure mode of operation, in which case only its associated SPU may access the local store 163-170 or a designated secured portion of the local store 163-170.

The CBE 100 may be a system-on-a-chip such that each of the elements depicted in FIG. 1 may be provided on a single microprocessor chip. Moreover, the CBE 100 is a heterogeneous processing environment in which each of the SPU's may receive different instructions from each of the other SPU's in the system. Moreover, the instruction set for the SPU's is different from that of the PPU, e.g., the PPU may execute Reduced Instruction Set Computer (RISC) based instructions while the SPU execute vectorized instructions.

The SPEs 120-134 are coupled to each other and to the L2 cache 114 via the EIB 196. In addition, the SPEs 120-134 are coupled to MIC 198 and BIC 197 via the EIB 196. The MIC 198 provides a communication interface to shared memory 199. The BIC 197 provides a communication interface between the CBE 100 and other external buses and devices.

The PPE 110 is a dual threaded PPE 110. The combination of this dual threaded PPE 110 and the eight SPEs 120-134 makes the CBE 100 capable of handling 10 simultaneous threads and over 128 outstanding memory requests. The PPE 110 acts as a controller for the other eight SPEs 120-134 which handle most of the computational workload. The PPE 110 may be used to run conventional operating systems while the SPEs 120-134 perform vectorized floating point code execution, for example.

The SPEs 120-134 comprise a synergistic processing unit (SPU) 140-154, memory flow control units 155-162, local memory or store 163-170, and an interface unit 180-194. The local memory or store 163-170, in one exemplary embodiment, comprises a 256 KB instruction and data memory which is visible to the PPE 110 and can be addressed directly by software.

The PPE 110 may load the SPEs 120-134 with small programs or threads, chaining the SPEs together to handle each step in a complex operation. For example, a set-top box incorporating the CBE 100 may load programs for reading a DVD, video and audio decoding, and display, and the data would be passed off from SPE to SPE until it finally ended up on the output display. At 4 GHz, each SPE 120-134 gives a theoretical 32 GFLOPS of performance with the PPE 110 having a similar level of performance.

The memory flow control units (MFCs) 155-162 serve as an interface for an SPU to the rest of the system and other elements. The MFCs 155-162 provide the primary mechanism for data transfer, protection, and synchronization between main storage and the local storages 163-170. There is logically an MFC for each SPU in a processor. Some implementations can share resources of a single MFC between multiple SPU's. In such a case, all the facilities and commands defined for the MFC must appear independent to software for each SPU. The effects of sharing an MFC are limited to implementation-dependent facilities and commands.

As mentioned above, in distributed memory architectures, such as the Cell Broadband Engine architecture illustrated in FIG. 1 above, software is primarily used to monitor memory references, schedule data transfers, and the like, to transparently manage the distributed memory architecture. Because of this software dependency, relatively simple source code that accesses the distributed memory architectures is trans-

formed into more complex code that includes a large amount of checking and management code for checking correct performance of, and managing, accesses to the distributed memory architecture. This checking and management code represents a large amount of overhead that degrades the performance of the system as a whole.

For example, in distributed memory architectures, such as the CBE **100**, there is very little if any hardware support to supply the data referenced in load/store operations and thus, every memory reference in the code executed in association with the distributed memory architecture is surrounded by checking code to ensure correctness in the memory access. Some checking code is devoted to checking if the required data is already in the cache storage, hereafter referred to as the “look-up” code. In case the look-up code determines that the data is not in the cache storage, i.e. a cache miss occurs, other code is executed that is devoted to determining the placement of the new data about to be brought into the cache storage, hereafter referred to as the “placement” code. Because of the space limitation in the cache storage, other code is introduced in order to send back to main memory the modified data no longer required by the computation.

Moreover, control information, e.g., dirty bits and the like, is utilized to keep track of what data has been modified and additional code is used to update this control information, hereafter referred to as “consistency” code. Furthermore, code for performing Direct Memory Access (DMA) operations is introduced to bring in data into the cache storage and to send back data to the main memory, hereafter referred to as “communication” code.

Using the CBE **100** as an example of a distributed memory architecture, because of the local address space in local stores **163-170** of the SPEs **120-134**, code is used to translate from the global address space to the local address space of the SPE **120-134**, hereafter referred to as “translation” code. Finally, code for synchronizing with the DMA engine is used to ensure that the data transfers are completed and the memory accesses can be safely executed, hereafter referred to as “synchronization” code.

FIG. **2** is an example of a portion of code and its transformation for ensuring proper access of a distributed memory architecture. As shown in FIG. **2**, a portion of C code **210** to be executed by an SPE is transformed, by a compiler (not shown), into transformed code **220** to correctly address each memory reference. The difference in code **210** and transformed code **220** illustrates the difficulty for a compiler to generate efficient code.

As shown in FIG. **2**, the original code **210** has two different memory references $v1[tmp]$ and $v2[i]$. In the transformed code **220**, two boxes **222** and **224** are shown to illustrate the corresponding portions of the transformed code **220** with the two memory references. Box **222** corresponds to the $v2[i]$ memory reference in the loop body of the original code **210** and box **224** corresponds to the $v1[tmp]$ memory reference in the loop body of the original code **210**.

Looking at the memory reference $v2[i]$ in box **222**, an if statement has been inserted in the transformed code **220** where the look-up code **230** determines the condition of the if statement. In case a cache miss occurs (trans_v2 equals NULL), the placement code **232** is executed and variable trans_v2 is updated with the base address of a cache line to where the recently missed data has been mapped. In case some data has to be sent back to main memory, the variable evict holds the base address of the cache line to be evicted. Another if statement controls the eviction event and the corresponding communication code **234** and synchronization code **236** is executed. Then additional communication code

240 and synchronization code **242** is executed to actually bring the data that was missed into the cache.

Finally, translation code **244** is provided that is composed of two computations. The first is a computation to determine the offset within the cache line and the second is code to add the offset to the cache line base address. Similar look-up, placement, communication, synchronization, and translation code has been introduced into the translated code **220** for the $v1[tmp]$ memory reference as well, as shown in box **224**.

The transformation of the original code **210** into the transformed code **220** in FIG. **2** is clearly not optimal in a number of different ways. First, the code transformation does not take into account the different locality that the two memory references are exposing. On one side, the $v2[i]$ reference shows excellent spatial locality, i.e. memory locations close to each other, such as in the same cache line, are likely to be accessed. Memory accesses or references, such as $v2[i]$ in the depicted example, that have a high degree of spatial locality are referred to herein as “regular” memory accesses or references. Such regular memory accesses or references may be characterized by having a constant stride, i.e. a constant number of locations in memory between successive elements, within a loop execution. Thus, “regular” memory accesses or references may be any memory reference in code that explicitly shows a strided access pattern at compile time (although the stride value itself need not be necessarily known at compile time).

Software caching techniques are a reliable solution for handling such regular memory references to achieve an efficient code. However, it would be desirable to remove some, if not all, the control code that surrounds a regular memory reference, e.g., the $v2[i]$ memory reference, by predicting the exact iteration where this memory access is going to cause a cache miss. For example, given the initial address of a memory access, the stride and the cache line size, it can be predicted how many iterations may be performed without having a cache miss. If it is known when a cache miss is going to be produced, one can execute the control code just for cache misses, i.e. at a cache line level, rather than for each iteration of a regular memory reference.

Thus, it would be desirable to have a code transformation that executes the control code at a cache line level, rather than a memory access level, so that this predictability may be exploited and used to optimize the control code. An immediate consequence of this, however, is that a configurable cache line size may be used so as to identify the largest cache line size that minimizes the execution of the control code. Moreover, since cache misses may not occur at the same iteration for every memory access in a computation, it is further desirable to have a mechanism to pin the cache lines in the cache storage until they are no longer in use by any memory reference, thereby excluding these cache lines from the cache’s replacement policy.

On another side, the memory access to $v1[tmp]$ does not expose any locality, i.e. there is no guarantee that $v1[x]$ is local in memory to $v1[y]$; there is no constant stride between successive $v1$ elements. Memory accesses or memory references that do not expose any locality are referred to herein as “irregular” memory accesses or references, i.e. memory accesses or memory references that are not “regular” memory accesses or references. For irregular memory accesses or references, it is reasonable to accept that the control code is going to be present at every iteration of a memory access/reference. However, depending on specific features in the software cache design, some optimizations may be done. First, it is not acceptable that the cache line size for this

memory access/reference type is similar to the one used in the v2[i] access. This is because the cache line size determines the amount of transferred data for a cache miss. If a very small element in the cache line is actually referenced, the ratio between the amount of transferred data and used data is very poor. Thus, it would be desirable to adapt the amount of transferred data according to the type of memory access, e.g., use a relatively small cache line size for memory references that do not expose locality and use a relatively larger cache line size for those memory reference that do expose locality. As will be described hereafter, the illustrative embodiments achieve such adaptability by permitting the compiler to determine where to map a memory reference with a high locality cache using a relatively larger cache line size and a transactional cache using a relatively smaller cache line size.

Second, the code transformation in FIG. 2 always forces an immediate synchronization between the execution flow and the programmed data transfers. This hinders any possible overlapping between computation and communication. The compiler should be allowed to reorganize the computation to maximize the chances for the overlapping of computation with communication. Such reorganization should try to group all the look-up, placement, and communication code and place it as far as possible from the actual use of the data and the synchronization before the data is used. However, it should be ensured that no possible conflict appears between the different placement actions. In this way, such reorganization makes control code overlap with data transfers originated in the control code.

The illustrative embodiments provide mechanisms for facilitating such transformations for the generation of efficient code for both types of memory access or references, i.e. regular and irregular memory accesses or references. The mechanisms of the illustrative embodiments provide a software cache design that is non-uniform with respect to how memory accesses/references are treated. The software cache design includes separate facilities for handling regular memory accesses or references and irregular memory accesses or references with these facilities being selectively utilized to handle the memory access/reference. A compiler or other entity involved in optimizing and generating executable code for one or more processors of a system implementing the software cache design of the illustrative embodiments, may take advantage of the mechanisms of the software cache design when optimizing and generating the executable code.

FIG. 3 is an exemplary block diagram of an overall software cache design architecture in accordance with one illustrative embodiment. As shown in FIG. 3, the software cache design 300 comprises three primary operational elements referred to as the high locality cache 310, the transactional cache 320, and the memory consistency engine 330. The high locality cache 310 is used to handle memory accesses/references identified by the compiler 340 as regular memory accesses/references. The transactional cache 320 is used to handle memory accesses/references identified by the compiler 340 as irregular memory accesses/references. The memory consistency engine 330 implements the necessary data structures to keep a relaxed consistency model. The compiler 340 operates to parse original computer code to identify memory accesses/references and classify them as regular or irregular memory accesses/references and inject control code for directing memory accesses/references to one of the high locality cache 310 or the transactional cache 320, as discussed in greater detail hereafter.

Depending on the locality that a memory access/reference exhibits, the software cache design implementation selects a specific mechanism to treat the memory access/reference. For

example, regular memory references may be mapped to the high locality cache 310 while irregular memory accesses/references may be mapped to the transactional cache 320. The actual mapping or selection of whether to utilize the high locality cache 310 or the transactional cache 320 may be performed, for example, by the compiler 340, such as by injecting explicit control code into the original code to direct software cache access through one of the high locality cache 310 or the transactional cache 320, in response to encountering the memory access/reference. The compiler 340 may execute on the PPU or an SPU, for example, and may provide code to the same or different SPU for execution.

As shown in FIG. 3, in response to a memory access/reference, the compiler 340 submits an address and data type associated with the memory access/reference to a selected one of the high locality cache 310 or the transactional cache 320. For every memory reference in code being handled by the compiler 340, the compiler 340 decides which path to follow to access the cache state or structures. If the memory reference corresponds to a regular memory reference, then the compiler emits control code causing the memory reference to use the path through the high locality cache 310 to access software cache structures. If the memory reference corresponds to an irregular memory reference, then control code emitted by the compiler causes the memory reference to access the software cache structures through the transactional cache.

The software cache design 300 supports a relaxed consistency. The memory consistency engine 330 is responsible for maintaining the memory consistency of the virtual memory, or in some cases real system memory, of the system, but some interaction exists between the high locality cache 310. The interaction between the high locality cache 310 and the memory consistency engine 330 comprises identifying what parts of a cache line have been modified by a memory access/reference and triggering write-back operations.

In particular, every time a cache line miss occurs in the high locality cache 310, as discussed hereafter, the high locality cache 310 informs the memory consistency engine 330 of which elements in the cache line are going to be modified. This is possible since only regular memory references are mapped to the high locality cache 310 by the compiler 340. The write-back operation is executed whenever a cache line is no longer referenced by regular memory references, e.g., when a reference counter of a cache line in the high locality cache 310 equals zero, as discussed hereafter.

Thus, for any one memory access/reference, the software cache is accessed through only one of the high locality cache or the transactional cache, again depending upon the determined locality of the memory access/reference. The high locality cache and transactional cache are consistent with each other with accessing of these caches being done by defining an address of the cache line in the cache storage. This hybrid approach to handling regular and irregular memory accesses/references is hierarchical in the sense that only the transactional cache is forced to check for the data in the high locality cache storage within a look-up process.

FIG. 4 is an exemplary block diagram of the high locality cache of the software cache design in accordance with one illustrative embodiment. The high locality cache 400 is designed to handle memory accesses/references that expose a high degree of spatial locality, i.e. regular memory accesses/references. The high locality cache 400 is designed to deliver acceptable hit/miss ratios, maximizing the overlapping between computation and communication. In one embodiment, the high locality cache 400 is fully associative and has a configurable cache line size. In another embodiment, the

high locality cache **400** has a set associativity that is sufficiently high for the type of loop processed. Namely, in such an embodiment, a cache with a set associativity of SA can process any loops in which the number of references is smaller than SA.

As shown in FIG. 4, the high locality cache **400** comprises a cache line size register **410**, a cache storage **420**, a cache line descriptor **430**, a cache look-up and translation table **440**, a cache directory **450**, an unused cache lines list **460**, and a cache hash mask register **470**. The cache line size register **410** and cache hash mask register **470** are software registers that store data representing a cache line size and a cache hash mask value, respectively. The cache storage **420** is a storage device, such as a memory, in which data that may be the subject of memory accesses/references may be located or moved into in response to a cache miss occurring. The other elements **430-450** and **460** may be data structures stored in registers, memories, or the like, associated with the high locality cache **400**. These elements **410-470** may be utilized by a compiler **480** when handling regular memory accesses/references in source code for generating executable code for execution by an associated processor implementing the software cache design. This compiler **480** may be the compiler **340** in FIG. 3, for example.

In one illustrative embodiment, the cache line size register **410** stores an unsigned integer value describing the size of the cache line (L). For example, this value may be a power of 2 and can range between 1024 and 4096.

The cache storage **420**, in one illustrative embodiment, may be dimensioned to 64 Kb and may be organized as N cache lines, where the value of N depends on the value stored in the cache line size register **410**. For example, the value for N may be computed by dividing the total cache storage, e.g., 64×1024 bytes, by the value in the cache line size register **410**. Therefore, in one illustrative embodiment, the cache storage **420** may store from 16 up to 128 cache lines.

The cache line descriptor **430** is a data structure that stores control information associated with a cache line. There are as many cache line descriptors **430** as cache lines that can be stored in the cache storage **420**. In one illustrative embodiment, the cache storage **420** may store up to 128 cache lines and thus, there are a corresponding 128 cache line descriptors **430**. In one illustrative embodiment, each of these cache line descriptors **430**, as shown in FIG. 4, includes such information as the base address, translation address, reference counter, directory link, placement link, cache line state, and communication tags.

The base address is the base address of the cache line in main memory of the system, e.g., shared memory **199** in FIG. 1. The translation address is the base address of the cache line in the cache storage **420** to where the cache line in main memory has been mapped. The cache line state describes whether or not the cache line stores modified data or not. The reference counter is an integer value that keeps track of how many regular references are referencing the cache line. The directory link is a pair of addresses pointing to the previous and next cache line descriptors in the list of the cache directory **450**. The placement link is a pair of addresses pointing to the previous and next entries in the unused cache lines list **460** in the case of the cache line being placed in the unused cache lines list **460**.

The communication tags are a pair of integer values, which in one illustrative embodiment have a range of [0 . . . 31], that are interpreted as the tags used in DMA operations for bringing the cache line into the cache storage **420** and for evicting the cache line. Taking the Cell Broadband Engine as an example of a distribute memory architecture, as noted above,

the memory flow controller is responsible for all data transfers between the SPE local store and the main memory, i.e. shared memory **199** in FIG. 1. The memory flow controller (MFC), in one implementation, supports 16 concurrent asynchronous requests of up to 16 Kb that are programmed through a DMA engine. DMA operations are identified by a DMA tag which may be an integer value from 0 to 31 that allows the application's code to track the status of the DMA operations. In order to synchronize with the DMA transfers, the MFC interface provides synchronization services that check or wait for previously issued DMA requests to complete. The MFC optimally schedules all DMA operations and is allowed to change the order the DMA operations are completed.

It is possible to explicitly indicate that a group of specific DMA operations have to be serialized. The MFC interface includes services to program "fenced" DMA transfers, i.e. DMA transfers whose order cannot be changed. Among all DMA operations holding the same DMA tag, fenced DMA operations are never reorganized, they never get executed before any prior DMA, nor are they executed after a later DMA. Thus, with regard to the software cache design of the illustrative embodiments, the communication tags in the cache line descriptors **430** store these DMA tags which may be used by the MFC interface.

The cache look-up and translation table **440** stores information for a look-up process and a translation process. Each row in the cache look-up and translation table **440** is assigned to a memory reference. In one illustrative embodiment, each row in the table contains three values used for the look-up and translation code, i.e. the base address of the cache line in the global address space, the base address of the cache line in the cache storage **420** address space, and a pointer to the cache line descriptor **430** being referenced by the regular memory reference. In another embodiment, the values kept from the look-up and translation code can be kept directly in hardware registers or in some other memory location such as local data on the stack of the current function.

The cache directory **450**, in one illustrative embodiment, is comprised of a plurality of double linked lists **452-458**, with no limitation on the number of cache lines that can be placed in any of the lists **452-458**. In one illustrative embodiment, the cache directory **450** is composed of 128 double-linked lists **452-458**. A hash function may be applied to the base address of a cache line in main memory and the output of this hash function is an identifier of one of the double-linked lists **452-458**. The double-linked lists **452-458** are used as a hash structure to speed up the look-up process, smoothing its impact on performance of the system and making possible the use of a fully associative software cache scheme. While double-linked lists are used in the illustrative embodiments, other mechanisms may be used for facilitating this hash look-up operation as will be apparent to those of ordinary skill in the art. For example, rather than using a double linked list, a single linked list may be utilized.

The cache unused lines list **460** is a list where cache line descriptors are placed when the cache line associated with the descriptor is no longer referenced by any regular memory reference. In one illustrative embodiment, the cache unused lines list **460** is a double-linked list that supports referencing to the first and last entries in the list. A reference counter for each cache line to maintain information identifying whether a cache line is being referenced by a regular memory reference. For example, a code transformation may be defined by the compiler that ensures that the reference counter is incremented/decremented whenever a memory reference no longer uses a cache line or recently starts references a cache

line, e.g. a regular memory reference is changing a cache line. If this reference counter reaches zero, then no regular memory references are referencing this cache line.

The cache hash mask register **470** stores an value used as a mask for identifying a list in the cache directory **450** in which to place a cache line descriptor **430** of a resident cache line. In one illustrative embodiment, the cache hash mask register **470** stores an unsigned integer value that operates as this mask. The operation of these elements **410-470** of the high locality cache **400** will now be described in greater detail with regard to one exemplary operational model.

An operational model of the high locality cache **400** is composed of the operations that execute upon the high locality cache structures **410-470** and implement the look-up, placement, communication, consistency, synchronization, and translation code mechanisms. The look-up operation, performed by the look-up code, is divided into two different phases. In a first phase, the look-up operation checks if an address generated by the regular memory access/reference corresponds to an entry in the cache line pointed to by the entry in the cache look-up and translation table **440** that is associated with the memory access/reference being looked-up. The entry in the cache look-up and translation table **440** is identified as corresponding to the regular memory access/reference due to the compiler having classified memory references into either regular or irregular categories and, for the regular memory references, assigning them to an integer value from 0 to a number of regular memory references minus one. This integer value is then used to associate and select a row in the cache look-up and translation table **440**.

In the case where such an entry exists in the cache look-up and translation table **440**, the look-up process continues checking if the generated address “falls out” from the cache line pointed out by the table entry. What is meant by the term “falls out” is that an address generated based on values stored in the cache line descriptor corresponding to the entry in the cache look-up and translation table **440**, i.e. the address generated based on the base address, the translation address, and a bit mask, does not match the address of the memory reference. For example, by taking the address generated by the memory reference and masking it with a bit mask from the cache line descriptor corresponding to the entry in the cache look-up and translation table **440**, the base address may be obtained which may then be compared to the base address of the cache line descriptor to determine if they match. If this is not the case, nothing else has to be done, and a cache hit is determined to have occurred. If the regular memory reference is “changing of cache line”, i.e. it has been detected that the base address corresponding to the generated address in the memory reference does not match the base address in the cache lookup and translation table, then the second phase of the look-up process is executed. That is, depending on the locality that a memory operation exposes, many instances of a memory operation can generate different memory addresses (referencing different data), but all of them belonging to the same cache line. Therefore, a memory operation is “changing of cache line” when the memory address it generates belongs to a different cache line from a previous memory address generated by the memory operation. In the case where no entry in the cache look-up and translation table **440** corresponds to the address generated by the regular memory access, the look-up process also goes on to the second phase.

The second look-up phase involves the cache directory **450** being accessed and a check as to whether the referenced cache line is already resident in the cache storage **420** or not. If a cache hit occurs, i.e. the referenced cache line is already resident in the cache storage **420**, then nothing else needs to

be done. If a cache miss occurs, i.e. the referenced cache line is not resident in the cache storage **420**, then placement and communication operations are executed.

As mentioned above, the reference counter in the cache line descriptor **430** for a particular cache line is used to keep track of how many regular references are referencing the cache line. Thus, this reference counter is incremented or decremented during the look-up process. That is, in one illustrative embodiment, the reference counter is decremented in phase one of the look-up process when a change of the cache line is detected based on the information in the cache look-up and translation table **440**. That is, if the two base addresses do not match, then the regular reference has generated an address that is no longer in the same cache line that the regular reference was previously referencing. Thus, the regular memory reference is “changing of cache line”. Thus, the reference counter is decremented when a regular memory reference no longer references the cache line. The incrementing of the reference counter is performed in the second phase of the look-up operation.

As mentioned above, if a cache miss occurs in the second phase of the look-up operation, a placement operation is performed by placement code. The placement code is based on the unused cache lines list **460** and the reference counters of the cache line descriptors **430**. The fact that a reference counter reaches a value of zero is interpreted by the compiler as the cache line being no longer in use and thus, the cache line is a candidate for future placement/replacement operations. Whenever the reference counter of a cache line descriptor **430** has a zero value, the cache line descriptor **430** is inserted into the unused cache lines list **460**. The insertion operation is done differently depending on the data in the cache line. Cache lines holding modified data, as determined from the cache line state information in the corresponding cache line descriptor **430**, are inserted as the last entry in the unused cache lines list **460**. Unmodified cache lines are inserted as the first entry in the unused cache lines list **460**. In response to a cache miss being detected, the cache line that is replaced, via the placement code, by placement of the cache line from main memory, is the cache line associated with the first cache line descriptor **430** in the unused cache lines list **460**.

As mentioned above, with regard to consistency code, the software cache operational model supports a relaxed consistency. The consistency code invokes the memory consistency engine **330** in FIG. 3 which is responsible for maintaining the memory consistency with some interaction between the high locality cache **310** and the memory consistency engine **330**. Essentially, since only regular memory references are mapped to the high locality cache **310**, every time a cache line miss occurs in the look-up code operation, the high locality cache **310** informs the memory consistency engine **330** about which elements in the cache line are going to be modified. Since the regular memory references are strided, given the regular memory address and its stride, it is possible to compute how many iterations (elements in the cache line) are going to be modified. This information is forwarded to the memory consistency engine **330** where appropriate data structures, e.g., dirty bits and the like, are updated to keep track of modified parts within every modified cache line.

Whenever the reference counter of a cache line is zero, and thus, the cache line descriptor **430** for that cache line is inserted into the unused cache lines list **460**, a write-back operation is executed, such as by way of a code transformation where the compiler injects appropriate code to perform the write-back operation. The write-back operation is only performed for those cache lines holding modified data, as

15

may be identified by the cache line state information in the cache line descriptor **430**. In one embodiment, the write-back operation involves an asynchronous DMA program, such as may execute in each of the SPEs of the Cell Broadband Engine architecture, for evicting the cache line. The communication tag, or DMA tag, used in the write-back operation is recorded in the cache line descriptor **430** of the evicted cache line. In another embodiment, a different communication mechanism can be used, such as a messages been exchanged between processors or between processor and memory.

With regard to communication code, the software cache design of the illustrative embodiments defines that all data transfer operations are programmed asynchronously. The software cache design reserves a first portion of the DMA tags, or communication tags in the cache line descriptor **430**, such as tags from 0 to 15, for data transfers that move data from main memory into the cache storage **420**. A second portion of the tags, e.g., tags 16 to 31, are reserved for data transfers moving data from the cache storage **420** to main memory. In both cases tags are assigned in a circular manner. Tags used in the communication code operations are recorded in the communication tags information of the cache line descriptor **430** for the particular cache line.

Regarding synchronization code, the software cache design of the illustrative embodiments records every communication tag used in any communication operation in the communication tag information of the cache line descriptors **430**. Thus, the synchronization operation is supported by the data in the communication tag information of the cache line descriptors **430**. It is sufficient to recover the communication tags associated with each cache line to synchronize and check that any pending data transfer of the cache line has completed. For example, in the SPE architecture of the Cell Broadband Engine, there are specific instructions in the SPE instruction architecture for synchronizing and checking that pending data transfers of a cache line have completed using DMA tags. Such mechanisms may be used to synchronize and check pending data transfers using the communication tags of the illustrative embodiments. In another embodiment, explicit synchronization primitives are used instead of the DMA tags. An array of synchronization variables is provided, and whenever a DMA command is performed using a given tag X in the prior embodiment, this alternative embodiment may first synchronize using the Xth synchronization variable in the given array of synchronization. Those skilled in the art would also understand that a more coarse, or possibly even a finer, synchronization scheme is also possible.

With regard to translation code, the cache look-up and translation table **440** contains an entry for the memory access/reference requiring the translation. This entry points to a cache line descriptor **430** where all the necessary data for the translation is available, i.e. base address and translation address. The translation code operation is composed of several arithmetic computations. The offset in the cache line being reference is computed by an AND operation, using the value in the cache line size register **410**, e.g., base address AND (cache line size register—1). The offset is added to the translation address in the cache line descriptor **430** to identify the portion of the cache line referenced by the memory access/reference.

FIGS. **5A** and **5B** are a flowchart outlining an exemplary operation of the high locality cache in accordance with one illustrative embodiment. As shown in FIGS. **5A** and **5B**, the high locality cache operation begins by checking, in the look-up and translation table, a cache line referenced by an address associated with the memory access/reference to determine if the cache line associated with the address of the memory

16

access/reference has been changed (step **510**). This step corresponds to phase one of the look-up process for the high locality cache described previously above. If the cache line referenced by the memory access/reference has not changed, a hit occurred and thus, nothing more needs to be done. If the cache line has been changed, then the reference counter needs to be updated.

If the result of the check is positive, i.e. the cache line has been changed, then the reference counter in the cache line descriptor for the referenced cache line is decremented (step **512**). A determination is then made as to whether the reference counter is zero (step **514**). If the result of this determination is negative, i.e. the reference counter for the cache line is zero, then the cache line needs to be placed in the unused cache lines list with a possible write-back operation being performed. If the result of this determination is positive, i.e. the reference counter for the cache line is NOT zero, then the operation proceeds to the second phase of the look-up process previously discussed above.

If a change is detected, a directory hash is performed (step **516**), to select one of the directory lists in the cache directory, the selected directory list is traversed, and a determination is made as to whether the cache line is resident in the selected directory list (step **518**). If the cache line is not resident in the directory list, then a software cache miss has occurred and a first entry in the unused cache line list is selected for replacement (step **520**). A determination is made as to whether the cache line corresponding to the first entry in the unused cache line list, i.e. the cache line was previously evicted (at a moment where its reference counter was zero) (step **522**). If so, the cache line that was evicted is synchronized with the main memory, e.g., using the communication tags of the cache line descriptor of the cache line that was evicted (step **524**). Thereafter, or if the cache line being evicted has not been modified, a DMA operation is programmed for moving the referenced missed cache line into the cache storage and a cache line descriptor for the cache line is generated/updated using the appropriate communication tags (step **526**). A directory hash is applied to the base address of the cache line in main memory (step **528**) and the cache line descriptor is inserted in the corresponding list in the cache directory (step **530**). The reference counter in the cache line descriptor is then incremented (step **532**) and the operation terminates.

If the cache line is resident in the directory list (step **518**), then a software cache hit has occurred and determination is made as to whether the reference counter for the cache line, in the cache line descriptor, is zero (step **534**). If the reference counter is zero, then the cache line is extracted from the unused cache lines list (step **536**). Thereafter, or if the reference counter is not zero, then the operation continues to step **532** where the reference counter is incremented and the operation then terminates.

If the reference counter is zero in step **514**, a determination is made as to whether the cache line no longer referenced by any regular memory access/reference has been modified (step **538**). If the cache line has not been modified, the cache line descriptor for the cache line no longer referenced by any memory access/reference is inserted into the unused cache lines list as the first entry in the unused cache lines list (step **540**). If the cache line no longer referenced by any memory access/reference has been modified, then a write-back operation is performed to synchronize the software cache with the main memory, a DMA operation is programmed for transferring the cache line data into main memory, and the communication tags for the DMA operation are recorded in the cache line descriptor for the cache line (step **542**). The cache line descriptor is then inserted into the unused cache lines list as a

last entry in the list (step 544). Thereafter, and after step 540, the operation continues to step 516.

Thus, the mechanisms of the illustrative embodiments provide a software cache design architecture in which regular memory accesses/references are handled in an efficient manner based on their relatively high spatial locality. That is the high locality cache is used to handle these types of regular memory accesses/references in source code. The software cache design architecture of the illustrative embodiments, as shown in FIG. 3 above, further includes mechanisms for handling irregular memory accesses/references in source code. Specifically, the transactional cache 320 in FIG. 3 is used to treat irregular memory accesses/references, i.e. memory access/references that do not expose any type of locality. Contrary to the high locality cache 310, the transactional cache 320 is not designed to deliver acceptable hit/miss ratios but instead to deliver very low overheads and maximize the overlap between computation and communication. FIG. 6 is an exemplary block diagram of a transactional cache, such as transactional cache 320 in FIG. 3, in accordance with one illustrative embodiment. The transactional cache 600 supports relaxed consistency and avoids the use of dirty bits and atomic DMA transfers, used in known distributed memory systems, by forcing the implementation to a "write through" scheme where every write operation always updates both the transactional cache and the main memory.

As shown in FIG. 6, the transactional cache 600 comprises a cache directory data structure 610, a cache storage data structure 620, a cache turn ticket data structure 630, and a cache placement index data structure 640. These data structures may be stored in registers, one or more memories, or the like, associated with a compiler 650 that controls and utilizes these data structures when handling irregular memory accesses/references in executable code executed by a processor implementing the software cache design of the illustrative embodiments. The compiler 650 may be the same compiler 650 as compiler 340 in FIG. 3, for example.

The cache directory 610, in one illustrative embodiment, is organized as a vector of 32 entries of 4-byte values. Each entry stores the base address of the cache line stored in the cache storage 620. The entry index determines the cache line in the cache storage 620 to hold the cache line data. The cache directory 610 memory addresses may be aligned to a 16-byte boundary, for example.

The cache storage 620, in one illustrative embodiment, is organized as a 4 KB capacity fully associative cache with a cache line size of 128 bytes (32 cache lines). The cache storage 620 memory addresses may be aligned to a 128-byte boundary, for example. In another embodiment, the transactional cache 620 has a set associativity that is sufficiently high for the type of loop processed. Namely, in such an embodiment, a cache with a set associativity of SA can process any loops in which the number of references is smaller than SA.

The cache turn ticket 630, in one illustrative embodiment, is a 1-byte unsigned integer used as a counter. The cache turn ticket 630 may be stored in a 16-byte vector and aligned to a 16-byte boundary, for example.

The cache placement index 640, in one illustrative embodiment, is a 1-byte unsigned integer used to determine what cache line in the cache storage 620 has to be used to serve the next software cache miss. The cache placement index 640 may be stored in a 16-byte vector and aligned to a 16-byte boundary, for example.

An operational model of a transactional cache 600, such as transaction cache 600, in accordance with one illustrative embodiment, will now be described. A "transaction" is defined herein as a set of data transfers and related computa-

tions involving the following order of three phases. In a first phase, only transfers from main memory to the local storage, e.g., a local store of a SPU, are considered. In a second phase, computations in the transaction take place using the data that has been brought into the local storage by the associated data transfer. In a third phase, results generated by the computations in the second phase are sent back to the main memory. Transactions are not forced to include all of the three phases but instead are only forced to abide by the ordering of these three phases.

At the beginning of a transaction, the cache turn ticket 630 is incremented and masked with the value 0x0000001. This ensures that the value in the cache turn ticket 630 switches between 0 and 1 as long as transactions are executed. The value in the cache turn ticket 630 determines which entries have to be used in the cache storage 620 to hold the data required by the transaction. For example, for a cache turn ticket 630 having a value of 0, entries from 0 to 15 are used. For a cache turn ticket 630 having a value of 1, entries from 16 to 31 are used.

At the beginning of a transaction, the cache placement index 640 is set to 0 or 16, according to the value in the cache turn ticket 630. As long as the memory accesses/references are treated within a transaction, the cache placement index 640 is incremented for every memory access/reference that causes a software cache miss. Moreover, at the beginning of a transaction, only the cache directory 610 entries to be used in the transaction are nil-initialized.

The operational model for the transactional cache 600 is composed of all the operations that execute upon the transactional cache structures and implement the look-up, placement, communication, consistency, synchronization, and translation mechanisms. Each of these types of operations are described hereafter.

By default, the look-up process in the transactional cache 600 is forced to check for the data in the high locality cache. If a software cache hit occurs, the data is supplied from the high locality cache structures. If a software cache miss occurs in the high locality cache, then the transactional cache directory 610 is checked. It should be noted that, depending on the available information, it may be possible to avoid this double checking, e.g., if the data mapped to the transactional cache storage 620 is known to not be mapped in the cache storage of the high locality cache, only one look-up operation is necessary. The look-up code is efficiently implemented using single-instruction-multiple-data (SIMD) instructions. The look-up code, in one illustrative embodiment, comprises 8 comparison instructions, each of these instructions comparing 4 entries in the cache directory 610 with the address associated with the memory access/reference. Thus, all 32 entries in the cache directory 610 are checked using these 8 SIMD instructions.

The placement code may comprise incrementing a value of the cache placement index 640 which determines which cache line has to be used for holding the new incoming data. After incrementing the cache placement index 640, the cache placement index 640 is masked with 0x0000001f in order to ensure this value never exceeds 32. The cache line to be used for serving the next cache miss is taken from the value of the cache placement index 640 before the increment occurs.

The communication code is directly mapped to asynchronous data transfers, e.g., by performing MFC calls that program a DMA operation. For a load operation, this communication code is placed in the first phase of a transaction which involves a data transfer of 128 bytes. This first phase of the transaction is only executed if the look-up code determined that a software cache miss occurred. For store operations, this

communication code is placed in the third phase of the transaction regardless of whether a software cache miss or a software cache hit occurred. Only the referenced datum in the store operation is transferred, not the whole cache line.

In order to ensure consistency within and across transactions, every DMA operation is tagged with an index of the cache line being used (from 0 to 31, for example), and a “fence” is placed right after the DMA operation, i.e. the MFC is not allowed to reorganize the DMA operations with the same tag. That is, those DMA operations referring to the same cache line cannot be reorganized due to the placement of the fence. Thus, the communication code for store operations are ordered as the store operations were originally coded and executed.

The write-back code is completely avoided by the illustrative embodiments. Neither dirty bits nor atomicity is required in the transactional cache implementation. In the third phase of the transaction, all modified data is sent back to main memory through the programming of a DMA operation as noted above. Thus, there is no need to maintain dirty bits or perform atomic DMA transfers of data from the software cache to the main memory.

The synchronizing code of every memory access/reference within a transaction is grouped in two different points. For load operations, the synchronize code is placed at the end of the first phase of the transaction. According to the cache turn ticket 630, the synchronizing code for load operations waits for DMA operations tagged with numbers 0 to 15 or 16 to 31, depending on which half of the 32 available cache lines are being used. For store operations, the synchronize code is placed at the beginning of the next transaction with the same value for the cache turn ticket, thereby synchronizing with the DMA operations tagged with numbers 0 to 15 or 16 to 31. Thus, the synchronization code only synchronizes with DMA transfers associated with one half of the total transactional cache storage 620 as determined by the cache turn ticket 630. This provides a double buffer effect.

The translation code for a load/store operation is composed of several arithmetic computations. First, the offset in the 128 byte cache line being referenced is computed by an AND operation using the mask 0x00000ff. The offset is then added to the base address of the cache line in the cache storage 620 pointed out by the placement code for the load/store operation. In this way, the address for the portion of the cache line referenced by the memory access/reference is generated by the translation code.

FIGS. 7A and 7B are a flowchart outlining an exemplary operation of a transactional cache in accordance with one illustrative embodiment. It should be appreciated that the operations shown in FIG. 7A are performed once per transaction. The operations shown in FIG. 7B are performed once per irregular reference within a transaction. As shown in FIG. 7A, at the beginning of a transaction, the cache turn ticket is incremented and masked (step 710). Based on the cache turn ticket value, the cache placement index is set to an initial value, e.g., 0 or 16 (step 715). The cache directory entries are nil-initialized (step 720).

For a next irregular reference within the transaction, a look-up operation is performed in the high locality cache to determine if the data corresponding to the transaction is present in the cache storage of the high locality cache (step 725). If the data is present in the cache storage of the high locality cache, then the data is supplied from this cache storage (step 730). If the data is not present in the cache storage of the high locality cache, then the transactional cache directory is checked to see if the data referenced in the transaction is present in the cache storage of the transactional cache (step

735). If the data is not present in the transactional cache's cache storage, then a placement operation is performed. If the data is present in the cache storage of the transactional cache, then the data may be returned from the cache storage (step 740). Alternatively, a placement operation may be performed as if a cache miss occurred and the data may be copied in to a newly allocated cache line.

The placement operation may comprise identifying the present value of the cache placement index as the index for the cache line to be used for serving the cache miss (step 745). The data corresponding to the transaction may then be retrieved from main or system memory and placed in the cache line of the cache storage of the transactional cache referenced by the index value of the cache placement index (step 750). The placement operation may then comprise incrementing the value of the cache placement index and then masking the value (step 755). The data may then be provided for use by the transaction from the cache storage of the transactional cache (step 760). The operation then terminates.

Having set forth the elements making up the software cache design of the illustrative embodiments, in a system in which this software cache design is utilized by the processors of the system, a compiler may perform executable code generation on source code in such a manner as to exploit the efficiencies provided by the software cache design. Specifically, knowing that the software cache mechanisms operate in the manner discussed above and have the above architecture, the compiler may perform code transformations and inject code into original computer source code so as to take advantage of these software cache mechanisms to generate more efficiently executable code.

For example, the compiler may target the execution of loops in the source code and is based on whether memory accesses/references in these loops are classified as regular or irregular memory accesses/references. The compiler may parse the original source code, or an intermediate representation of this source code, for a loop and determine if memory accesses/references in the loop have a constant stride within the loop execution or not. Regular memory accesses/references may be determined to be those memory accesses/references in the loop that have a constant stride within the loop execution. All other memory accesses/references may be considered to be irregular memory accesses/references. As mentioned above, regular memory accesses/references are mapped to the high locality cache of the software cache design while irregular memory accesses/references are mapped to the transactional cache.

The code generation by the compiler may comprise three different phases. FIG. 8 is an exemplary diagram of an executable code generation by a compiler in accordance with one illustrative embodiment. As shown in FIG. 8, in a first phase, memory accesses/references within the loop are classified as regular or irregular in the manner previously described. Every regular memory access/reference may be identified with an associated integer value ranging from zero up to the number of regular memory accesses/references in the loop minus 1 (since the integer values start at zero). This identifier is used as an index to access the cache look-up and translation table of the high locality cache during the look-up process of the high-locality cache operational model. In the depicted example, a first memory reference v2[i] is assigned the identifier 0 and a second memory reference v1[tmp] is assigned the identifier 1.

In a second phase, the compiler transforms the loop into a nest of two loops that basically perform a dynamic sub-chunking of the iteration space of the original loop, i.e. splitting the original iteration space into smaller “chunks”, the

number of which and their sizes are determined dynamically at runtime. The definition of every chunk is done according to the changes of referenced cache lines on all regular memory references.

As shown in FIG. 8, during the second phase, the compiler has introduced a “while” loop responsible for the sub-chunking. Variables `_lb_01` and `ub_01` hold the lower and upper bounds of the loop. At the beginning of the “while” loop’s body, for every regular memory reference a set of four statements has been introduced: a call to `_LOOKUP` function, an if statement that conditionally executes `_MMAP` function, a call to `_NEXT_MISS` function, and finally a call to `_UPDATE_MEM_CONSISTENCY`.

The `_LOOKUP` function checks if the cache line referenced in the memory access is causing a change of cache line, according to the content of the look-up and translation table (this matches phase one in the look-up process discussed above). In case a change has been produced, the `_MMAP` function decrements the reference counter of the recently unreferenced cache line and, if the result is equal to zero, a write-back operation is executed. After that, the directory is accessed to determined if the cache line change produced a software cache miss or a software cache hit. In case a software cache miss occurs, the placement policy is applied and the corresponding data transfer is programmed asynchronously. In both cases, the reference counter of the recently referenced cache line is incremented.

The function `_NEXT_MISS` is responsible for computing the number of iterations that can be done for a particular memory access without changing the cache line. Consecutive executions of this function always keeps the minimum value for all the checked memory accesses. The dynamic sub-chunking is controlled by the `_NEXT_MISS` function, defining the chunks of iterations according to the iterations where a regular memory reference is going to change the cache line. It should be noted that not all regular memory references are going to change cache lines at the same iteration and that a change of a cache line does not imply a cache miss.

The function `_UPDATE_MEM_CONSISTENCY` is responsible for operating, through the memory consistency engine 330 in FIG. 3, to update the data structures responsible for the memory consistency, e.g., dirty bits and the like. Since the number of iterations for the next chunk has been computed by the `_NEXT_MISS` function, it is possible to compute which elements are going to be modified in the chunk execution for the regular memory accesses. The function `_SYNCHRONIZE` blocks the execution and waits until all data transfers associated with the software cache miss have completed. The original loop appears now with an iteration space defined by the sub-chunking process (variables `_start_01` and `_end_01`). The memory accesses have been replaced by the `_LD` and `_ST` functions which are responsible for the corresponding address translation and actual memory access.

Finally, during this second phase, the compiler selects an appropriate cache line size. In a case where all regular memory references in a loop point to different cache lines, the selected cache line size needs to ensure that all of them will succeed in allocating a cache line. In accordance with one illustrative embodiment, it should be noted that for a cache line size of 512 bytes, the compiler is able to treat 128 memory references. In case the loop presents more than 128 regular memory references, then only 128 are treated as actual regular memory references with the reset being mapped in to the transactional cache.

In a third phase, the compiler applies transaction code generation at only the basic block level. It is not considered a

code reorganization that attempts to deal with control flow dependencies and thus, the compiler is only allowed to introduce transactions within a basic block. The code generation described with regard to this third phase is based on the following defined operations. The `_OPEN_TRANSACTION` operation is responsible for initializing the cache data structures for a transaction. The output of this function is a range of tags to be used by all the data transfers within the transaction about to execute. The `_GET` operation is responsible for the look-up, placement, and the communication code for bringing one cache line into the transactional cache storage. The `_PUT` operation is responsible for transferring one datum in the transactional cache storage to main memory. The `_TSYNCHRONIZE` operation is responsible for the synchronization code, according to the output of the `_OPEN_TRANSACTION` operation.

FIG. 8 shows, in the third phase, the code generation for transactions. In order to increase the number of memory operations, the loop has been unrolled. Without losing any accuracy, the loop, in this example, is unrolled four times to simplify the explanations of the code structure. A larger unroll degree could be selected if necessary. In the same direction, only the resulting loop is showed, not the epilogue where the remaining iterations should be executed in case the number of iterations is not multiple of the unroll degree. Two transactions are defined, each one composed by two references to `v1[tmp]`.

FIG. 9 is an exemplary diagram illustrating how the transactional cache structures are modified along one iteration of the unrolled loop of FIG. 8. Only the code devoted to the cache operations is shown for simplicity of the explanation. With reference to FIG. 9, initially, the cache turn ticket is supposed to be zero with each transaction being mapped to one half of the total transactional cache storage.

A first transaction is mapped to entries from 0 to 15 in the transactional cache storage. The two `_GET` operations are placed in entries 0 and 1, respectively. Tags 0 and 1 are assigned for the corresponding data transfers. The second transaction is defined and its two `_GET` operations are mapped to entries 16 and 17 respectively. Tags 16 and 17 are assigned to the corresponding data transfers. It has been assumed that both references missed in the transactional cache storage. It should be noted that the synchronization for the first transaction is going to synchronize with communications tagged with tags in the range [0 . . . 15]. The second transaction will synchronize with communications tagged in the range [16 . . . 31]. Notice that the synchronization code for first transaction is delayed by the introduction of code for the second transaction. With that, it is possible to diminish the number of lost cycles waiting for data. Similarly, the computation for the first transaction delays the execution of the synchronization code for the second transaction, causing the same effect.

The first transaction evicts the modified data using tags 0 and 1. The second transaction executes similarly, but using tags 16 and 17. It should be noted that when another iteration begins, before any transaction is defined, it has to be ensured that all evicted data is already in main memory. At the beginning of each transaction, a synchronization operation is performed with all possible pending data transfers tagged with tags corresponding to the cache lines in the half assigned to the transaction. In the example, the first transaction synchronizes with tags [0 . . . 15], but again, between the communication code that used those tags in the previous iteration and the current synchronization point, the compiler has been able to place the evicted code for transaction two. Similarly, the definition of the first transaction executes between the initial

23

synchronization of the second transaction and the evicted code that used the same tags that are going to be used in the second transaction.

The memory consistency engine contains the necessary data structures to maintain a relaxed consistency model. For every cache line in the high locality cache **310**, information about what data has been modified is maintained using a dirty bits data structure. Whenever a cache line has to be evicted, the write-back operation or process is performed which comprises three steps. First, the cache line in the main memory is read. A merge operation is then applied between the cache line in the software cache storage and the cache line in recently transferred from main memory. The output of the merge is then sent back to the main memory. All data transfers of this type are synchronous and atomic.

FIG. 10 is a flowchart outlining an exemplary operation of a compiler in accordance with one illustrative embodiment. As shown in FIG. 10, the compiler first classifies memory references of the original computer code into either a regular or irregular class of memory references (step **1010**). Regular memory references are assigned indexes into the cache lookup and translation table (step **1020**). Loops in the original computer code are transformed into a nest of two loops performing dynamic sub-chunking of the original iteration space of the loop (step **1030**). Transaction code generation is then applied at the basic block level of these loops (step **1040**) and the operation terminates.

Thus, the architecture provides a software cache design and mechanisms of a software cache for handling memory access/references based on whether or not the memory accesses/references have a high spatial locality, i.e. are regular memory accesses/references with a constant stride. With the software cache design set forth above, the amount of transferred data, and the cache line sizes, are adapted according to the type of memory access, i.e. regular or irregular. Moreover, with the use of the software cache design, a compiler is allowed to reorganize the source code to maximize the chances for the overlapping of computation with communication. Such reorganization may group all the look-up, placement, and communication code and may place it as far as possible from the actual use of the data and the synchronization before the data is used.

With the above architecture in place, further optimizations may be made, for regular memory references targeting the high locality cache, to efficiently transform an original loop in original code so as to have little, if any, cache overhead in the main computation loop of the transformed loop. These further optimizations reduce cache overhead in the main computation loop by (1) coalescing memory references together and treating them as a single unit; (2) for streams of references in which references cannot be coalesced, inserting code for performing cache lookup operations only for leading memory references in a stream of memory references instead of for every memory reference, and performing a free operation for only a trailing memory reference in a stream of memory reference; and (3) ordering cache lookup operations according to a decreasing likelihood of a cache miss such that memory references that are known will result in a cache miss are performed first. Each of these mechanisms will be described in greater detail hereafter.

FIG. 11A illustrates a very simple portion of a code loop while FIG. 11B illustrates the basic template for transformed code corresponding to the code loop in FIG. 11A in accordance with one illustrative embodiment. In the example, the original “for” code loop in FIG. 11A is transformed into a while-for loop, having an outer “while” loop **1160** and an inner “for” loop **1150**, as shown in FIG. 11B, in a similar

24

manner as shown previously in the transformation illustrated in FIG. 8. In FIG. 11B, the transformed code loop includes a first portion of initialization code **1110**, a second portion of lookup code **1120**, a third portion of reference code **1130**, and a fourth portion of cleanup code **1140**.

In the depicted example, *h* refers to the handle of the cache line in which the data is present for a particular memory reference. For example, “*h*=Lookup(ref *R*)” which checks to see if the memory reference *R* is present in the software cache and, if so, returns a corresponding handle which is a pointer to the base address of the cache line in which the data is present. If the data is not present in the cache line, then the handle is set to 0. As shown in FIG. 11B, initially, the handles for the memory references are initialized to 0.

The function “Map(ref *R*)” checks to see if the memory reference *R* is present in the cache and if not, a new cache line is installed in the cache. A reference counter for that cache line, which counts the number of consumers of the cache line, i.e. the number of memory references hitting that cache line, is incremented accordingly. The function “Free(handle *h*)” decreases the reference counter associated with the cache line pointed to by the handle *h* and frees the cache line for replacement when this reference counter reaches zero. The “Barrier” operation blocks operations to a cache line until all direct memory access (DMA) operations (issued by the Map function) targeting that cache line are completed. The “NumInLine(ref *R*, handle *h*, stride *S*)” function computes the remaining number of memory accesses in a cache line pointed to by handle *h* assuming a stride *S*, e.g., *a*[2*i*] has a stride of 2**size*(*a*[0]). By default, 0 is returned on a null handle.

The lookup code portion **1120** first determines if the number of remaining memory accesses in the cache line is zero or not. If the number of remaining memory access is zero, then a call to “Free(*h*)” releases the handle *h* so that the cache line is freed and a next cache line may be loaded in its place. This occurs when there are no more memory references that access the current cache line. The handle *h* is then set to the base address of the cache line for the memory reference *a*[*xi*+*y*]. If the number of remaining memory accesses is greater than zero, then a trip value is calculated indicating the number of iterations the inner “for” loop **1150** can iterate over without a cache miss. This value is used to determine the upper bound (UB) used in the “for” loop.

The reference code portion **1130** calculates the position in the cache for the data that is required as the combination of the base address associated with the handle *h*, the address of the memory reference, and an offset. The cleanup code portion **1140** essentially decrements the reference counter for the cache line corresponding to the handle *h* once the memory reference *a*[*xi*+*y*] has been resolved using the reference code portion **1130**.

Given the above template shown in FIG. 11B, as mentioned above, one optimization that may be made by the illustrative embodiments is to coalesce memory references such that multiple memory references in a stream may be treated as a single memory reference for purposes of cache lookup operations. That is, given two references *a*[*x*1+*y*1] and *a*[*x*2+*y*2], if it can be proven that NumInLine(ref) is guaranteed to be the same for both memory references then the two references can be coalesced into one. In order to prove that NumInLine(ref) is the same for both memory references, the addresses of the memory references can be written in terms of a multiple of a divisor (referred to as *n* here) of the cache line size *N*. That is, in order to prove that NumInLine(ref) is the same for both memory references, integer values *n*, *u*, *v*, *w*1, and *w*2 are found such that the two memory reference addresses may be

25

written as a common multiple of n (namely $u+v*i$ below) plus two possibly distinct offsets (namely $w1$ and $w2$ below) that are less than n :

$$a[x1*i+y1]=\&a+d(x1*i+y1)=n(u+v*i)+w1, w1<n$$

$$a[x2*i+y2]=\&a+d(x2*i+y2)=n(u+v*i)+w2, w2<n$$

where n is a divisor of N (namely $n*m=N$ for some integers n and m), $\&a$ is the base address for the cache line, $d(x1*i+y1)$ and $d(x2*i+y2)$ are the offset from the base address to the address for the memory reference, i is the loop index variable, and u , v , $w1$, and $w2$ are integer numbers that are defined such as to satisfy both $\&a+d(x1*i+y1)=n(u+v*i)+w1$ and $\&a+d(x2*i+y2)=n(u+v*i)+w2$ with $w1$ and $w2$ both strictly smaller than n . Given this representation of the memory references, if one reference is determined to be present in the current cache line, then the other memory reference is also determined to be present in the current cache line. This is because both memory references are a combination of the term “ $n(u+v*i)$ ” plus an offset $w1$ or $w2$ which is less than a fraction of the cache line size N and hence are within a single cache line.

Essentially, the coalescing of memory references generates pairs of memory references in which, if one memory reference is present in a cache line, the other memory reference is guaranteed to be within the cache line, as illustrated in FIG. 12. As shown in FIG. 12, a first memory reference is represented by a shaded block and a second memory reference is represented by a white block. Each of the cache lines 1210 and 1220 is comprised of pairs of memory references such that if one memory reference 1212 is present in a cache line, its coalesced pair memory reference 1214 is also guaranteed to be present in the cache line. Such coalescing can be extended to more than two memory references. For example, consider the three references $a[x1*i+y1]$, $a[x2*i+y2]$, and $a[x3*i+y3]$. Assuming the same notation as above, one wants to determine if these three addresses can be re-written as, respectively, $n(u+v*i)+w1$, $n(u+v*i)+w2$, and $n(u+v*i)+w3$, where n is a divisor of N , each of n , u , v , $w1$, $w2$, $w3$ are integer value variables, and $w1$, $w2$, and $w3$ are strictly smaller than n . Higher degrees can be constructed by induction. It should be also noted that by definition, two references that have exactly the same address can always be coalesced if the cache line has the same number of elements of the same references.

Streams of memory references may be generated by gathering coalesced memory references into groups or streams of memory references where the combined memory references are known to span less than N bytes, i.e. the size of a single cache line. For instances where there are two or more memory references coalesced together, and the two or more memory references span less than N bytes, then it is known for sure that the two or more memory references cannot span more than two cache lines. In general, they may be all located within a single cache line, or they may span two consecutive cache lines, but never more than two consecutive cache lines. For memory references that cannot be coalesced, i.e. a single memory reference, a single cache line per stream is utilized. Thus, in the example depicted in FIG. 12, the references $a[i-1]$ 1230, $a[i+1]$ 1240, and $a[i+3]$ 1250 may constitute a memory reference stream that would currently touch a combination of the first cache line 1210 and the second cache line 1220.

In the above manner, memory references in the transformed code may be coalesced and gathered into memory reference streams. Having generated memory reference streams, code insertion is performed for each memory refer-

26

ence stream so as to perform only a single `Map()` call, i.e. a single installation of the cache line into the cache, for the leading memory reference of the memory reference stream, and a single freeing of the cache line for replacement for the trailing memory reference of the memory reference stream, rather than performing such operations for each memory reference.

Memory references in a memory reference stream that spans less than a cache line size are in at most two consecutive cache lines. Thus, for memory references in between the leading and trailing edge memory references of the memory reference stream, it need only be determined which cache line handle of the two cache lines over which the memory reference stream spans should be used for resolving the memory reference. In order to implement this memory reference stream optimization, the compiler may insert code for streams of two or more memory references in accordance with the table shown in FIG. 13.

As shown in FIG. 13, for the leading edge memory reference $a[i+3]$ (column 1310) of the memory reference stream shown in FIG. 12, initialization code 1340 is inserted to install a cache line corresponding to the leading edge memory reference cache line handle $h1$. The lookup code 1350 that is inserted for the leading edge determines if the number of remaining memory reference accesses in the cache line pointed to by the handle $h1$ is zero or not and, if zero, the handle $h1$ for a next cache line is looked up. A trip value is calculated that identifies the number of iterations of the inner “for” loop that can be iterated over without exceeding the cache line size. The cleanup code 1360 decrements the reference counter associated with the handle $h1$.

For the trailing edge memory reference $a[i-1]$ (column 1320) of the memory reference stream shown in FIG. 12, initialization code 1340 is inserted to perform a lookup of the base address corresponding to the trailing edge memory reference cache handle $h3$ in the software cache and to determine if the cache line corresponding to the handle $h3$ is in the software cache. If not, the cache line corresponding to the handle $h3$ is loaded into the software cache. The lookup code 1350 determines if the number of available memory accesses in the cache line is greater than zero and if not, the cache line corresponding to the handle $h3$ is freed for replacement, the trailing edge memory reference cache line handle $h3$ is set to the leading edge memory reference cache line handle $h1$, and a trip value is calculated. The cleanup code 1360 determines if the handle $h1$ is not equal to $h3$ and if they are not equal, decrements the reference counter associated with the cache line pointed to by handle $h3$.

For each of the memory references between the leading edge memory reference and trailing edge memory reference, i.e. the middle memory references (column 1330), initialization code 1340 is inserted for performing a lookup of the middle memory reference handle $h2$. Again, because these memory references are middle memory references, they are guaranteed to be either in the first installed cache line corresponding to the handle $h1$ or the second installed cache line corresponding to $h3$. Therefore, it is not necessary to determine if the cache line is in the software cache or not since it is guaranteed that the cache line is in the software cache by virtue of the code inserted by the compiler for the leading and trailing memory references of the memory reference stream.

The lookup code 1350 for middle memory references determines if the number of available memory accesses in the cache line is greater than zero or not and, if not, the middle memory reference cache line handle $h2$ is set to the leading edge memory reference cache line handle $h1$. Note that each middle reference necessitates code as shown in the third

column of the table in FIG. 13. Specifically, each middle reference has its own private version of the h2 variable. For example, if there were three middle references, there would be three distinct h2 variables, e.g. h21, h22, and h23. A trip value is then calculated. Recall that the goal of the trip value is to indicate the maximum number of iterations that can be executed in the innermost for loop, e.g., for loop 1150 in FIG. 11B, without encountering any misses. There is no cleanup code 1360 for the middle memory references.

Consider the example in FIG. 12. Assuming a stride of only one loop (namely, the i loop index variable is incremented by 1), it can be seen that NumInLine(a[i-1]) is 3, as there are 3 data between where the reference a[i-1] currently points in the cache line 1210 until the end of cache line 1210. Similarly, NumInLine(a[i+1]) is currently 1, as a[i+1] points to the last element in cache line 1210 and NumInLine(a[i+3])=7. The minimum of NumInLine over each of the references in the loop is computed, namely the minimum of 3, 1, and 7 equals 1. This indicates that, at most, 1 iteration may be executed without having at least one of the references in the loop experience a miss. Since it is undesirable to invoke cache miss handling code in the innermost for loop, e.g., loop 1150 in FIG. 11B, this means that no more than 1 iteration in the innermost for loop 1150 may be iterated, so as to return to the lookup code 1120 in FIG. 11B which will install the proper handle and re-compute a new upper bound UB indicating the maximum number of iterations to be executed in the innermost for loop 1150. Note that the upper bound UB also ensures that the trip count of the original loop is not exceeded (e.g., N in FIGS. 11A and 11B; note also that N here is not associated with the cache line size—it is an arbitrary upper bound).

Consider now the performance of one illustrative embodiment in the steady state, namely discarding the first and last iterations in which extra initialization and cleanup code may be executed, e.g., code located in location 1110 and 1140 in FIG. 11B. It can be seen from FIG. 13 that, for a steady state, there is only one Map() operation for bringing in a cache line into the software cache, i.e. the Map() operation in the initialization code for the leading edge memory reference. Moreover, there is only 1 Free() operation that frees the cache line for replacement, i.e. in the lookup code for the trailing edge memory reference (the other “Free()” operations serve to decrement the reference counters and do not result in the cache line being freed for replacement). The majority of memory references, i.e. the middle memory references, do not require a Map() or Free() operation and thus, the overhead associated with these operations is eliminated for the majority of memory references. Thus, the overhead for processing a stream of memory references during execution of the optimized code, through the insertion of code by the compiler in the above manner, becomes O(1) instead of O(x) where x is the number of memory references in the stream. Note also that, for streams that consists of a single memory reference, that reference is treated as a leading edge only. In other words, for a single memory reference stream, the illustrative embodiments only generate code associated with a leading reference.

Thus, the mechanisms of the illustrative embodiments may reduce overhead associated with processing memory references in code loops of original code by coalescing memory references and grouping them into memory reference streams that do not span more than N bytes, i.e. more than a single cache line size in address range. In addition, further optimization may be achieved by controlling the order in which memory references are checked so as to achieve an optimum

overlap of computation operations and communication operations, i.e. operations for loading a cache line into the software cache.

FIG. 14 illustrates how the order in which a memory reference is checked may influence computation/communication overlap. FIG. 14 refers to the order in which the lookup code, e.g., the code located in location 1120 in FIG. 11B, and corresponding code associated with the lookup code in the “lookup” row of the table in FIG. 13 for reference streams, are performed. As shown in FIG. 14, in a first stream of memory reference checks 1410, there are three memory reference checks 1411-1413 that result in a cache hit in the software cache followed by a memory reference check 1414 that results in a cache miss. Following the cache miss there is another memory reference check 1415 that results in a cache hit and then the memory barrier 1416 is reached, at which point no further operations are performed until the DMA operations associated with the cache miss are completed. The rectangle 1417 represents a number of processor cycles required to complete the DMA operations for retrieving the data corresponding to the memory reference 1414 that resulted in a cache miss.

While there is some overlap of the computation operations associated with the memory reference check 1415 and the communication operations of the memory reference check 1414, there is a large amount of time following the memory barrier 1416 in which the processor is idle and no computation operations are overlapped with the communication operations. That is, no computations are being performed while the processor waits for the data required by the cache miss to be loaded from main memory into the software cache. Thus, there are wasted processor cycles in this region.

In the second stream of memory reference checks 1420, there is a first memory reference check 1421 that results in a cache hit followed by a memory reference check 1422 that results in a cache miss. The cache miss is followed by three memory reference checks 1423-1425 that each result in a cache hit and then the memory barrier 1426. As shown by the rectangle 1427, the communication operations associated with handling the cache miss 1422 completely overlap the computation operations associated with the memory reference checks 1423-1425 before encountering the memory barrier. Thus, there is optimum overlap of computation/communication in this instance and no processor cycles are wasted waiting for DMA operations to complete.

Therefore, from the examples shown in FIG. 14 it is clear that the order in which memory reference checks are performed has an impact on the performance of the processor. It is optimum to schedule the memory reference checks that are known, or are at least most likely, to result in a cache miss, to be earlier in the memory reference stream than those that are not known, or at least are less likely, to result in a cache miss. This of course requires that one know which memory references are most likely to result in a cache miss.

With the coalescing of memory references and grouping of memory references into streams having an address range that is less than the size of a cache line, it turns out that only the leading edge memory references of such a stream are likely to result in an actual software cache miss. This is because once a cache line is loaded into the software cache, as a result of a leading edge memory reference check, all the other memory references in the stream are guaranteed to be within that cache line by the nature of the coalescing and grouping described previously. As a result, it is most beneficial with the mechanisms of the illustrative embodiments to organize memory reference checks such that the leading edge memory references of the various memory reference streams are checked

first in processing order so that their communication operations may overlap computation operations for other memory references.

Thus, with the above optimizations in mind, FIG. 15 illustrates an example operation for optimizing memory reference checks for regular memory references in a code loop of original code in accordance with one illustrative embodiment. The operation outlined in FIG. 15 assumes that memory references have already been processed in the manner described previously above to categorize them as either regular or irregular, and to insert code for directing memory references to one of the high locality cache or the transactional cache based on this categorization. It should be appreciated that the operation outlined in FIG. 15 may operate in conjunction with, and at approximately a same time, as the operation for inserting code to direct memory references to one of the high locality cache or the transactional cache as described above.

As shown in FIG. 15, the operation starts by coalescing regular memory references that can be coalesced in accordance with the criteria described previously (step 1510). The coalesced memory references are gathered into streams where the memory references span less than a cache line size N bytes (step 1520). Thereafter, the original code loop is transformed into a “while-for” loop, in accordance with the template described previously, for example (step 1530). Each stream is then processed such that initialization, lookup, and cleanup code is inserted according to the table shown in FIG. 13, for example, for a leading edge memory reference, trailing edge memory reference, and middle memory references (step 1540). Thereafter, the memory reference lookup sequences are scheduled such that memory references having a highest likelihood of a cache miss are scheduled earlier than those with a lower likelihood of a cache miss, in a descending order (step 1550). The optimized code may then be output by the compiler for use in generation an executable code that may be executed by one or more processors of the data processing system (step 1560). The operation then terminates.

In one illustrative embodiment, in step 1550, the code corresponding to the row “lookup” in the table of FIG. 13 are scheduled in the manner described above. Each reference has some lookup code which is considered as a non-divisible unit of work. The multiple units of work corresponding to the multiple references may be moved in the execution schedule as necessary to achieve optimum execution. In one illustrative embodiment, these units of work are ordered so as to execute the memory references having the highest likelihood of a cache miss earlier than units of work that are less likely to have a cache miss.

While FIG. 15 shows that the optimized code is output for generation of executable code in step 1560, it should be appreciated that additional optimizations may be applied to the code as are generally known in the art. For example, constant propagation, if-conversion, strength reduction, partial redundancy elimination, and other traditional optimizations may be applied to the transformed and optimized code output by the compiler in step 1560. It should further be appreciated that while the above mechanism is described in terms of the optimizations being applied to regular memory references, not all regular memory references may need to be optimized in this manner. To the contrary, some memory references, such as local memory references, may not need to be optimized in this manner.

Moreover, it should also be appreciated that while the above description refers to specific sizes and numbers of elements, this description is only exemplary and is not intended to set forth or imply any limitations with regard to these elements. For example, various memory sizes, cache

sizes, numbers of cache lines, etc. are set forth above with regard to illustrative embodiments, however the present invention is not limited to these sizes or numbers. To the contrary, these are only exemplary and are used to facilitate an understanding of the invention rather than specifying any limitations of the present invention. Other sizes and numbers of the various elements set forth above with regard to the illustrative embodiments may be used without departing from the spirit and scope of the present invention.

As noted above, it should be appreciated that the illustrative embodiments may take the form of an entirely hardware embodiment, an entirely software embodiment or an embodiment containing both hardware and software elements. In one exemplary embodiment, the mechanisms of the illustrative embodiments are implemented in software or program code, which includes but is not limited to firmware, resident software, microcode, etc.

A data processing system suitable for storing and/or executing program code will include at least one processor coupled directly or indirectly to memory elements through a system bus. The memory elements can include local memory employed during actual execution of the program code, bulk storage, and cache memories which provide temporary storage of at least some program code in order to reduce the number of times code must be retrieved from bulk storage during execution.

Input/output or I/O devices (including but not limited to keyboards, displays, pointing devices, etc.) can be coupled to the system either directly or through intervening I/O controllers. Network adapters may also be coupled to the system to enable the data processing system to become coupled to other data processing systems or remote printers or storage devices through intervening private or public networks. Modems, cable modems and Ethernet cards are just a few of the currently available types of network adapters.

The description of the present invention has been presented for purposes of illustration and description, and is not intended to be exhaustive or limited to the invention in the form disclosed. Many modifications and variations will be apparent to those of ordinary skill in the art. The embodiment was chosen and described in order to best explain the principles of the invention, the practical application, and to enable others of ordinary skill in the art to understand the invention for various embodiments with various modifications as are suited to the particular use contemplated.

What is claimed is:

1. A method, in a data processing system, for optimizing regular memory references in original computer code, comprising:

parsing the original computer code to identify memory references in the original computer code;

classifying the memory references in the original computer code as either regular memory references or irregular memory references, wherein accesses to a software cache by regular memory references are controlled by a high locality cache mechanism;

transforming the original computer code, by a compiler, to generate transformed computer code; and

outputting the transformed computer code for generating executable code to be executed on a computing device, wherein transforming the original computer code comprises:

grouping regular memory references into one or more memory reference streams, each memory reference stream having a leading memory reference, a trailing memory reference, and one or more middle memory references; and

31

inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations associated with the leading memory reference and trailing memory reference in a different manner from initialization, lookup, and cleanup operations for the one or more middle memory references, wherein regular memory references are memory references with high spatial locality and have a constant strided access pattern, and wherein irregular memory references are memory references that do not expose a high spatial locality or constant strided access pattern, wherein inserting instructions in the original computer code includes inserting instructions for the one or more middle memory references to determine which cache handle of the two cache lines to use for each of the one or more middle memory references.

2. The method of claim 1, wherein grouping regular memory references into one or more memory reference streams comprises generating pairs of memory references in which, if one memory reference is present in a cache line, the other memory reference is guaranteed to be within the cache line.

3. The method of claim 1, wherein grouping regular memory references into one or more memory reference streams comprises gathering the regular memory references into one or more memory reference streams where coalesced regular memory references of a memory reference stream span less than the size of a single cache line.

4. The method of claim 3, wherein coalesced regular memory references of a memory reference stream at most span two cache lines in a software cache.

5. The method of claim 1, wherein inserting instructions in the original computer code comprises inserting instructions to perform a single lookup operation for a cache line in the software cache for only the leading memory reference of each memory reference stream and inserting instructions to perform a single freeing operation for freeing the cache line in the software cache for replacement for only the trailing memory reference of each memory reference stream, and wherein insertion of instructions to perform the lookup operation and freeing operation are not performed for the one or more middle memory references.

6. The method of claim 1, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each leading edge memory reference of each of the one or more memory reference streams:

initialization code to install a cache line corresponding to a cache line handle h1 associated with the leading edge memory reference;

lookup code that determines if a number of remaining memory reference accesses in the cache line associated with the cache line handle h1 is zero or not and for looking up a cache line handle of a next cache line in the software cache if the number of remaining memory reference accesses is zero;

code for calculating a trip value that identifies a number of iterations of an inner loop to iterate over without exceeding a cache line size of the installed cache line; and cleanup code that decrements a reference counter associated with the cache line handle h1.

7. The method of claim 6, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each trailing edge memory reference of each of the one or more memory reference streams:

32

initialization code to perform a lookup of a base address corresponding to a cache line handle h3, associated with the trailing edge memory reference, in the software cache and to determine if a cache line corresponding to the cache line handle h3 is in the software cache and load the cache line corresponding to the cache line handle h3 into the software cache if it is determined to not be in the software cache;

lookup code that determines if a number of available memory accesses in the cache line corresponding to the cache line handle h3 is greater than zero and, if not, freeing the cache line corresponding to the cache line handle h3 for replacement; and

cleanup code that determines if the cache line handle h1 is not equal to the cache line handle h3 and, if they are not equal, decrements a reference counter associated with the cache line handle h3.

8. The method of claim 7, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each middle memory reference of each of the one or more memory reference streams:

initialization code to perform a lookup operation of a cache line handle h2 associated with the middle memory reference;

lookup code that determines if a number of available memory accesses in a cache line associated with the cache line handle h2 is greater than zero or not and, if not, the cache line handle h2 is set to the cache line handle h1; and

code for calculating a new trip value indicating a maximum number of iterations that can be executed in the inner loop without exceeding a cache line size of the installed cache line.

9. The method of claim 1, further comprising: reorganizing code in the transformed computer code to perform memory reference checks in an order in which memory reference checks that are most likely to result in a software cache miss are executed before memory reference checks that are less likely to result in a software cache miss.

10. A computer program product comprising a computer recordable storage device having a computer readable program recorded thereon, wherein the computer readable program, when executed on a computing device, causes the computing device to:

parse the original computer code to identify memory references in the original computer code;

classify the memory references in the original computer code as either regular memory references or irregular memory references, wherein accesses to a software cache by regular memory references are controlled by a high locality cache mechanism;

transform the original computer code, by a compiler, to generate transformed computer code; and

output the transformed computer code for generating executable code to be executed on a computing device, wherein the computer readable program causes the computing device to transform the original computer code by:

grouping regular memory references into one or more memory reference streams, each memory reference stream having a leading memory reference, a trailing memory reference, and one or more middle memory references; and

inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup opera-

33

tions associated with the leading memory reference and trailing memory reference in a different manner from initialization, lookup, and cleanup operations for the one or more middle memory references, wherein regular memory references are memory references with high spatial locality and have a constant strided access pattern, and wherein irregular memory references are memory references that do not expose a high spatial locality or constant strided access pattern, wherein inserting instructions in the original computer code includes inserting instructions for the one or more middle memory references to determine which cache handle of the two cache lines to use for each of the one or more middle memory references.

11. The computer program product of claim 10, wherein grouping regular memory references into one or more memory reference streams comprises generating pairs of memory references in which, if one memory reference is present in a cache line, the other memory reference is guaranteed to be within the cache line.

12. The computer program product of claim 10, wherein grouping regular memory references into one or more memory reference streams comprises gathering the regular memory references into one or more memory reference streams where coalesced regular memory references of a memory reference stream span less than the size of a single cache line.

13. The computer program product of claim 12, wherein coalesced regular memory references of a memory reference stream at most span two cache lines in a software cache.

14. The computer program product of claim 10, wherein inserting instructions in the original computer code comprises inserting instructions to perform a single lookup operation for a cache line in the software cache for only the leading memory reference of each memory reference stream and inserting instructions to perform a single freeing operation for freeing the cache line in the software cache for replacement for only the trailing memory reference of each memory reference stream, and wherein insertion of instructions to perform the lookup operation and freeing operation are not performed for the one or more middle memory references.

15. The computer program product of claim 10, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each leading edge memory reference of each of the one or more memory reference streams:

initialization code to install a cache line corresponding to a cache line handle h1 associated with the leading edge memory reference;

lookup code that determines if a number of remaining memory reference accesses in the cache line associated with the cache line handle h1 is zero or not and for looking up a cache line handle of a next cache line in the software cache if the number of remaining memory reference accesses is zero;

code for calculating a trip value that identifies a number of iterations of an inner loop to iterate over without exceeding a cache line size of the installed cache line; and cleanup code that decrements a reference counter associated with the cache line handle h1.

16. The computer program product of claim 15, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each trailing edge memory reference of each of the one or more memory reference streams:

initialization code to perform a lookup of a base address corresponding to a cache line handle h3, associated with

34

the trailing edge memory reference, in the software cache and to determine if a cache line corresponding to the cache line handle h3 is in the software cache and load the cache line corresponding to the cache line handle h3 into the software cache if it is determined to not be in the software cache;

lookup code that determines if a number of available memory accesses in the cache line corresponding to the cache line handle h3 is greater than zero and, if not, freeing the cache line corresponding to the cache line handle h3 for replacement; and

cleanup code that determines if the cache line handle h1 is not equal to the cache line handle h3 and, if they are not equal, decrements a reference counter associated with the cache line handle h3.

17. The computer program product of claim 16, wherein inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations comprises, for each middle memory reference of each of the one or more memory reference streams:

initialization code to perform a lookup operation of a cache line handle h2 associated with the middle memory reference;

lookup code that determines if a number of available memory accesses in a cache line associated with the cache line handle h2 is greater than zero or not and, if not, the cache line handle h2 is set to the cache line handle h1; and

code for calculating a new trip value indicating a maximum number of iterations that can be executed in the inner loop without exceeding a cache line size of the installed cache line.

18. The computer program product of claim 10, wherein the computer readable program further causes the computing device to:

reorganize code in the transformed computer code to perform memory reference checks in an order in which memory reference checks that are most likely to result in a software cache miss are executed before memory reference checks that are less likely to result in a software cache miss.

19. An apparatus, comprising:

a processor; and

a memory coupled to the processor, wherein the memory comprises instructions which, when executed by the processor, cause the processor to:

parse the original computer code to identify memory references in the original computer code;

classify the memory references in the original computer code as either regular memory references or irregular memory references, wherein accesses to a software cache by regular memory references are controlled by a high locality cache mechanism;

transform the original computer code, by a compiler, to generate transformed computer code; and

output the transformed computer code for generating executable code to be executed on a computing device, wherein the instructions causes the processor to transform the original computer code by:

grouping regular memory references into one or more memory reference streams, each memory reference stream having a leading memory reference, a trailing memory reference, and one or more middle memory references; and

inserting, into the original computer code, instructions to execute initialization, lookup, and cleanup operations associated with the leading memory reference

and trailing memory reference in a different manner from initialization, lookup, and cleanup operations for the one or more middle memory references, wherein regular memory references are memory references with high spatial locality and have a constant strided access pattern, and wherein irregular memory references are memory references that do not expose a high spatial locality or constant strided access pattern, wherein inserting instructions in the original computer code includes inserting instructions for the one or more middle memory references to determine which cache handle of the two cache lines to use for each of the one or more middle memory references.

* * * * *