Instance-wise Points-to Analysis for Loop-based Dependence Testing*

Peng Wu [†] Paul Feautrier [‡] David Padua [§] Zehra Sura [§]

[†] IBM T.J. Watson Research Center
P.O.Box 218
Yorktown Heights, NY 10598

pengwu@us.ibm.com

[‡] A3 Project
INRIA Rocquencourt
78153 Le Chesnay, France

paul.feautrier@inria.fr

[§] Department of Computer Science
University of Illinois
Urbana, IL 61801

{padua,zsura}@cs.uiuc.edu

Keywords

pointer analysis, dependence analysis, heap analysis, Java, pointer arrays

General Terms

languages, algorithms

Categories and Subject Descriptors

D.3.4 [**Processors**]: compilers, optimizations

ABSTRACT

We present a points-to analysis that aims at enabling loop-based dependence analysis in the presence of Java references. The analysis is based on an abstraction called *element-wise points-to* (ewpt) *mapping*. An ewpt mapping summarizes, in a compact representation, the relation between a pointer and the heap object it points to, for every *instance* of the pointer inside a loop and for every *array element* directly accessible through this pointer. Such instance-wise and element-wise information is especially important for loop-based dependence analyses and for a language where multi-dimensional arrays are implemented as arrays of pointers. We describe an iterative algorithm to compute ewpt mappings. We also present techniques to remove objects from ewpt mappings for destructive updates.

The points-to algorithm was implemented and evaluated on a set of benchmark programs. We demonstrate that ewpt information can significantly improve the precision of dependence analysis. In many cases, the dependence analysis reports no false dependences due to array accesses.

---

1. INTRODUCTION

If Java is to be used for high performance computing, we must enable classical loop optimizations for it. Due to the rigid exception semantics of Java and its pervasive use of pointers, two problems need to be solved before loops can be optimized: finding accurate loop-based dependences in the presence of pointers and identifying large exception-free regions.

In the presence of pointers, a loop-based dependence test needs to relate pointers from *different program points* as well as pointers from *different loop iterations*. However, existing pointer information abstractions are either unable or not precise enough to satisfy these two requirements. Let us examine the two most commonly used pointer abstractions: alias relation and store-based points-to set.

- Alias relation, often computed as alias pairs or access paths, indicates whether two pointers may refer to the same memory location at a given program point. The limitation of alias pairs is that they only capture alias relations among pointers at the *same* program execution point. Therefore, they are not suitable to representing pointer information for dependence analyses which require to relate pointers across different execution points [9].

- A store-based points-to set[1] captures the set of memory locations [4] or anchors [9] that a pointer may point to at any given *static* program point. Although points-to sets can relate pointers from different program points, when dealing with pointers inside loops, they can not accurately capture the aliasing among pointers from different iterations. The following code fragment illustrates the scenario where points-to sets are not adequate for loop-based dependence analysis:

```
1    for (i = 0; i < m; i++) {
2        p = q;
3        p.x = ···;
4        q = new Object();
5    }
```

In this example, there is no output-dependence over statement 3 across different iterations of loop i. To

---

[1]Some pointer analyses use points-to sets to model the relation between pointers and access paths. We refer to such points-to relation as store-less points-to sets.

discover this, it is necessary to know that `p` at statement 3 points to a different object on each iteration of the loop. To the best of our knowledge, no existing pointer analysis is able to compute and represent points-to information for each loop iteration instance of a pointer.


Figure 1: 2-dimensional Java array

In Java, multi-dimensional arrays are implemented as trees of one-dimensional arrays (see Fig. 1). This feature introduces new types of pointers whose points-to sets must be captured precisely. Consider an access `p[i][j]` in a double-nested `i-j` loop. In order to compute the points-to set of `p[i]` for any value of `i`, one needs to capture precise points-to sets for all elements of `p`.

This paper presents an instance-wise points-to analysis to summarizes precise points-to information for all instances of a pointer inside a loop and for all elements of a pointer array. The algorithm was implemented and evaluated on a small set of benchmarks. The results demonstrate that our points-to analysis can significantly improve the precision of dependence analysis.

Contribution. In summary, the paper makes the following contributions:

- It proposes the *element-wise points-to mapping* abstraction that summarizes points-to information for pointer instances and array element pointers in a single form. It also introduces a precise heap naming scheme which include loop iteration vectors in the heap names.

- It proposes a store-based pointer analysis that is suitable for loop-based dependence tests and for analyzing pointer arrays.

- It presents a technique to remove objects from `ewpt` mappings for destructive updates. This technique can help to identify redundant null-pointer checks on heap-residing pointers.

Organization. The rest of the paper is organized as follows. Section 2 gives an overview of the analysis. Section 3 describes our assumptions on the use of some language features and introduces basic notations. Section 4 presents the transfer functions and Section 5 gives the iterative algorithm. Section 6 describes how to remove objects from `ewpt` mappings for destructive updates. Section 7 discusses the applications of the analysis and experimental results. Section 8 outlines some related work and Section 9 concludes.

2. OVERVIEW OF THE METHOD

In this method, heap objects are named by `new` statement and iteration vectors that uniquely identify the statement instance that allocates the heap object (i.e., instance-wise properties). The method is based on the `ewpt` mapping abstraction. Treating all references uniformly as array references of some dimensions, an `ewpt` mapping captures objects pointed to by individual elements of an array as a mapping (i.e., element-wise properties).

Consider the control-flow graph in Fig. 2 and program point 4 in iteration `i` (denoted as $4(i)$). The object pointed to by an arbitrary array element `a[x]` at $4(i)$ can be represented by the following `ewpt` mapping,

$$\text{PointsTo}(a[x]) = \{\langle \texttt{null}, x > i \rangle, \langle N_{\texttt{t}}[x], x \leq i \rangle\},$$

where

- $N_{\texttt{t}}[x]$ denotes the object allocated by statement `t` at iteration `x`.

- $\langle N_{\texttt{t}}[x], x \leq i \rangle$ represents a set that, for a given x, evaluates to $N_{\texttt{t}}[x]$ when $x \leq i$, and evaluates to an empty set at other points.

- $\langle \texttt{null}, x > i \rangle$ represents a set that evaluates to the `null` object when $x > i$, and evaluates to an empty set at other points.

The points-to algorithm is an iterative fixed point algorithm. For illustration purposes, we show next and in Table 1 the step-by-step output of `ewpt` mappings when applying our algorithm to the example in Fig. 2. Program points are labeled numerically along the control-flow edges in Fig. 2.

Point 1, `a = new Complex[n]`. Variable `a` points to $N_{\texttt{s}}$ which represents the object allocated by `s`. When first allocated, elements of $N_{\texttt{s}}$ are initialized to `null`, hence, `a[x]` points to `null` for $0 \leq x < n$.

We assume that array bounds conditions, such as $0 \leq x < n$, are implicit in `ewpt` mappings. For conciseness, the rest of the paper uses `null` instead of $\langle \texttt{null}, 0 \leq x < n \rangle$ in `ewpt` mappings.

Point 2, `for (...)`, **first iteration.**

Point 3, `b = new Complex()`. Variable `b` points to $N_{\texttt{t}}[i]$ that is the object allocated by statement instance $\texttt{t}(i)$.

Point 4, `a[i] = b`. This statement makes element `a[i]` point to $N_{\texttt{t}}[i]$, while other elements of `a` still point to `null`. Hence,

$$\text{PointsTo}(a[x]) = \texttt{null}; \langle N_{\texttt{t}}[i], x = i \rangle$$

The `ewpt` mapping above is then converted to an equivalent form

$$\text{PointsTo}(a[x]) = \texttt{null}; \langle N_{\texttt{t}}[x], x = i \rangle.$$

This transformation is necessary to preserve the precision of the algorithm during the next step (point 2, second iteration).[2]

---

[2]The two representation, $\langle N_{\texttt{t}}[i], x = i \rangle$ and $\langle N_{\texttt{t}}[x], x = i \rangle$ will become $\langle N_{\texttt{t}}[i^-], x = i^- \rangle$ and $\langle N_{\texttt{t}}[x], x = i^- \rangle$, during the next step. The latter representation (with minimum occurrences of `i`) preserves better precision.

Figure 2: Constructing 1-dimensional array of `Complex` Objects

Point 2, `for (...)`, second iteration. Since loop variable `i` is incremented, occurrences of `i` in previous ewpt mappings are replaced by i^-, which denotes the value of any iteration before `i`. This transformation is called *aging*.

Point 3, `b = new Complex()`. Variable `b` points to $N_t[i]$. Note that, object $N_t[i]$ allocated at the previous iteration of the loop became $N_t[i^-]$ after aging.

Point 4, `a[i] = b`. The statement makes `a[i]` point to $N_t[i]$, and leaves other elements of the array unchanged. Hence,

$$\text{PointsTo}(a[x])$$
$$= \texttt{null}; \langle N_t[x], x = i^- \rangle; \langle N_t[x], x = i \rangle$$
$$= \texttt{null}; \langle N_t[x], x \in \{i, i^-\} \rangle.$$

Point 2, 3, and 4, third iteration. The fixed point is reached.

Point 5. ewpt mappings at point 4 and 1 are merged. Occurrences of `i` are substituted by its last value $(n-1)$ since `i` is not live after the loop. This transformation is called *binding*.

This example reveals several important features of the pointer analysis.

- Heap objects are named by allocation statement instances. This naming scheme automatically encodes instance-wise information into the points-to information representation. In addition, using names such as $\langle N_t[x], x \leq i \rangle$, ewpt mappings are able to precisely represent the fact that an object is attached at a particular region of an array.

- Loop counters may occur in the ewpt mappings and are handled by aging and binding. Aging may introduce approximation but is needed for the convergence of the algorithm. It is the widening operator of the iterative algorithm.

- The analysis does not remove any heap names from ewpt mappings for heap assignments (i.e., no destructive updates). As a result, `null` is always present in PointsTo($a[x]$). A technique to enable kills for destructive updates will be presented in Section 6.

3. BASIC DEFINITIONS

This section describes our assumptions on the use of some language features and introduces basic notations.

3.1 Input Program Format

Without loss of generality, we assume that reference assignments of the source program comply to the forms given in Fig. 3. More complicated accesses can be converted into these forms by introducing temporaries. In fact, reference accesses supported by Java byte-code comply with these forms.

```
Stack Assignment
    1.  p = null
    2.  p = new Cls ()
    3.  p = new Cls[m_1] ... [m_k] ... []
    4.  p = q
    5.  p = q.a
    6.  p = q[e]

Heap Assignment
    7.  p.a = q
    8.  p[e] = q

Method Invocation
    9.  p = q.foo(···);
    10. q.bar(···);
```

Figure 3: Reference assignments

Control statements are loops and conditionals. For `while` loops, we create a loop counter before the analysis. These conventions allow the use of notation $s(\vec{i})$, for statement instances, where s is a statement label and $\vec{i}$ is the iteration vector of s.

3.2 Location, Object and Reference

A *location* is a place in memory where a primitive data type is stored. An *object* is an aggregation of locations and is named by a `new` statement instance. More precisely,

- Given an allocation statement instance $s(\vec{i})$, $N_s[\vec{i}]$ denotes the object created by $s(\vec{i})$.

  One complication comes from the allocation of multidimensional arrays. In general, a statement s in form of `p = new Cls[m_1]···[m_k]··· [ ]_n` is equivalent to,

```
p = new Cls[m1][]···[];
for (i1 = 0; i1 < m1; i1++) {
  p[i1] = new Cls[m2][]···[];
  for (i2 = 0; i2 <m2; i2++) {
    p[i1][i2] = new Cls[m3][]···[];
      ···
```

Iter.	PrgmPts	PointsTo(a)	PointsTo(b)	PointsTo($a[x]$)
	1	N_s	null	null
1	2	N_s	null	null
	3	N_s	$N_t[i]$	null
	4	N_s	$N_t[i]$	null; $\langle N_t[i], x=i\rangle \Rightarrow$ null; $\langle N_t[x], x=i\rangle$
2	2	N_s	$N_t[i^-]$	null; $\langle N_t[x], x=i^-\rangle$
	3	N_s	$N_t[i]$	null; $\langle N_t[x], x=i^-\rangle$
	4	N_s	$N_t[i]$	null; $\langle N_t[x], x\in\{i^-,i\}\rangle$
3	2	N_s	$N_t[i^-]$	null; $\langle N_t[x], x=i^-\rangle$
	3	N_s	$N_t[i]$	null; $\langle N_t[x], x=i^-\rangle$
	4	N_s	$N_t[i]$	null; $\langle N_t[x], x\in\{i^-,i\}\rangle$
	5	N_s	null; $N_t[n-1]$	null; $\langle N_t[x], 0\le x\le n-1\rangle$

Table 1: The step-by-step output of the analysis

To name these objects uniquely, for statement s, we create pseudo names s_ℓ, ℓ ranging from 1 to k, to represent the allocation site of objects `new Cls[m`$_\ell$`] [ ]`$\cdots$`[ ]` in the above loop. We "virtually" replace s by the above loop by expanding the current iteration space $\vec{\imath}$ by ℓ dimensions: $N_{s_\ell}[\vec{\imath},\vec{x}]$, where $\vec{x}$ is a vector of length ℓ, represents the object allocated by statement instance $s_\ell(\vec{\imath},\vec{x})$ in the expanded iteration space.

- $N_s[\vec{\imath}].a$ denotes the location occupied by field a of object $N_s[\vec{\imath}]$.

A *reference* is a strongly typed pointer to an object. It never points inside an object. The *dimension* of a reference is the number of "[ ]" in its type declaration. The dimension of a non-array reference is 0.

Consider a reference p of n dimensions. REACH(p) denotes all objects that can be reached from p. Objects in REACH(p) are further divided into two parts:

Skeleton part contains all objects pointed to by p with less than n subscripts, i.e., $p[x_1]\cdots[x_k]$ where $k < n$ (see Fig. 2).

Outer part contains all objects that can be *reached* from p with n subscripts, i.e., $p[x_1]\cdots[x_n]$.

Intuitively, the skeleton part of p captures the backbone of the array referenced by p, whereas the outer part captures the leaf objects. Objects in the skeleton part are necessarily array objects.

3.3 Element-wise Points-to Mapping

Given a reference p of n dimensions, the following *element-wise points-to mappings* are defined for p:

- We define mapping p_k for $0 \le k < n$. The domain of p_k is a set of k-tuples that contains all subscripts of p of length k. The value of $p_k(\vec{x})$, where $\vec{x}$ is a k-tuple, is the set of all objects that $p[\vec{x}]$ may point to. Due to Java's strong typing, these objects must have the same type as $p[\vec{x}]$, and they necessarily belong to the skeleton part of p.

- If the last dimension of p is of reference type, we define an additional mapping p_n. For any n-tuple, $\vec{x}$, the value of $p_n(\vec{x})$ is the set of all objects that may be *reached* through $p[\vec{x}]$. These objects can be of any type and dimension (since they may not be directly pointed to by $p[\vec{x}]$), and they necessarily belong to the outer part of p.

We call p_k the ewpt mapping of p at level k. Consider the sets computed in Table 1. There, PointsTo(a) is a_0, and PointsTo($a[x]$) is a_1.

4. TRANSFER FUNCTIONS

The points-to algorithm views every reference assignment as a transfer function of ewpt mappings. For ease of understanding, this section presents the transfer functions without specifying the representation of ewpt mappings and the implementation of the operations used. A concrete implementation of the transfer functions will be presented in Section 5.

4.1 Stack Assignments

Let s be the current statement and $i_1,\ldots,i_d$ be the counters of the loops surrounding s. Assume that p has ewpt mappings from level 0 to level n. The transfer function of each of the stack assignments is as follows:

- `p = null`: Since accesses attempted via a null pointer generate an exception, $p[x_1]\cdots[x_k]$ refers to nothing, i.e., it maps to the empty set. Hence,

$$p_0() = \{\text{null}\} \qquad (1)$$
$$p_k(x_1,\ldots,x_k) = \varnothing, 1 \le k \le n.$$

- `p = new Cls()`: p is necessarily 0-dimensional. It has only one ewpt mapping, p_0. Hence,

$$p_0() = \{N_s[i_1,\ldots,i_d]\}. \qquad (2)$$

- `p = new Cls[m`$_1$`]`$\cdots$`[m`$_\ell$`][ ]`$_{\ell+1}\cdots$`[ ]`$_n$: This statement first allocates a n-dimensional array, assigning it to `p`. Then, it allocates $(n-1)$-dimensional arrays, m_1 of them, and assign them to `p[x`$_1$`]`. Repeating the same process, until it allocates $(n-\ell+1)$-dimensional arrays, and assigning them to $p[x_1,\ldots,x_{\ell-1}]$. Elements of these $(n-\ell+1)$-dimensional arrays (referenced by $p[x_1,\ldots,x_{\ell-1}]$) are initialized to `null`. Elements of

`null` arrays are $\varnothing$. Hence,

$$
\begin{aligned}
p_0() &= \{N_{\mathsf{s}_0}[i_1, \ldots, i_d]\} \quad\quad\quad\quad\quad (3)\\
p_k(x_1, \ldots, x_k) &= \{N_{\mathsf{s}_k}[i_1, \ldots, i_d, x_1, \ldots, x_k]\}, \\
&\quad\quad 1 \le k \le \ell - 1 \\
p_\ell(x_1, \ldots, x_\ell) &= \{\texttt{null}\} \\
p_k(x_1, \ldots, x_k) &= \varnothing, \ell + 1 \le k \le n.
\end{aligned}
$$

- **p = q**: $p[x_1]\cdots[x_k]$ points to the same objects as $q[x_1]\cdots[x_k]$ does. Hence,

$$
p_k(x_1, \ldots, x_k) = q_k(x_1, \ldots, x_k), \ 0 \le k \le n. \quad (4)
$$

- **p = q.a**: $p[x_1]\cdots[x_k]$ points to the same objects as $q.a[x_1]\cdots[x_k]$ does. The legality of `q.a` implies that `q` is of 0 dimension. Then, $q_0()$ contains all objects that can be reached by `q`, including those pointed to by `q.a`. This transfer function is conservative.

$$
p_k(x_1, \ldots, x_k) = q_0(), \ 0 \le k \le n. \quad (5)
$$

Since objects in $p_k(x_1, \ldots, x_k)$ must have the same type as `p` and must be of $(n-k)$ dimensions, a refinement of the above transfer function is to assign to p_k those objects in q_0 that are of $(n-k)$ dimensions.

- **p = q[e]**: The type rules ensure that `q` has one more dimension than `p`. Hence,

$$
\begin{aligned}
p_k(x_1, \ldots, x_k) = \\
q_{k+1}(e, x_1, \ldots, x_k), \ 0 \le k \le n. \quad (6)
\end{aligned}
$$

4.2 Heap Assignments

Since a heap location may be accessed through different references, one heap assignment may change the ewpt mappings of several references at the same time. For example, consider the assignment in Fig. 4. Since the location of `p[e]` can be reached from both `p` and r, after `p[e] = q`, both p_1 and r_2 will change.


Figure 4: Heap assignments

Statement of the form `p.a = q`. Consider any reference r that can reach the location of `p.a`. Suppose that r is of n dimensions. Since `p` cannot possibly point to an array, if `p.a` can be reached from r, it must be in the outer part of r (i.e., r_n). It may seem at first that the reachability condition is $r_n(x_1, \ldots, x_n) \cap p_0() \ne \varnothing$. However, if the intersection is `null`, `p.a` cannot be reached by r because `null.a` is not a valid location. Hence, the correct reachability condition is

$$
r_n(x_1, \ldots, x_n) \cap p_0() \nsubseteq \{\texttt{null}\}.
$$

The assignment may also remove objects from r_n; however, there is not enough information to find them. Conservatively, we add $\text{REACH}(q)$ to all mappings that can reach the location of `p.a`.

All in all, let r'_n be r_n after the assignment. The transfer function is

$$
r'_n(x_1, \ldots, x_n) =
$$
$$
r_n(x_1, \ldots, x_n) \cup \{\langle \text{REACH}(q), \Sigma \rangle\} \quad (7)
$$

where Σ is the system of constraints obtained by imposing the condition

$$
r_n(x_1, \ldots, x_n) \cap p_0() \nsubseteq \{\texttt{null}\}.
$$

For example, if Σ is computed by imposing the condition $N_{\mathsf{s}}[x_1, x_2] \cap N_{\mathsf{s}}[n, n] \nsubseteq \{\texttt{null}\}$, after simplification, we have $\Sigma \equiv \{x_1 = n, x_2 = n\}$.

Statement of the form `p[e] = q`. Consider any reference r that may reach the location of `p[e]`. Suppose that r is of n dimensions, `p` is of m dimensions, then `q` must be of $(m-1)$ dimensions.

- If the location of `p[e]` belongs to the skeleton part of r, `p` must point to one of the subarrays of r, call it $\mathsf{r}[u_1]\cdots[u_h]$.[3] Then, the following condition must be satisfied:

$$
r_h(u_1, \ldots, u_h) \cap p_0() \nsubseteq \{\texttt{null}\}.
$$

In this case, the effect of `p[e] = q` is to replace subarrays of $\mathsf{r}[u_1]\cdots[u_h][e]$ with subarrays of `q`. That is to replace $\mathsf{r}[u_1]\cdots[u_h][e][x_{h+2}]\cdots[x_{h+1+k}]$ by $\mathsf{q}[x_{h+2}]\cdots[x_{h+1+k}]$. All in all, for all k such that $1 + h + k \le m$,

$$
\begin{aligned}
r'_{1+h+k}(x_1, \ldots, x_{h+1+k}) = r_{1+h+k}(x_1, \ldots, x_{1+h+k}) \\
\cup \{\langle q_k(x_{h+2}, \ldots, x_{h+1+k}), \Sigma \rangle\} \quad (8)
\end{aligned}
$$

where Σ is the system of contraints got by imposing the condition

$$
r_h(x_1, \ldots, x_h) \cap p_0() \nsubseteq \{\texttt{null}\} \land x_{h+1} = e.
$$

- If the location of `p[e]` belongs to the outer part of r, $\text{REACH}(q)$ is added to all $r_n(x_1, \ldots, x_n)$ that can reach the location of `p[e]`, which is the same as the case of `p.a = q`. Hence, the transfer function is (7).

4.3 Handling Loop Counters

There are two transformations to be applied on loop counters in ewpt mappings. The first one is aging, which is applied at every program point that follows a back-edge of a loop. It reflects the fact that i has been advanced. Consider a loop counter i. Aging would replace i in the mappings by $i - 1$. However, this may generate an infinite number of terms, such as i, $i - 1$, $i - 1 - 1$, etc., during the iterative process. To limit the number of forms that aging may generate, we introduce a symbol i^- that represents the value of the loop counter in any iteration before i. Then, during aging, occurrences of i in the mappings are replaced by i^-. This may introduce approximation in the algorithm.

---

[3] Since the dimension of this subarray (i.e., $n - h$) must be the same as the dimension of `p` (i.e., m), we can further determine h: $h = n - m$ provided that $n \ge m$.

A better scheme is to k-limit the number of forms that a loop variable can be aged into. Readers can refer to [20] for details.

The second transformation is binding, which is applied at any program point that follows a loop exit edge. This is to ensure that ewpt mappings at different program points contain only variables that are in scope at those points. For simplicity of presentation, loop counters are initialized to 0 and have a step of 1. Consider a loop variable i. If the last iteration of i is n, during binding, occurrences of i in the mappings are replaced by n, and occurrences of i^- are replaced by a range $[0, n-1]$. In the case of unknown loop bounds, i and i^- are replaced by $*$ that represents any positive integer value.

4.4 Inter-procedural Analysis

We would like to extend our method to programs with method invocations:

```
p = q.foo(···);
q.bar(···);
```

The second case covers both void methods and methods returning a non-pointer type. We assume that virtual methods and overloaded methods have been resolved by a previous pass of the compiler. We also assume that each method uses distinct local variables and distinct formal parameters. All of this can be obtained by qualification.

In the iterative algorithm, whenever a method call is encountered, we jump to the first statement of the method after executing a *prelude*; similarly, when encountering a return statement, we execute a *postlude* and jump to the statement following the invocation. In essence, the interprocedural analysis is both context-sensitive and flow-sensitive.

- During the prelude stage, actual parameters in the calling context are mapped to the formal parameters in the callee context. This can be done by executing "virtual" assignments:

  formal = actual

  Since the ewpt mappings of the actual parameters may contain loop counters from the calling context, which is not visible from the callee, we replace such loop counters in the ewpt mappings of formal parameters by "$*$".[4]

- The postlude of a method assigns the value returned by the method, if any, to the left-hand side reference in the invocation.

Since heap objects referenced by the caller may be passed to the callee, any structural update to these objects needs to be reflected to the caller. In other words, the transfer functions of heap assignment (in Section 4.2) need to be augmented to reflect the interprocedural context. This is done by making the ewpt mappings of the callers visible to

---

[4]A more precise handling of method invocation would be to virtually inline the callee into the calling context, i.e., virtually inlining the callee's iteration space into the caller's. However, such a scheme would fail in the presence of recursive method calls. In fact, this gets down to a more fundamental problem of devising a store-based analysis for recursive programs.

---

the callee. For example, transfer functions (7) and (8) now apply to ewpt mappings of all references from both the callee method and methods of its calling chain.

In a really modular analysis, methods should be associated not to ewpt, but to ewpt *transformers*. The design of such an analysis is left for future work.

5. THE ITERATIVE ALGORITHM

The iterative framework requires transfer functions and a meet operation. The meet operation is set-union. The transfer functions are described in Section 4. This section gives a concrete implementation of these transfer functions by defining the representation of ewpt mappings and providing the operations in the transfer functions based on this representation.

5.1 Symbolic Name

The algorithm operates on tables called ewpt tables. Each entry of an ewpt table contains a set of heap names in the following format:

$$\langle N_{\mathsf{s}}[a_1, \ldots, a_d], x_1 = a_{d+1}, \cdots, x_k = a_{d+k}, \Sigma_A, \Sigma_B \rangle$$

- s is a statement and d is the nesting level of s;
- $x_1, \ldots, x_k$ are the parameters of any ewpt mapping at level k where k is called the *rank* of the heap name;
- $A = \{a_1, \ldots, a_{d+k}\}$ is a set of $(d+k)$ bound variables;
- B is another set of bound variables that can be renamed at will. The size of B is bounded by $2d$ (see the definition of Σ below).
- Σ_A is a system of constraints over A with *at most* one constraint per $a \in A$. Let e be a subscript expression in the program and i be the counter of a loop surrounding s. Constraints in Σ_A take one of the following forms:

  - $a = e[i \leftarrow b]$ where $b \in B$.
  - $a = e[i \leftarrow u]$ where u is the last value of i.
  - $e[i \leftarrow 0] \le a \le e[i \leftarrow u-1]$ with the same conventions.

  The notation $e[x \leftarrow y]$ stands for substituting occurrences of x in e by y.

- Σ_B is a system of constraints over B. For each $b \in B$, the constraint is either of the form $b \in i$ or $b \in i^-$ where i is the counter of a loop surrounding s. Redundant variables in B are removed: 1) when both $b \in i$ and $c \in i$ belong to Σ_B, b in Σ_B is replaced by c, and b is removed from B; 2) when b does not occur in Σ_A, it is removed from B. These rules ensure that B contains no more than $2d$ variables.

The above representation is called the *standard format*. Heap names in the standard format are called *symbolic names*.

5.2 Transfer Functions

There is a one-to-one correspondence between any ewpt mapping p_k and ewpt[p,k]. That is symbolic names in ewpt[p,k] represents the value of $p_k(x_1, \ldots, x_k)$. For instance, the following table entry

$$\mathtt{ewpt[p,1]} = \{\mathtt{null}; \langle N_{\mathsf{s}}[a_1], x_1 = a_1, a_1 \le n \rangle\}$$

represents the mapping

$$p_1(x) = \{\mathtt{null}; \langle N_{\mathsf{s}}[x], x \le n \rangle\}.$$

`p = null`	$\mathbf{ewpt}[p,0] \leftarrow \{\mathtt{null}\}$ $1 \le k \le n : \mathbf{ewpt}[p,k] \leftarrow \varnothing$
`p = new Cls()`	$\mathbf{ewpt}[p,0] \leftarrow \langle N_\mathrm{s}[a_1,\ldots,a_d], a_1 = b_1, \ldots, a_d = b_d, b_1 \in i_1, \ldots, b_d \in i_d \rangle$
`p = new Cls[m_1]···` `[m_ℓ][ ]_{ℓ+1}···[ ]_n`	$\mathbf{ewpt}[p,0] \leftarrow \langle N_\mathrm{s}[a_1,\ldots,a_d], a_1 = b_1, \ldots, a_d = b_d, b_1 \in i_1, \ldots, b_d \in i_d \rangle$ $1 \le k \le \ell - 1 : \mathbf{ewpt}[p,k] \leftarrow \langle N_{\mathrm{s}_k}[a_1,\ldots,a_d,x_1,\ldots,x_k], a_1 = b_1, \ldots, a_d = b_d, b_1 \in i_1, \ldots, b_d \in i_d \rangle$ $\mathbf{ewpt}[p,\ell] \leftarrow \{\mathtt{null}\}$ $\ell + 1 \le k \le n : \mathbf{ewpt}[p,k] \leftarrow \varnothing$
`p = q`	$0 \le k \le n : \mathbf{ewpt}[p,k] \leftarrow \mathbf{ewpt}[q,k]$
`p = q.a`	$0 \le k \le n : \mathbf{ewpt}[p,k] \leftarrow \mathbf{ewpt}[q,0]$
`p = q[e]`	$0 \le k \le n : \mathbf{ewpt}[p,k] \leftarrow \mathbf{ewpt}[q,k+1](e)$
`p.a = q`	$\forall r, n = rank(r) :$ $\mathbf{ewpt}[r,n] \leftarrow \mathbf{ewpt}[r,n] \sqcup \langle \mathrm{REACH}(q), \mathbf{ewpt}[r,n] \cap \mathbf{ewpt}[p,0] \not\subseteq \{\mathtt{null}\} \rangle$
`p[e] = q`	$\forall r, n = rank(r), m = rank(p), \text{if } h = n - m \ge 0:$ $\mathbf{ewpt}[r,h+1+k] \leftarrow \mathbf{ewpt}[r,h+1+k] \sqcup \langle \mathbf{ewpt}[q,k](x_{h+2},\ldots,x_{h+1+k}),$ $\qquad \mathbf{ewpt}[r,h] \cap \mathbf{ewpt}[p,0] \not\subseteq \{\mathtt{null}\}, x_h = e \rangle$ $\mathbf{ewpt}[r,n] \leftarrow \mathbf{ewpt}[r,n] \sqcup \langle \mathrm{REACH}(q), \mathbf{ewpt}[r,n] \cap \mathbf{ewpt}[p,0] \not\subseteq \{\mathtt{null}\} \rangle$

Table 2: Transfer functions of ewpt tables

Derived directly from (1) - (8), Table 2 gives the transfer functions based on the `ewpt` table representation where s denotes the current statement and $i_1, \ldots, i_d$ denote the counters of the loops that surround s. Table 2 also uses several operations applied to sets of symbolic names. These operations are defined below. All but the first one are defined on symbolic names, but can be extended to sets in the usual way.

Let A_k be a set of symbolic names and ν_k be a symbolic name. Both are of rank k.

$\nu_k(e)$ binds a subscript expression, e, to the first parameter of ν_k, x_1. The resulting symbolic name is of rank $(k-1)$. Given $\nu_k = \langle N_\mathrm{s}[a_1, \ldots, a_d], \Sigma_k \rangle$, this operation eliminates x_1 from ν_k by adding the constraint $x_1 = e$ to Σ_k. The computation is performed in three steps:

1. If the system is unfeasible, it returns $\varnothing$.

2. Else, from the new constraint $x_1 = e$, we deduce $a_{d+1} = e[i \leftarrow b]$. This constraint is added to Σ_A iff a_{d+1} was not constrained in ν_k. If it was constrained, then there are two constraints for a_{d+1}, which is forbidden in the standard format. Therefore, we chose to ignore one of them. This is in fact a widening operation. Since the new constraint always defines a singleton set while the old one may define a range, we chose to replace the old constraint with the new one.

3. Finally, the parameters of ν_k are shifted: $[x_2 \leftarrow x_1, \cdots, x_k \leftarrow x_{k-1}]$. This substitution will generate a symbolic name of rank $(k-1)$.

To give an example, suppose

$$\nu_1 = \langle N_\mathrm{s}[a,b], x_1 = a, x_2 = b, a \le i, b \le j \},$$

then,

$$\nu_1(i+1) = \varnothing$$
$$\nu_1(i-1) = \langle N_\mathrm{s}[a,b], x_1 = b, a = i-1, b \le j \rangle.$$

This operation is also used to compute the points-to set of an array element from `ewpt` mappings, for instance, to compute read-sets/write-sets of a reference access for dependence analysis.

$\nu_k(x_{m+1}, \ldots, x_{m+1+k})$ substitutes x_i by x_{i+m} in ν_k for any $1 \le i \le k$. The resulting symbolic name is of rank $(k+h+1)$. For instance, given

$$v_1 = \langle N_\mathrm{s}[a], x_1 = a, a = i \rangle,$$

then,

$$v_1(x_2) = \langle N_\mathrm{s}[a], x_2 = a, a = i \rangle.$$

$A_k \sqcup \nu_k$ adds ν_k to set A_k. The resulting set is simplified by removing common names. Furthermore, if a symbolic name subsumes another, the latter is removed from the resulting set. For instance, given

$$\nu_1 = \langle N_\mathrm{s}[c], x_1 = c, c < i \rangle$$
$$A_1 = \{\langle N_\mathrm{s}[a], x_1 = a, a \le i \rangle\},$$

then,

$$A_1 \sqcup \nu_1 = \{\langle N_\mathrm{s}[d], x_1 = d, d \le i \rangle\}.$$

In this example, the resulting set represents the same mapping as A_1 (after intermediate variables are renamed).

$\nu_k \cap \mu_0 \not\subseteq \{\mathtt{null}\}$ solves a system of constraints over $x_1, \ldots, x_k$: $\nu_k(x_1, \ldots, x_k) \cap \mu_0 \not\subseteq \{\mathtt{null}\}$. Consider

$$\nu_k = \langle N_\mathrm{s}[a_1, \ldots, a_d], \Sigma_k \rangle$$
$$\mu_0 = \langle N_\mathrm{t}[a_1, \ldots, a_d], T_0 \rangle.$$

The operation returns false when $s \ne t$, i.e., objects in ν_k and μ_0 are created by different statements. Otherwise, it returns $\Sigma' = \Sigma_k \cup T_0$. Again, Σ' needs to be simplified: unfeasible constraints are represented by false; and if some variable in A has more than one constraints, the most precise one is kept and the others are discarded.

Consider the following example,

$$\nu_2(x_1, x_2) = \langle N_\mathrm{s}[a_1, a_2], x_1 = a_1,$$
$$x_2 = a_2, a_1 \le i, a_2 < j \rangle$$
$$\mu_0() = \langle N_\mathrm{s}[c_1, c_2], c_1 = i, c_1 = j \rangle$$
$$\mu_0'() = \langle N_\mathrm{s}[c_1, c_2], c_1 = i, c_2 = * \rangle.$$

then,

$$\nu_2 \cap \mu_0 \not\subseteq \{\texttt{null}\} \quad \Rightarrow \quad \text{false}$$
$$\nu_2 \cap \mu_0' \not\subseteq \{\texttt{null}\} \quad \Rightarrow \quad \{x_1 = a_1, x_2 < a_2,$$
$$a_1 = i, a_2 < j\}.$$

Reach(q) computes all objects that can be reached from reference variable q. It can be computed as the union of all **ewpt** entries of q, removing constraints over the parameters of those symbolic names.

Finally, the handling of loop counters is explained. Consider a loop counter i with an upper bound u. Aging is performed by replacing i by i^- in all symbolic names. Binding is performed as follows,

- For each $b \in i$ in B, b in Σ_A is replaced by u;
- Consider any b with a constraint $b \in i^-$ in B. If b occurs in any constraint $a = e$ in Σ_A, then $a = e$ is replaced by $e[b \leftarrow 0] \le a \le e[b \leftarrow u-1]$.

It can be checked that operations on symbolic names preserve the standard format. Hence, the number of symbolic names in a given program is finite. This is the key for the convergence proof of the algorithm. Due to space limitations, this proof is omitted. It can be found in full in [20].

5.3 Cost Analysis

A rough estimate of the complexity of the analysis is the product of the following factors:

- the number of program points;
- the number of **ewpt** mappings at each point;
- the number of iterations to reach a fixed point;
- the number of symbolic names in each mapping.

The second factor can be easily computed. Let d_i be the dimension of the i-th reference in the program. The number of **ewpt** mappings is

$$N_e = \sum_i (d_i + 1).$$

To estimate the iteration count, let us consider aging first. Since each loop counter i in the **ewpt** mappings is aged to i^- at the end of a loop body, it may take up to 2 iterations to reach a fixed point plus one more iteration to test it. The iteration count also depends on how fast modifications can be propagated. While forward modifications are propagated instantly, it may need several iterations to do backward propagations. Consider the following example:

```
0    for(i=0; i<n; i++) {
1        p = q;
2        q = r;
3        r = new Cls();
4    }
```

3 iterations are needed until the effect of statement 3 is propagated back to `p`. An upper bound of the number of iterations in a m-nested loop would be m times the number of pointer assignments in the loop body.

We have already proved in [20] that the number of symbolic names is finite. Although likely to be an overestimate, this gives an upper bound on the number of symbolic names

in an **ewpt** mapping. In practice, we can limit the size of symbolic names in an **ewpt** mapping to reduce the cost of the analysis.

One may wonder about the complexity of the operations on symbolic names. The complexity of these operations can be treated as a constant, although they may be quite expensive in practice. It is worth to mention, however, that symbolic names in **ewpt** mappings at level 0 are essentially points-to sets. Operations on them are inexpensive set union and intersection. They become expensive only when **ewpt** mappings at a level greater than 0 (for true array element accesses) are involved.

6. ADDITIONAL KILLING

This section presents a technique to remove objects from **ewpt** mappings in heap assignments. There are two reasons why the iterative algorithm performs no kill on heap assignments:

- To kill on heap assignment, one needs *must-alias* information, whereas **ewpt** mappings capture *may-points-to* information.
- To kill an object from a particular array element may involve restricting the constraint part of symbolic names. Then, aging that replaces a set by its super-set (i.e., replacing $i-1$ by i^-) is not safe any more.

The additional killing is performed after the iterative algorithm is terminated. As opposed to *widening*, this is the *narrowing* part of the points-to analysis. The technique is based on proving the solidity of objects. An object is **solid** if all its reference fields are not **null**. Here, we focus on how to remove **null** from **ewpt** mappings. However, the technique can be easily extended to a general killing scheme. Consider the example in Fig. 2. At program point 5, the following **ewpt** mappings are computed,

$$a_0() = N_{\tt s}$$
$$a_1(x) = \{\texttt{null}; \langle N_{\tt t}[x], 0 \le x < n\rangle\}.$$

Since a points to $N_{\tt s}$, $a[x]$ and $N_{\tt s}.[x]$ ought to point to the same objects. Hence, if object $N_{\tt s}$ is solid, $a[x]$, for any possible x, must not be **null**. This means that **null** can be removed from a_1. In general, if an **ewpt** mapping, p_k, contains only solid objects, **null** can be removed from p_{k+1}.

The solidity of an object can be proved in two steps:

Step One Pointer assignments in the program are converted to pseudo assignments. For every statement in form of $l = r$, the following pseudo assignment ($\leftarrow$) is constructed,

$$\text{Location}(l) \leftarrow \text{PointsTo}(r),$$

where $\text{Location}(l)$ and $\text{PointsTo}(r)$ are computed using the **ewpt** mappings from the previous iterative algorithm. A pseudo assignment is a **must-assignment** if its left-hand-side is a singleton non-null location.

For example, the following pseudo code is generated from the code in Fig. 2,

```
a ← Ns;
for (i = 0; i<n; i++) {
    b ← Nt[i];
    Ns.[i] ← Nt[i];
}
```

Step Two To prove that an object is solid, one considers all pseudo statements that assign to a field of the object. A must-assignment *generates* a non-null field if the right-hand-side (*rhs*) of the assignment contains no `null`. An assignment whose *rhs* contains `null` *kills* a non-null field. For array elements, the Gen and Kill sets are summarized as intervals over loops. Finally, if every field of an object does not point to `null`, the object is solid.

7. APPLICATIONS AND RESULTS

The transfer functions were implemented in Java and `javac` was augmented to drive the fixed point computation. We evaluate the analysis and demonstrate that `ewpt` mappings can be used to improve dependence test, loop parallelization, and exception optimization.

7.1 The Benchmarks

The experimental results are reported based on running the analysis over eight Java programs, as given in Table 3. All benchmark programs use either multi-dimensional arrays or one dimensional object arrays. Six of them are numerical codes that would benefit from classical loop optimizations. Program `listtbl` is an artificial example that constructs an array of linked lists in a loop, and is included in the benchmark because it shows an interesting pattern of reference assignments.

The inter-procedural scheme of the analysis was not implemented; instead, we inlined method calls and commented out system calls that had no side-effects in the programs. Table 3 gives lengths of the programs before and after inlining.[5]

7.2 Cost and Precision

The points-to analysis is applied to the seven benchmark programs. Analysis time is measured on an Ultra SPARC5 with a 270MHz processor, using `java` from SUN JDK1.2.2 with `jit` enabled. Table 4 summarizes the measurement for each program: "Prgm Pts" gives the number of program points where `ewpt` mappings were computed; "time" gives the actual analysis time; and "ewpt/javac" gives the percentage of the analysis time in a plain `javac` compilation.

Program	Prgm Pts	Analysis Time	
		time (ms)	ewpt/javac
`listtbl`	10	9	0.1%
`cmatmul`	25	162	2.1%
`shallow`	73	259	3.4%
`cholesky`	19	195	2.6%
`sor`	18	184	2.5%
`lufact`	40	168	2.9%
`moldyn`	53	77	1.0%
`euler`	299	2440	25%

Table 4: Analysis cost

Overall, the analysis time of the first seven programs is fairly small (0.1%- 3.4% of a plain `javac` compilation). The analysis time of `euler` (25% of `javac` compilation) is much

[5]`lufact` is smaller after inlining because of some dead code elimination done during inlining.

higher because `euler` involves switching the four sub-arrays of a 2-dimensional array `ug`. As a result, the `ewpt` mappings of `ug` contains 18 symbolic names, whereas, in other benchmarks, most `ewpt` mappings contain about 2 symbolic names. This suggests that a reasonable k-limiting on the size of `ewpt` mappings can help reduce analysis cost on irregular assignment patterns.

To measure the precision, we checked the output `ewpt` mappings obtained. For array references, the `ewpt` mappings were fairly precise. In particular, the mappings of `ug` in `euler` capture correct points-to information due to the switching of elements. Furthermore, using the "killing" technique, the analysis is able to remove all of redundant `null` from the mappings of array elements.

7.3 Dependence Analysis

Three versions of dependence tests were implemented using different pointer information. All of them use the Omega library [15] to determine the dependences.

- `type` collects read- and write-sets as sets of types. Two accesses are reported dependent if their types are compatible and at least one of them is a write. Two array accesses `a[x]` and `b[y]` are dependent if `a` and `b` are of compatible types, and if `x` and `y` may denote the same offset. The latter condition is determined by the Omega library.

- `flat` assumes no aliasing between different array elements (i.e., arrays are flat FORTRAN-like arrays) and no aliasing between arrays of different names. For non-array references, it uses a type-based analysis as in the previous test. Note that, `flat` is based on an unsafe assumption about array aliasing, hence it gives a lower bound of the number of array-induced dependences in the program.

- `ewpt` computes read- and write-sets as sets of heap locations from `ewpt` mappings. Details can be found in [20]. Dependence testing on heap locations is the same as that on FORTRAN arrays.

Table 5 gives the statistics collected for the three implementations. Only dependences due to conflicts of heap locations were reported, and at most one dependence was reported between any pair of statements. Note that `ewpt` reports fewer dependences than `flat`. This is because although `flat` assumes perfect information about arrays, it is quite conservative about non-array objects; whereas `ewpt` has information on both types of objects.

It is worth mentioning that using conventional points-to analyses in place of `type` would not improve the dependence test significantly, due to their lack of ability to disambiguate different elements of an array.

Table 5 also shows the number of parallel loops detected by `ewpt` and `type`. The actual number of parallel loops is given in `real`. We assume that conflicts due to stack locations can be handled by techniques such as scalar privatization. Overall, `ewpt` is able to detect all actual parallel loops, whereas parallel loops detected by `type` are mostly initialization loops.

7.4 Exception Analysis

A Java virtual machine automatically performs two runtime checks – the null-pointer and the array bounds check –

Program	Description	Lines (inlined)	Source
`listtbl`	constructing an array of linked lists	15	-
`cmatmul`	complex matrix multiplication	47	-
`cholesky`	cholesky factorization of a matrix	38	IBM
`shallow`	complex shallow-water simulation	197 (218)	IBM
`sor`	successive over-relaxation routine	40	Java Grande
`lufact`	LU factorization routine	287 (153)	Java Grande
`moldyn`	molecular dynamics simulation	234	Java Grande
`euler`	computational fluid dynamics	915 (2028)	Java Grande

Table 3: The benchmarks

Program	Dependences			Parallel Loops		
	type	flat	ewpt	type	ewpt	real
`listtbl`	5	3	2	0	1	1
`cmatmul`	8	3	3	3	7	7
`cholesky`	10	4	4	5	6	6
`shallow`	1092	152	152	6	17	17
`sor`	6	5	5	3	4	4
`lufact`	72	45	45	9	11	11
`moldyn`	2	0	0	17	19	19
`euler`	12559	2489	2489	36	55	55

Table 5: Dependences and parallel loops

Program	bound-check		null-check		safe loop	
	ewpt	type	ewpt	type	ewpt	real
`listtbl`	0	2	0	3	2	2
`cmatmul`	0	13	0	13	7	7
`cholesky`	0	19	0	19	8	8
`shallow`	0	278	0	278	20	20
`sor`	0	17	0	17	7	7
`lufact`	22	57	0	57	6	15
`moldyn`	2	2	0	8	22	24
`euler`	12	507	0	705	57	60

Table 6: Redundant checks and safe loops

for each indirect load. Restricted by the exception semantics of Java, an instruction cannot be moved across a point where exceptions may occur. Therefore, it is important for high-level optimizers to identify regions that are free of exceptions. One particular problem is to eliminate redundant null-pointer checks for pointers residing in array cells (e.g., `p[i]`).[6] Since ewpt mappings can tell whether an access is null, or is out-of-bounds, they can be exploited to improve Java exception analysis. For instance, if mapping p_1 contains no null, the compiler can prove that access to `p[i]` can not throw any null-pointer exception.

Table 6 gives the number of array bounds and null-pointer checks identified as redundant as well as the number of exception-free loops (safe loops). Some redundant array bounds checks are undetected because our implementation is not able to compare symbolic expressions. On the other hand, all null-pointer checks in the benchmarks have been identified as redundant due to the analysis' ability to remove null from ewpt mappings.

---

[6]Pointers stored on the heap are heap-residing pointers; those stored on the stack are stack-residing pointers.

8. RELATED WORK

We compare our work with research in three areas: analyses of heap-directed pointers, dependence analyses in the presence of pointers, and Java exception analysis.

Pointer Analysis. There are two approaches to compute properties for heap-directed pointers. The first one is referred to as *store-based* because heap locations are named statically, using either the pointer type [7, 18, 2], or the allocation site [19]. Our points-to analysis is store-based, but it uses a more precise naming scheme. We name objects by their allocation statement *instances*. In many cases, naming heap objects using allocation sites is a good cost/precision trade-off. However, for loop-based dependence tests it is important to distinguish different objects created by the same allocation site. Furthermore, our points-to analysis is able to compute points-to information for individual array elements that most others cannot.

Another store-based approach is Rugina et. al. [16]. The bulk of this paper is concerned with the determination of access regions in arrays. The usual region analysis is generalized to cases in which bounds must be expressed as polynomials in the parameters. Our aim here is quite different. We focus on the way pointers refer to dynamically allocated arrays. We cannot expect addresses created by `new` statements to be representable by polynomials. Conversely, when an array is allocated, our analysis can only build polyhedral regions, and Rugina et. al. work is more general than ours.

As opposed to the store-based approach, a store-less analysis directly computes alias properties without naming heap objects. These properties could be alias pair information [6, 8, 3] or shape information [12, 17, 9]. The shape of a pointer tells us whether a pointer refers to a list, a tree, etc. It is the "store-less" counterpart of the element-wise information that our analysis captures. Although the work by Deutsch [6] is store-less, our work shares one common feature with his. Both analyses try to summarize properties of unbounded objects in one symbolic form. His work uses symbolic access paths, whereas ours uses ewpt mappings (or symbolic names).

Dependence Testing with Pointers. Dependence tests that are based on store-less pointer analyses [11, 10, 13] represent read/write sets as sets of access paths. Access paths from different program points cannot be compared as they do not have any associated points-to information. Therefore, these dependence tests are applicable only to loops with no pointer assignments.

Like our points-to analysis, Ghiya and Hendren [9] also address pointer analysis in the context of dependence test-

ing. Their analysis aims to enhance store-less schemes so that they can be used for dependence tests. Our scheme aims to improve the precision of store-based analyses for iteration-based dependence tests. They are able to handle pointers to flat arrays, but not array elements of pointer types since their points-to analysis is not element-wise. In addition, they cannot handle dependences in a loop with pointer assignments as their scheme is not instance-wise.

Chambers *et al.* [2] address dependence analysis for Java in the presence of exceptions, multi-threading, and dynamic class loading. They use a type-based points-to analysis. They do not focus on getting precise information for loop iterations or array elements.

Overall, no previous work can decide that the following loop carries no dependence over statement 3.

```
1    for (i = 0; i<n; i++) {
2      p = new ...;
3      p.a = ...;
4    }
```

The fact that `p` points to a new object at each iteration cannot be abstracted by either shape or alias information. Our analysis can do this because of the instance-wise naming of objects.

Exception Analysis. Bodik, Gupta, and Sarkar [1] proposed a method to eliminate exception checks based on partial redundancy elimination. Lacking precise pointer information, their method only remove partially redundant checks for heap-residing references. Our points-to analysis, on the other hand, is able to directly remove redundant checks based on `ewpt` mappings.

Moreira, Gupta, and Midkiff [14] exploited exception-free regions for numerical Java programs. However, they did not address the techniques to identify such regions.

9. CONCLUSION AND FUTURE WORK

This paper presents an iterative algorithm that computes points-to information for instances of references and elements of reference arrays as `ewpt` mapping. Such pointer information can be used to enable a precise loop-based dependence test in the presence of Java references. We also propose a technique to perform kills on `ewpt` mappings for heap assignments. This technique can help identify redundant null-pointer checks for heap-residing pointers. We obtain promising results with a reasonable cost when using `ewpt` information on dependence analysis and exception analysis.

This work can be extended in many directions. The interprocedural algorithm is neither efficient nor precise. One possible solution is that when an `ewpt` mapping is likely to be changed, instead of carrying it around, abstract it into boolean properties such as the injectivity of the mapping (or combness as coined in [5]), which are much more lightweight and reach convergence faster. The challenge, then, is how to combine a store-based scheme with a store-less abstraction to improve the precision of the inter-procedural analysis. Consider the following example:

```
    ...                  bar(Cls[] f) {
  for (..i..)              for (..j..)
s:  bar(a[i]);         t:    f[j] = new Cls();
    ...                  }
```

A precise naming scheme needs to distinguish instances of `t` from different invocations of `bar()` (i.e., different instances of `s`). One method is to nest the iteration space of `t` inside that of `s`, i.e., $s, t(i, j)$, then name the object allocated by $t(j)$ in the calling context of $s(i)$ as $N_{s,t}[i, j]$. The challenge then is to devise such a store-based naming scheme for recursive programs.

We would also like to explore other application of the analysis such as heap optimizations. Since `ewpt` mappings capture precise links between references, objects and allocation sites, it is possible to use them for garbage collection, object layout optimization (e.g., flattening arrays), and object privatization.

10. REFERENCES

[1] Rastislav Bodik, Ragiv Gupta, and Vivek Sarkar. ABCD: Eliminating array bounds checks on demand. In *ACM Symp. on Programming Language Design and Implementation*, June 2000.

[2] C. Chambers, I. Pechtchanski, V. Sarkar, M. Serrano, and H. Srinivasan. Dependence analysis for Java. In *Workshop on Languages and Compilers for Parallel Computing*, August 1999.

[3] Ben-Chung Cheng and Wen mei Hwu. Modular interprocedural pointer analysis using access paths: design, implementation, and evaluation. In *ACM Symp. on Programming Language Design and Implementation*, pages 57–69, June 2000.

[4] Jong-Deok Choi, Michael Burke, and Paul Carini. Efficient flow-sensitive interprocedural computation of pointer-induced aliases and side effects. In *ACM Symp. on Principles of Programming Languages*, pages 232–245, 1993.

[5] Albert Cohen, Peng Wu, and David Padua. Pointer analysis for monotonic container traversals. Technical Report CSRD 1586, University of Illinois at Urbana-Champaign, January 2001.

[6] Alain Deutsch. Interprocedural may-alias analysis for pointers: Beyond k-limiting. In *ACM Symp. on Programming Language Design and Implementation*, June 1994.

[7] A. Diwan, K.S. McKinley, and E.B. Moss. Type-based alias analysis. In *ACM Symp. on Programming Language Design and Implementation*, June 1998.

[8] Rakesh Ghiya and Laurie J. Hendren. Connection analysis: A practical interprocedural heap analysis for C. In *Workshop on Languages and Compilers for Parallel Computing*. Springer-Verlag, 1995.

[9] Rakesh Ghiya and Laurie J. Hendren. Putting pointer analysis to work. In *ACM Symp. on Principles of Programming Languages*, January 1998.

[10] Laurie J. Hendren and Alexandru Nicolau. Parallelizing programs with recursive data structures. In *IEEE Trans. on Parallel and Distributed Computing*, January 1990.

[11] Joseph Hummel, Laurie J. Hendren, and Alex Nicolau. A general data dependence test for dynamic, pointer-based data structures. In *ACM Symp. on Programming Language Design and Implementation*, June 1994.

[12] Joseph Hummel, Laurie J. Hendren, and Alex Nicolau. A language for conveying the alising properties of dynamic, pointer-based data structures. In *Inthernational Parallel Processing Symposium*, pages 208–216, April 1994.

[13] J. Larus and P. Hilfinger. Detecting conflicts between structure accesses. In *ACM Symp. on Programming Language Design and Implementation*, Atlanta, GA, June 1988.

[14] J. E. Moreira, S. P. Midkiff, and M. Gupta. From flop to megaflops: Java for technical computing. *ACM Trans. on Programming Languages and Systems*, 22(2):265–295, March 2000.

[15] William Pugh. The Omega test: A fast and practical integer programming algorithm for dependence analysis. In *Supercomputing*, 1991.

[16] Radu Rugina and Martin Rinart. Symbolic bound analysis of pointers, array indices and accessed memory regions. In *PLDI'2000*. ACM, 2000.

[17] M. Sagiv, T. Reps, and R. Wilhelm. Solving shape analysis problems in languages with destructive updating. In *ACM Symp. on Principles of Programming Languages*, January 1996.

[18] B. Steensgaard. Points-to analysis in almost linear time. In *ACM Symp. on Principles of Programming Languages*, January 1996.

[19] R. Wilson and M. Lam. Efficient context-sensitive pointer analysis for C programs. In *ACM Symp. on Programming Language Design and Implementation*, June 1995.

[20] Peng Wu. Analyses of pointers, induction variables, and container objects for dependence testing. Technical Report UIUCDCS-R-2001-2209, University of Illinois at Urbana-Champaign, May 2001. Ph.D Thesis.