

ANALYZING THREADS FOR SHARED
MEMORY CONSISTENCY

BY

ZEHRA NOMAN SURA

B.E., Nagpur University, 1998

M.S., University of Illinois at Urbana-Champaign, 2001

DISSERTATION

Submitted in partial fulfillment of the requirements
for the degree of Doctor of Philosophy in Computer Science
in the Graduate College of the
University of Illinois at Urbana-Champaign, 2004

Urbana, Illinois

© Copyright by Zehra Noman Sura, 2004

Abstract

Languages allowing explicitly parallel, multithreaded programming (e.g. Java and C#) need to specify a memory consistency model to define program behavior. The memory consistency model defines constraints on the order of memory accesses in systems with shared memory. The design of a memory consistency model affects ease of parallel programming as well as system performance. Compiler analysis can be used to mitigate the performance impact of a memory consistency model that imposes strong constraints on shared memory access orders. In this work, we explore the capability of a compiler to analyze what restrictions are imposed by a memory consistency model for the program being compiled.

Our compiler analysis targets Java bytecodes. It focuses on two components: delay set analysis and synchronization analysis. Delay set analysis determines the order of shared memory accesses that must be respected within each individual thread of execution in the source program. We describe a simplified analysis algorithm that is applicable to programs with general thread structure (MIMD programs), and has polynomial time worst-case complexity. This algorithm uses synchronization analysis to improve the accuracy of the results. Synchronization analysis determines the order of shared memory accesses already enforced by synchronization in the source program. We describe a dataflow analysis algorithm for synchronization analysis that is efficient to compute, and that improves precision over previously known methods.

The analysis techniques described are used in the implementation of a virtual machine that guarantees sequentially consistent execution of Java bytecodes. This implementation is used to show the effectiveness of our analysis algorithms. On many benchmark programs, the performance of programs on our system is close to 100% of the performance of the same programs executing under a relaxed memory model. Specifically, we

observe an average slowdown of 10% on an Intel Xeon platform, with slowdowns of 7% or less for 7 out of 10 benchmarks. On an IBM Power3 platform, we observe an average slowdown of 26%, with slowdowns of 7% or less for 8 out of 10 benchmarks.

This work is dedicated to my family.

Acknowledgments

This work owes its existence in part to all members of the Pensieve Compiler group: Prof. David Padua, Prof. Sam Midkiff, Jaejin Lee, David Wong, and Xing Fang.

Prof. David Padua and Prof. Sam Midkiff have inspired, guided, encouraged, supported, scrutinized, and corrected this work. But most of all, I am grateful to them for the things I have learnt over the past few years.

The many discussions with David Wong have been very helpful, both in developing the techniques presented here, as well as implementing them. I am thankful for his critiques, and for his support throughout this work.

Special thanks are due to Xing Fang for implementing the memory barrier insertion phase and accomodating my feature requests.

Prof. David Padua has provided the opportunities for my work, and the freedom to experiment. I appreciate his patience and support, and I am glad for his influence in shaping the way I think.

I thank Prof. Marc Snir, Prof. Sarita Adve, and Prof. Laxmikant Kale for their interest, encouragement, insights, and guidance. Their involvement has improved the quality of this work.

I would also like to acknowledge Peng Wu. It was working with her that I first realized the joys of research.

Table of Contents

List of Tables	xi
List of Figures	xii
1 Introduction	1
1.1 Memory Consistency Models	2
1.1.1 Sequential Consistency Model	3
1.1.2 Relaxed Consistency Models	4
1.1.3 Choosing a Consistency Model	5
1.2 Objectives and Strategy	6
1.2.1 Compiling for a Memory Consistency Model	7
1.2.2 Our Compiler System	9
1.2.3 Assumptions	12
1.3 Presentation Outline	13
2 Delay Set Analysis	14
2.1 Problem Statement	14
2.2 Background	15
2.3 Simplified Delay Set Analysis Applied to Real Programs	19
2.4 Illustrative Examples	23
2.5 Implementation	25
2.5.1 Nodes in the Access Order Graph	25
2.5.2 Edges in the Access Order Graph	25

2.5.3	Algorithm	26
2.5.4	Complexity	28
2.5.5	Type Resolution	28
2.6	Optimizing for Object Constructors	29
2.7	Related Work	32
3	Synchronization Analysis	34
3.1	Problem Statement	34
3.2	Thread Structure Analysis	36
3.2.1	Matching Thread Start and Join Calls	37
3.2.2	Algorithm	39
3.2.3	Using Thread Structure Information in Delay Set Analysis	45
3.2.3.1	Eliminating Conflict Edges	45
3.2.3.2	Ordering Program Accesses	46
3.2.3.3	An Example	49
3.3	Event-based Synchronization	50
3.3.1	Algorithm	51
3.3.2	Using Event-based Synchronization in Delay Set Analysis	54
3.3.2.1	Limited Use in Java Programs	55
3.4	Lock-based Synchronization	56
3.4.1	Algorithm	57
3.4.2	Using Lock Information in Delay Set Analysis	60
3.5	Implementation	61
3.5.1	Exceptions	63
3.5.2	Space Complexity	65
3.6	Related Work	66
4	Compiler Framework	68
4.1	Java Terminology	68
4.2	The Jikes Research Virtual Machine and Compiler	69

4.3	Our Memory-model Aware Compiler	71
4.3.1	Thread Escape Analysis	72
4.3.1.1	Algorithm	72
4.3.1.2	Optimization for Fast Analysis	77
4.3.2	Alias Analysis	81
4.3.3	Inhibiting Code Transformations	82
4.3.4	Memory Barrier Insertion	85
4.3.5	Specialization of the Jikes RVM Code	87
4.3.6	Dynamic Class Loading	89
4.3.6.1	Incremental Analysis	89
4.3.6.2	Method Invalidation	93
4.4	Related Work	95
5	Experimental Evaluation	97
5.1	Benchmark Programs	97
5.2	Performance Impact	101
5.2.1	Execution Times	102
5.2.2	Effect of Escape Analysis	107
5.2.3	Accuracy of Delay Set Analysis	109
5.2.4	Memory Barrier Instruction Counts	111
5.2.5	Optimizations and Memory Usage	113
5.2.6	Effect of Synchronization Analysis	116
5.3	Analysis Times	117
5.3.1	Effect of Adaptive Re-compilation for Optimization	120
6	Conclusion	123
6.1	Contributions	123
6.2	Open Problems	125

Bibliography	127
Vita	139

List of Tables

5.1	Benchmark Characteristics	101
5.2	Slowdowns for the Intel Xeon Platform	104
5.3	Slowdowns for the IBM Power3 Platform	106
5.4	Delay Edge Counts for the Intel Xeon Platform	109
5.5	Delay Edge Counts for the Power3 Platform	110
5.6	Memory Fences Inserted and Executed on the Intel Xeon Platform	111
5.7	Sync Instructions Inserted and Executed on the Power3 Platform	112
5.8	Optimizations Inhibited and Memory Overhead	113
5.9	Analysis Times in Seconds for the Intel Xeon Platform	118
5.10	Analysis Times in Seconds for the Power3 Platform	120
5.11	Number of Methods Re-compiled for Adaptive Optimization	121
5.12	Analysis Times in Seconds With No Re-compilation for Optimization . .	121

List of Figures

1.1	Unordered Shared Memory Accesses	2
1.2	Busy-wait Synchronization Example	6
1.3	Memory Barrier Example	9
1.4	Components of Our Memory-Model Aware Compiler	10
2.1	Access Order Graphs for a Multithreaded Program	17
2.2	Re-orderings Allowed Under Sequential Consistency	17
2.3	Simplified Cycle Detection	22
2.4	Example to Illustrate Cumulative Ordering	23
2.5	Example to Illustrate a Non-minimal Cycle	24
2.6	Example to Illustrate Conservatism of Our Simplified Approach	24
2.7	Algorithm for Simplified Delay Set Analysis	27
2.8	Delay Edges Related to the Point an Object Escapes	29
2.9	Effect of Delay Set Analysis Optimization for Constructor Methods	31
3.1	Use of Synchronization to Break Delay Cycles	35
3.2	Orders Implied by Program Thread Structure	37
3.3	Matching <code>join()</code> for a <code>start()</code> in a Loop	38
3.4	Effect of Thread Start and Join Calls	41
3.5	<i>Gen</i> and <i>Kill</i> Functions for Thread Structure Analysis	42
3.6	Algorithm to Compute Orders for Thread Starts and Joins	43
3.7	Orders Inferred From Thread Structure Analysis	48
3.8	Example to Infer Orders From Thread Structure Synchronization	50

3.9	Algorithm to Compute Orders for Event-based Synchronization	53
3.10	Algorithm to Compute Sets of Acquired Locks	57
3.11	Possible Cycles for a Conflict Edge Between Synchronized Accesses . . .	60
3.12	Flow Graph With Normal Control Flow and Exception Control Flow . .	64
3.13	Factored Control Flow Graph in the Jikes RVM	64
4.1	Performance of the Jikes RVM and Commercial JVMs	70
4.2	Different Calling Contexts of Methods for Escape Analysis	77
4.3	Escape Analysis Algorithm	78
4.4	Effect of Statements in Escape Analysis Algorithm	79
4.5	Optimizing Escape Analysis for Fast Convergence	80
4.6	Loop Invariant Code Motion Example	83
4.7	Loop Unrolling Example	84
4.8	Common Subexpression and Load Elimination Example	84
4.9	Example for Memory Barrier Insertion	87
4.10	Determining the Set of Methods to Invalidate	94
5.1	Slowdowns for Sequential Consistency on Intel Xeon	103
5.2	Slowdowns for Sequential Consistency on IBM Power3	106
5.3	Slowdowns for SC Using the Jikes RVM Escape Analysis	108
5.4	Example for Loop Invariant Code Motion	114
5.5	Example for Load Elimination	115
5.6	Effect of Thread Structure Analysis and Locking Synchronization	116

Chapter 1

Introduction

In shared memory multiprocessing systems, multiple threads of execution can communicate with one another by reading and writing common memory locations. When explicit synchronization does not define a total order for all shared memory accesses in a program, the result of an execution can vary depending on the actual order in which shared memory accesses are performed by different threads. For example, consider the system in Figure 1.1. Processors 1 and 2 both access the same memory locations *X* and *Y*, and issue memory access requests independent of one another. Memory locations *X* and *Y* initially contain zero. Since there are no dependences between the instructions executed by each individual processor, the instructions *A*, *B*, *C*, and *D* may be completed in memory in any order. Depending on what this order is, the variables *t1* and *t2* may contain any combination of 0 and 1 values when both threads terminate.

A memory consistency model defines the order of shared memory accesses that different threads must appear to observe. In this work, we are interested in a compiler that takes the memory consistency model into account and generates code that ensures required shared memory access orders are maintained during execution.

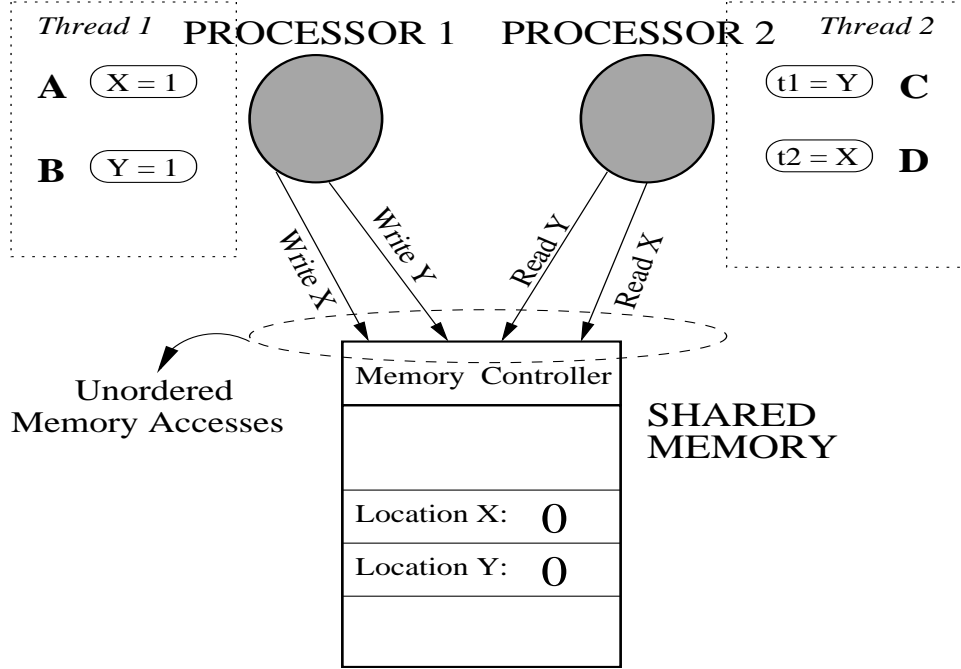


Figure 1.1: Unordered Shared Memory Accesses

1.1 Memory Consistency Models

A memory consistency model constrains the order of execution of shared memory accesses in a multithreaded program, and helps determine what constitutes a correct execution.

A memory access is an instruction that reads from or writes into some memory location(s). The value contained in a memory location may change when an instruction writes to that location. Throughout the execution of a program, each memory location will assume a sequence of values, one after another. Each value in this sequence is the result of the execution of an instruction that writes to the memory location. We refer to the sequence of values corresponding to a memory location L as $Sequence(L)$.

Definition 1.1.1 *Given a memory location L and a value V contained in $Sequence(L)$, we say V is **available to a processor** when all subsequent reads of L by the processor return a value that does not precede V in $Sequence(L)$.*

Definition 1.1.2 *Given a memory location L , an access to L is **complete** if the value that it reads from or writes into L is available to all processors in the system.*

Two memory accesses may be re-ordered if the second access is issued by a processor before the first access is *complete*. In a shared memory system, system software can cause shared memory accesses to be re-ordered when it compiles programs into machine code. Also, hardware can re-order shared memory accesses when it executes machine code.

Definition 1.1.3 *A **hardware memory model** defines shared memory access orders that are guaranteed to be preserved when the hardware executes a program.*

Definition 1.1.4 *A **programming language memory model** defines shared memory access orders that are guaranteed to be preserved on any system that executes a program written in that language.*

1.1.1 Sequential Consistency Model

Sequential consistency [Lam79] is a memory consistency model that is considered to be the simplest and most intuitive. It imposes both atomicity and ordering constraints on shared memory accesses.

Definition 1.1.5 *An access A is **atomic** if, while it is in progress, other accesses that concurrently execute with A cannot modify or observe intermediate states of the memory location accessed by A . [Sin96]*

As defined in [SS88], a *program segment* is the code executed by a single thread of a multithreaded program. Sequential consistency requires that:

“The outcome of an execution of a parallel code is as if all the instructions were executed sequentially and atomically. Instructions in the same program segment are executed in the sequential order specified by this segment; the order of execution of instructions belonging to distinct segments is arbitrary” [SS88].

Sequential consistency is a *strong* memory consistency model since it imposes strong constraints on re-ordering shared memory accesses within a program segment. If the system in Figure 1.1 is sequentially consistent, then all orders for shared memory instructions A , B , C , and D are no longer allowed. Specifically, any order in which both B occurs before C , and D occurs before A is prohibited, since it results in $t1=1$ and $t2=0$, which is not a sequentially consistent outcome.

1.1.2 Relaxed Consistency Models

Relaxed memory consistency models relax the requirement that all instructions in a program segment appear to execute in the sequential order in which they are specified. These models allow some instructions within a program segment to complete out of order.

An example of a relaxed consistency model is *weak consistency* [DSB88]. Weak consistency distinguishes each access as a regular data access or a special synchronization access. It requires that [Sin96]:

1. *All synchronization accesses must obey sequential consistency semantics.*
2. *All write accesses that occur in the program code before a synchronization access must be completed before the synchronization access is allowed.*
3. *All synchronization accesses that occur in the program code before a non-synchronization access must be completed before the non-synchronization access is allowed.*
4. *All orders due to data dependences within a thread are enforced. Thus, two accesses in a thread must appear to complete in the sequential order specified by the program source if they both access the same memory location, and at least one of them is a write.*

The IBM Power3 architecture supports weak consistency and, as a result, allows any order of shared memory instructions A , B , C , and D from the example in Figure 1.1.

1.1.3 Choosing a Consistency Model

Relaxed consistency models impose fewer constraints than sequential consistency on the order of shared memory accesses. This allows more instruction re-ordering, increasing the potential for instruction level parallelism and better performance. However, it is usually easier to understand the sequence of events in program code when using sequential consistency, since it allows fewer valid re-orderings of shared memory accesses. Thus, it seems likely that using sequential consistency will improve productivity over relaxed consistency models, both in development and maintenance of multithreaded program code.

Until the last decade, memory consistency models were mostly studied in the context of hardware architectures. They were of concern only to system programmers and designers, and architects. They had to cater to performance more than programmability. Most multiprocessor systems implement some relaxed memory consistency model [AG96] as the hardware memory model.

With the popularity of languages like Java and C# that incorporate explicit semantics of memory consistency models, programming language memory models have become an issue for a large part of the programmer community, and for language and compiler designers. The issues facing programmers with different memory consistency models are illustrated by the *busy-wait* construct which is used to synchronize accesses to `Data` in Figure 1.2. The code fragment in Figure 1.2 provides busy-wait synchronization in a sequentially consistent system. However, for most relaxed memory consistency systems, this busy-wait synchronization will not work as expected. For a language that supports a relaxed memory consistency model, the compiler can move `Flag = 1` prior to “`Data = ...`”, breaking the desired synchronization. Even if the compiler does not break the synchronization with an optimization, the target architecture may break the synchronization if it implements a relaxed consistency model that re-orders the two write instructions.

\vdots Data = ...; Flag = 1;	\vdots while (Flag==0) wait; ... = Data;
--------------------------------------	--

Figure 1.2: Busy-wait Synchronization Example

With widespread use of multithreaded programming, the trade-offs between productivity and performance have taken on increased importance. In [Hil98], Hill addresses these concerns, and advocates sacrificing some performance for ease-of-use. Sequential consistency may be the ideal choice for a programming language memory model because of its ease-of-use. However, we do not know how this choice impacts performance. This is because sequential consistency does not explicitly prohibit re-ordering of shared memory accesses in a thread; it only requires that an execution maintain the illusion that no shared memory accesses in a thread have been re-ordered. A compiler can analyze the program to determine re-orderings that do not break the illusion of sequential consistency. Then, the performance of program execution will depend on the accuracy with which these allowed re-orderings are determined, and how they compare with the number of re-orderings allowed under a relaxed consistency model.

1.2 Objectives and Strategy

Our aim in this thesis is to minimize the loss of performance when sequential consistency is the programming language memory model. We develop analysis techniques for this purpose and apply them in a just-in-time compiler for Java bytecodes. We target Java since it is a general-purpose programming language in widespread use, that also allows multithreaded programming.

For a system that implements a given programming language memory model in software, performance of the generated code tends to degrade with an increase in the number

of shared memory access orders that must be enforced. This is because there is a potential decrease in the amount of instruction-level parallelism when a greater number of accesses are ordered. Also, the memory barrier instructions that are used to enforce these orders have an intrinsic cost. The results presented in Section 5.2.1 and Section 5.2.3 support this. On an Intel Xeon-based system, a 90% reduction in the number of orders to enforce results in a performance improvement of 26 times on average. On an IBM Power3-based system, a 42% reduction in the number of orders to enforce results in a performance improvement of 84% on average. It is therefore desirable to minimize the number of orders that must be enforced. The best method known to identify orders that must be enforced is *delay set analysis* [SS88]. Delay set analysis considers thread interactions through shared memory, and determines the minimum number of orders of shared memory accesses that must be respected within each individual thread in the source program.

Precise delay set analysis was first described by Shasha and Snir in [SS88]. Their analysis assumes straight-line code where the location in memory corresponding to each access is unambiguous, and accesses performed by *each* thread in the program are individually accounted for. However, as discussed in Section 2.3, a Java compiler that performs delay set analysis must be able to handle branching control flow, imperfect memory disambiguation, and an unknown or very large number of program threads. Moreover, Krishnamurthy and Yelick [KY96] prove that Shasha and Snir’s precise delay set analysis is NP-complete, and the execution time is exponential in the number of threads.

Our strategy is to use a simplified delay set analysis that takes into account explicit synchronization programmed by the user.

1.2.1 Compiling for a Memory Consistency Model

Well-synchronized programs include explicit synchronization such that the set of valid shared memory access orders is the same for sequential consistency and relaxed con-

sistency models. In practice, most programs are written to be well-synchronized. This means that a compiler with perfect analysis capability should be able to produce code for a system that uses sequential consistency as the programming language memory model, such that the performance of this code is close to that of code generated for a system that uses a relaxed consistency model. However, analyses are not perfect. So, performance depends on the precision with which a compiler can determine the shared memory access orders that must be enforced to honor the programming language memory model.

The example in Figure 1.2 illustrates the challenges faced by a memory model aware compiler. To generate correct code, it must:

1. Inhibit itself from performing some code optimizations: These are optimizations that affect the values a processor reads from or writes into shared memory, and cause the resulting execution to be invalid according to the programming language memory model. Examples of such optimizations are dead code elimination, common subexpression elimination, register allocation, loop invariant code motion, and loop unrolling [MP90]. In Figure 1.2, allocating `Flag` to a register will cause the `while` loop to never terminate when the value of `Flag` assigned to the register is zero.
2. Enforce the programming language memory model on the target architecture: To generate correct code efficiently, the compiler must determine which accesses in each thread are to memory locations also accessed in another thread, and which of those accesses must be completed in the same order specified by the program code. These orders can be enforced by appropriately inserting synchronization instructions in the machine code generated.

Languages and hardware provide constructs that the user, or compiler, may use to prevent re-ordering of shared memory accesses – such as `synchronized` in Java [GJS96] and machine instructions like `fences` on the Intel Xeon architecture [IA3], or `syncs` on the IBM Power3 architecture [PPCb]. In this thesis, we refer to hardware instructions for memory synchronization as *memory barriers*.

<pre> : Data = ...; sync Flag = 1; </pre>	<pre> : while (Flag==0) wait; sync ...= Data; </pre>
---	--

Figure 1.3: Memory Barrier Example

Definition 1.2.1 *A **memory barrier** is a hardware instruction used to synchronize shared memory accesses. It defines two sets of memory accesses:*

1. *A set of accesses that are issued by a processor before it issues the memory barrier.*
2. *A set of accesses that will be issued by a processor after it has issued the memory barrier.*

The memory barrier instruction guarantees that all accesses in the first set complete before any access in the second set begins execution.

Figure 1.3 shows the code in Figure 1.2 with **sync** instructions inserted to enforce sequential consistency. These instructions guarantee that the set of *all* memory accesses performed in a thread before the **sync** are complete before the thread executes *any* memory access after the **sync**.

1.2.2 Our Compiler System

In Figure 1.4, we give an overview of the components in our compiler system. We implement it using the Jikes Research Virtual Machine from IBM, described in Section 4.2. The components in the figure are the analyses and transformations that we add to support sequential consistency. Given a source program, the *program analysis* component determines shared memory access orders that must be enforced so that sequential consistency is not violated. The *code transformations* component uses the analysis results

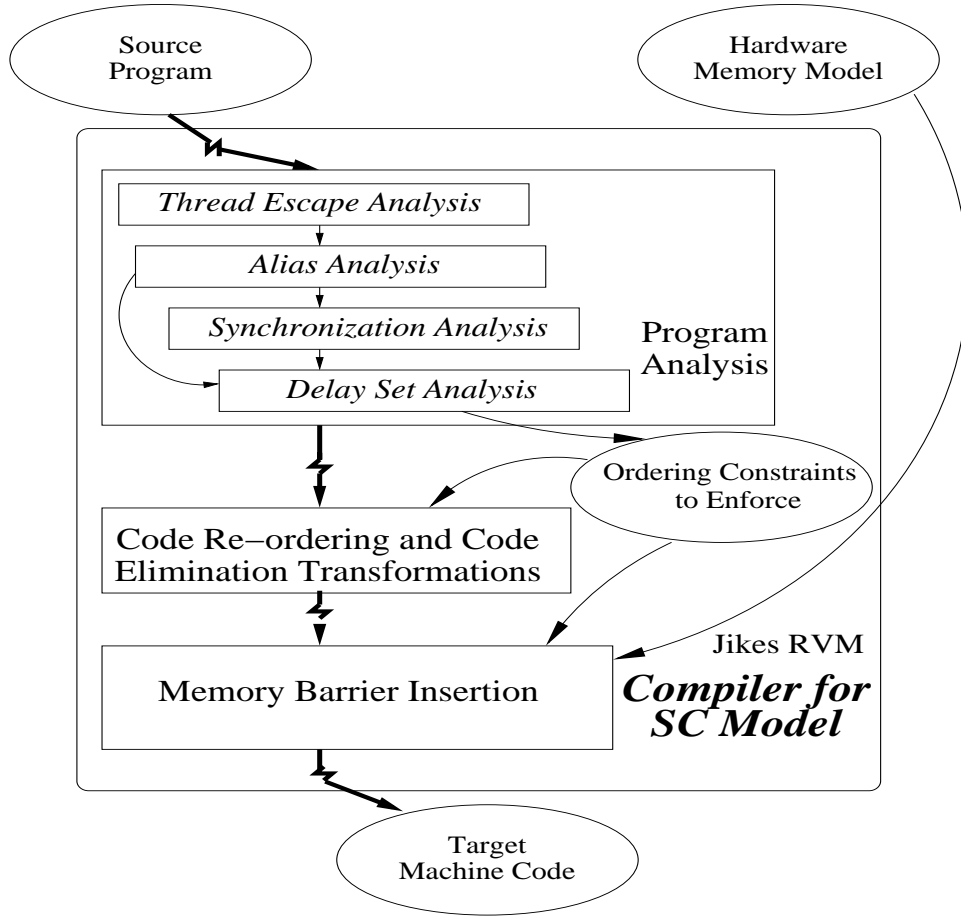


Figure 1.4: Components of Our Memory-Model Aware Compiler

to judiciously optimize the program without changing any of the required access orders. In particular, we examine each optimization and augment it so that when compiling for sequential consistency, an optimization instance is not performed if it violates an access order that needs to be enforced. The *memory barrier insertion* component also uses the analysis results, along with knowledge of the hardware memory model, to generate code that enforces the required access orders. It does this by emitting memory barriers for access orders that are not enforced by the hardware memory model but are required by sequential consistency [FLM03].

Program Analyses

Delay set analysis determines orders within a thread that must be enforced to honor the semantics of the memory consistency model. We describe our delay set analysis algorithm in Chapter 2. To perform delay set analysis, we require several other supporting analyses: synchronization analysis, thread escape analysis, and alias analysis.

Synchronization analysis determines orders that are enforced by explicit synchronization in the program. This helps improve the precision of delay set analysis. Our synchronization analysis is described in detail in Chapter 3.

Thread escape analysis identifies shared memory accesses that reference an object that may be accessed by more than one thread. Only these accesses need to be considered when performing delay set analysis. We describe our thread escape analysis in Section 4.3.1.

Alias analysis determines if two accesses may refer to the same memory location. This information is required by delay set analysis. We describe our approach to alias analysis in Section 4.3.2.

1.2.3 Assumptions

In our work, we focus on the high-level analysis that determines pairs of accesses that must not be re-ordered to enforce memory consistency. In previous work [FLM03], the back-end memory barrier insertion phase (Section 4.3.4) was implemented in the Jikes Research Virtual Machine. We re-use this implementation in our experiments to enforce the orders determined by our high-level analysis.

Besides ordering constraints, sequential consistency requires atomic execution of shared memory accesses (Section 1.1.1). We assume the following in our implementation:

1. The hardware architecture supports atomic access of multiple locations in memory that are read or written by a single instruction. Java bytecodes are specified such that each instruction reads or writes a single primitive or reference value. All such values are either 8-bit, 16-bit, 32-bit, or 64-bit wide. The IBM Power3 system and the Intel Xeon system that we use in our experiments by default provide atomic access for 8-bit, 16-bit, 32-bit and 64-bit wide locations that are correctly aligned on 8-bit, 16-bit, 32-bit and 64-bit address boundaries, respectively [PPCb, IA3]. This covers atomic access for all instructions in Java bytecodes that access memory.
2. The hardware architecture makes available memory barrier instructions that can be used to:
 - (a) order two accesses within a processor: a memory barrier between two accesses prevents the second access from being issued by the processor until it has finished executing the first access.
 - (b) provide total ordering of an access across multiple processors: a memory barrier inserted between two accesses ensures that the first access *completes* with respect to all processors in the system before the second access is issued. This is used to honor the property of sequential consistency that specifies a total ordering over all instructions in the program.

This requirement is satisfied for the systems we use in our experiments: by `sync` instructions on the IBM Power3 system [PPCa], and `mfence`, `lfence`, and `sfence` instructions on the Intel Xeon system [IA3].

1.3 Presentation Outline

This thesis is structured as follows:

- In Chapter 2, we describe delay set analysis and our simplified algorithm for it.
- In Chapter 3, we discuss how we perform synchronization analysis, and how we use the analysis results in delay set analysis.
- In Chapter 4, we present the components required in our memory-model aware compiler system, and describe their implementation.
- In Chapter 5, we report on the performance of our system when sequential consistency is used as the programming language memory model, as opposed to weak consistency.
- Finally, in Chapter 6, we present our conclusions.

Chapter 2

Delay Set Analysis

2.1 Problem Statement

If shared memory accesses are re-ordered, the resulting program computation may not be equivalent to the program computation without the re-ordering.

Definition 2.1.1 *“Two computations are **equivalent** if, on the same inputs, they produce identical values for output variables at the time output statements are executed and the output statements are executed in the same order.”*[AK02]

A memory consistency model defines constraints on the order of shared memory accesses. Delay set analysis determines those orders that must be respected within each individual thread in the source program, to ensure that the result of a program execution is always valid according to the memory consistency model. The analysis results in a *delay set*, i.e. a set of ordered pairs of shared memory accesses such that the second access in each pair must be delayed until the first access has completed.

2.2 Background

In [SS88], Shasha and Snir show how to find the minimal delay set that gives orders that must be enforced to guarantee all executions are valid for sequential consistency. Their analysis is precise since they assume:

- straight-line code with no branching control flow,
- the target location in memory for each access is unambiguously known when the analysis is performed, and
- the number of threads in the program is known when the analysis is performed.

In this section, we explain how this precise delay set analysis is performed.

Sequential consistency requires atomic execution of multiple accesses corresponding to the same program instruction. The analysis described in [SS88] explicitly determines the minimal set of orders that need to be enforced for atomicity requirements to be satisfied. In our work, we aim to provide sequential consistency for a Java virtual machine. As discussed in Section 1.2.3, each instruction (or bytecode) executed by a Java virtual machine accesses at most one location in memory¹. We assume the hardware architecture supports atomic access of locations accessed by any single Java instruction. Thus, we do not need the portion of the analysis in [SS88] that determines orders that guarantee atomicity of multiple accesses corresponding to a single instruction.

Shasha and Snir’s precise analysis uses a graph that we call the *access order graph*.

Definition 2.2.1 *An access order graph is a graph $G = (V, P, C)$, where:*

¹The *synchronized* keyword in Java provides programmers the facility to specify that some set of instructions must execute atomically with respect to some other set of instructions. By default, our compiler explicitly enforces all orders that are specified in the program using this *synchronized* keyword. As described in Section 3.4, our analysis exploits these orders to reduce the number of orders in the delay set that it computes.

- V is the set of nodes in the graph. There is one node for each shared memory access performed by each thread in the program.
- P is the set of all program edges in the graph.
- C is the set of all conflict edges in the graph.

Definition 2.2.2 A **program edge** is a directed edge from node A to node B , where A and B represent shared memory accesses performed by the same thread, and program semantics require A to appear to complete before B .

Definition 2.2.3 A **conflict edge** is an edge between two nodes that access the same memory location, and at least one of them writes to the location being accessed.

Program edges capture the ordering requirements specified by the programming language memory model, as well as those specified by the control flow semantics of a program². When sequential consistency is the memory model, program edges exist between an access and all other accesses that follow it in sequential program order.

Conflict edges represent points where instructions in the program may communicate through shared memory, and the execution of one instruction may affect another. Conflict edges within the same thread of execution are directed data dependences, whose order is determined by traditional dependence analysis. These orders are enforced by default during execution on the platforms that we use for our experiments. In general, conflict edges between accesses in different threads are not directed. Threads execute concurrently, and either one of two accesses related by a conflict edge may be executed first. The outcome of an execution may differ depending on which of these two accesses executes first.

Figure 2.1 shows a multithreaded program represented as an access order graph. Solid edges are program edges and dashed edges are conflict edges. Part (a) of the figure shows

²The precise exception model in Java also contributes to its control flow semantics.

the access order graph for a program assuming sequential consistency, so that all accesses within a thread are ordered. Part (b) of the figure shows the access order graph for the same program assuming weak consistency, so that only accesses to the same memory location are related by edges in the graph.

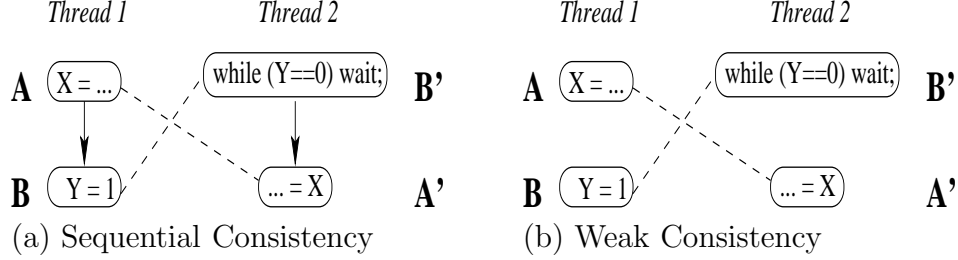


Figure 2.1: Access Order Graphs for a Multithreaded Program

Consider a program edge from an access A to an access B . This program edge says that A must appear to complete before B in a valid program execution. However, during execution, the thread performing accesses A and B may re-order them if this change is not observable, i.e. the program computation with the re-ordering is equivalent to the program computation without the re-ordering. Thus, assuming sequential consistency for both programs shown in Figure 2.2, the two accesses A and B in the code in part (a) may be re-ordered, but this is not the case for the code in part (b).

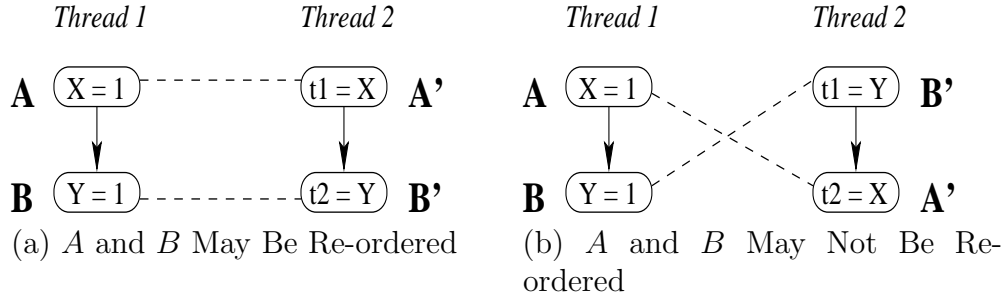


Figure 2.2: Re-orderings Allowed Under Sequential Consistency

Theorem 2.2.1 *For a program edge from A to B , re-ordering A and B may affect the program computation only if the access order graph contains a path from B to A that begins and ends with a conflict edge.*

Suppose there are no intra-thread dependences between A and B , and B executes first. If no other instruction accesses memory location A or B , then this re-ordering has no effect on program execution. The only way the execution of B may be noticed by another instruction and affect program execution is if there is a conflict edge between B and some access, say B' , that accesses the same location as B . Similarly, the execution of A may be noticed by another instruction only if there is some other conflicting access A' . Thus, for the re-ordering of A and B to be evident, there must be an execution of the program in which access B occurs before access B' , B' occurs before A' , and A' occurs before A . In terms of the access order graph, this translates to a path from B to A , such that this path begins and ends with a conflict edge. Therefore, if the access order graph contains a program edge from A to B , as well as a path from B to A that begins and ends with a conflict edge (as in Figure 2.2(b) but not in Figure 2.2(a)), then re-ordering the two accesses A and B may affect the program computation.

Note that the delay edge A to B and the path from B to A form a cycle in the access order graph. Shasha and Snir's algorithm to find delay pairs looks for such cycles in the access order graph. Theorem 3.9 of [SS88] shows that it is sufficient to consider only “critical” cycles that satisfy the following properties:

1. The cycle contains at most two accesses from any thread; these accesses occur at successive locations in the cycle.
2. The cycle contains either zero, two, or three accesses to any given memory location; these accesses occur in consecutive locations in the cycle.

Each program edge in a “critical” cycle is a *delay edge*, and the ordered pair of accesses corresponding to this program edge is a *delay pair*. Delay pairs represent the orders within a thread that must be enforced for any program execution to honor the semantics of the programming language memory consistency model.

From the first property of critical cycles, we can infer that conflict edges that begin and end the path from B to A in Theorem 2.2.1 must be edges between accesses in

different threads. This is because there is a program edge between A and B which means A and B are in the same thread. They are successive accesses in the cycle from A to B and back to A . So, for the cycle to be critical, all other accesses in the path from B back to A must be in a thread different from the thread that contains accesses A and B . Based on this, we can refine Theorem 2.2.1:

Theorem 2.2.2 *For a program edge from A to B , re-ordering A and B may affect the program computation only if the access order graph contains a path from B to A that begins and ends with a conflict edge, and these conflict edges are between accesses in different threads.*

Corollary *For a program edge from A to B , if no path exists from B to A that begins and ends with a conflict edge such that these conflict edges are between accesses in different threads, then the program edge from A to B is not a delay edge.*

2.3 Simplified Delay Set Analysis Applied to Real Programs

Precise delay set analysis finds the minimal set of delay edges that need to be enforced in a program, but it is not possible in a compiler for a general-purpose language. In real applications, the access order graph cannot always be accurately constructed. There are several reasons for this:

- The number of dynamic threads spawned during execution may be indeterminate at analysis time, or simply too large to individually account for each thread in the access order graph. Therefore, a single graph node may represent an access in the program when it is performed by multiple threads, instead of an individual node corresponding to each thread that performs the access.

- Most real programs include branching control flow such as loops and if-then constructs. Due to this, a graph node may represent multiple instances of shared memory accesses that may be performed during a program execution.
- For a language such as Java that includes references, a variable may contain a reference to different memory locations at different points in the execution. Two accesses using the same reference variable may not access the same memory location. Conversely, two accesses using different reference variables may in fact access the same memory location. Pointer analysis is needed to disambiguate memory locations accessed. Also, array dependence analysis is needed to determine if two array accesses refer to the element at the same array index. Perfect memory disambiguation is not possible for all programs at compile time because the location accessed at some point may depend on the specific execution path that a program follows, or it may depend on external inputs, or the subscript expression for an array access may be too complex to analyze. In such cases, the analysis must conservatively assume that an access may refer to any memory location in a set of possible locations. So a graph node may represent accesses to multiple shared memory locations.

Thus, nodes in the access order graph constructed by our compiler may represent multiple accesses in the program execution.

Definition 2.3.1 *Given a node N in the access order graph, an **instance** of N is a dynamic access that may be performed in some program execution, and this access is represented by N in the access order graph.*

Our access order graph conservatively includes a program edge (or a conflict edge) between nodes A and B if there exists a program edge (or a conflict edge) between some instance of A and some instance of B .

Even if the access order graph can be accurately constructed, precise delay set analysis is expensive to apply in practical systems if the number of threads in a program is large.

In [KY96], Krishnamurthy and Yelick show that Shasha and Snir’s precise delay set analysis is NP-complete by reducing the Hamiltonian path problem to the problem of finding the minimal delay set. They prove the following theorem:

Theorem 2.3.1 *“Given a directed graph G with n vertices, we can construct a parallel program P for n processors such that there exists a Hamiltonian path in G if and only if there exists a simple cycle in (the access order graph for) P .”*

Note that all critical cycles are simple cycles since they contain at most two consecutive accesses from a single thread. Thus, the execution time for precise delay set analysis is exponential in the number of threads in the program.

Our goal is to perform delay set analysis that is approximate and fast, but precise enough to generate code that performs well for sequential consistency. In our simplified analysis, we do not find cycles in the access order graph to determine the set of all possible delay edges. Instead, we consider each program edge one at a time, and apply a simple, conservative test to determine if it *cannot* be a delay edge. This test checks to see if the end-points of a potential cycle exist. If they do, we conservatively assume that the complete cycle exists, without verifying if it actually does exist in the access order graph.

We say a node M may occur before (or after) a node N if some instance of M may occur before (or after) some instance of N in a program execution. To test if a program edge from node A to node B cannot be a delay edge (see Figure 2.3), we determine if there exist two other nodes X and Y such that:

1. a conflict edge exists between A and X , and
2. a conflict edge exists between B and Y , and
3. some instance of A and some instance of X may occur in different threads, and
4. some instance of B and some instance of Y may occur in different threads, and

5. A may occur after X and Y, and
6. B may occur before X and Y, and
7. Y may occur before X.

If X and Y exist, then we conservatively assume that a path from B to A exists that completes a cycle in the access order graph, and so the program edge from A to B is a delay edge. If no such nodes X and Y exist, then it follows from the Corollary to Theorem 2.2.2 that the program edge from A to B cannot be a delay edge.




Figure 2.3: Simplified Cycle Detection

Note that if A may occur after X, and Y may occur before X, this does not necessarily mean that A may occur after Y. This is because our analysis is conservative. It may be the case that A always occurs only before X and only before Y, but our imprecise analysis can only determine that A always occurs before Y, and it cannot determine the relative order of A and X. Therefore, we test each of the orders enumerated above, including whether A may occur after Y, and whether B may occur before X.

Without synchronization analysis, the check for a delay edge is not very effective since it only tests for the first two conditions enumerated above. We use synchronization analysis, described in Chapter 3, to refine delay set analysis. Synchronization analysis determines threads that may concurrently execute in the program, and orders that are guaranteed between accesses in the program due to explicit synchronization programmed by the user. This information is used to test all the conditions for a delay edge enumerated above.

2.4 Illustrative Examples

Example 1

Figure 2.4 shows an example similar to one presented in [AG96]. This example illustrates transitive ordering required by sequential consistency, which assumes a total order over all accesses in a program execution. If access *B* of Thread 2 observes the value 1 written by *A*, and *D* in Thread 3 observes the value 1 written by *C* in Thread 2, then the value observed by *E* in Thread 3 must be the one written by *A*, or an access to *X* that executes after *A*. That is, the orders *A* to *B*, *B* to *C*, *C* to *D*, and *D* to *E* together imply the order *A* to *E*.




Figure 2.4: Example to Illustrate Cumulative Ordering

There is a cycle (A, B, C, D, E, A) in the access order graph. The program edges in this cycle from *B* to *C* and from *D* to *E*, are the delay pairs that need to be ordered for sequential consistency. Enforcing these delay pairs must also enforce the transitive order *A* to *E* in an execution where *A* executes before *B* and *C* executes before *D*. For this, enforcing the delay pair from *B* to *C* must ensure that *C* is delayed until the value returned by *B* is available to all processors, not just the processor that executes Thread 2.

Example 2

The access order graph in Figure 2.5 has two cycles: (D, A, B, C, E, F, D) and (D, A, B, C, D) . The nodes in the smaller cycle form a subsequence of the nodes in the bigger cycle. As

shown in [SS88], only the program edges in the smaller cycle are delay edges that need to be enforced. Thus, the program edge from E to F is not a delay edge. In our simplified delay set analysis, we check only for potential cycles, and do not determine any cycle in its entirety. Thus, our analysis will conservatively include the program edge from E to F as a delay edge.




Figure 2.5: Example to Illustrate a Non-minimal Cycle

Example 3

The access order graph in Figure 2.6 further illustrates the conservatism of our simplified analysis. This graph has no cycles but, assuming all threads execute concurrently, our analysis will determine the program edge from A to B to be a delay edge. This is because of the conflict edges between A and D , and B and E .




Figure 2.6: Example to Illustrate Conservatism of Our Simplified Approach

2.5 Implementation

2.5.1 Nodes in the Access Order Graph

There are two kinds of nodes in our access order graph: shared memory accesses and method calls. Thread escape analysis, described in Section 4.3.1, is used to determine instructions that refer to memory locations accessible by multiple threads in the program. Method call nodes represent all accesses to shared memory that may be performed as a result of the method invocation, i.e. directly in the method called, or indirectly in other methods invoked by the method called. When determining conflict edges for a method call node, we consider conflict edges due to any such access.

2.5.2 Edges in the Access Order Graph

When compiling a method, the control flow graph determines the program edges in that method. These program edges are between nodes that correspond to call instructions and shared memory access instructions in the method. Exceptions in Java affect the flow of control. In our implementation, the control flow graph includes edges to capture control flow for exceptions in the program³. Thus, program edges also include ordering requirements due to exceptions. We assume program edges exist between an instruction and all its successors in the control flow graph for the method. Each intra-procedural program edge is tested to check if it is a delay edge that must be enforced in the code generated.

A conflict edge exists between two nodes if both nodes may access the same memory location. We take advantage of user-defined types in Java to partition shared memory such that if two accesses refer to objects of the same type, then we assume they may access the same location in memory. Since Java is a strongly typed language, two accesses

³Refer to Section 3.5.1 for a discussion of how exceptions are handled in our analysis implementation.

to objects of different types cannot be to the same memory location. This strategy is effective because it groups together memory locations that are likely to have similar access properties, and it avoids the need to perform expensive alias analysis.

2.5.3 Algorithm

We implement our delay set analysis as part of the compiler described in Chapter 4. We analyze the program source and record a summary of information for each method. This information implicitly determines the access order graph needed for delay set analysis. Figure 2.7 shows how we compute this summary information for each method, and how we use this information to test for delay edges. When compiling a method, we individually test each program edge that is determined from the control flow graph of the method to check if it is a delay edge.

When we test for a delay edge from a node A to a node B , we explicitly determine conflict edges that involve at least one of A or B . We do not search for individual nodes that conflict with A or B , but look for methods that may contain accesses corresponding to conflicting nodes. Thus, there are at most m candidates in the set *ConflictMethods* in Figure 2.7, where m is the number of methods in the program. The method information we compute has a single entry to summarize all shared memory accesses in the method that read from an object of a specific type. Similarly, there is a single entry to summarize all shared memory accesses in the method that write to an object of a specific type.

The delay set analysis algorithm presented here is very simple. In Section 3.2.3, 3.3.2, and 3.4.2, we show how information about program synchronization can be used to make this analysis more accurate.

To Determine Method Summary Information:

1. Compute \forall method M , $DirectWrites(M)$ to be the set of all types T such that some access in M writes to a location of type T .
2. Compute \forall method M , $DirectReads(M)$ to be the set of all types T such that some access in M reads a location of type T .
3. Compute \forall method M , sets $AllWrites(M)$ and $AllReads(M)$ as follows:
 - (a) \forall method M , $AllWrites(M) = DirectWrites(M)$
 - (b) \forall method M , $AllReads(M) = DirectReads(M)$
 - (c) Repeat while any set $AllWrites(M)$ or $AllReads(M)$ changes:
 - \forall method M ,
 - \forall method N such that M calls N ,
 - i) $AllWrites(M) \cup = AllWrites(N)$
 - ii) $AllReads(M) \cup = AllReads(N)$

To Test For A Delay Edge From Node A to Node B :

Define, for an instruction S that corresponds to a node in the access order graph,

$ConflictMethods(S)$

$= \{M \mid T \in DirectWrites(M)\};$

S is a read access to a location of type T .

$= \{M \mid T \in DirectWrites(M) \text{ or } T \in DirectReads(M)\};$

S is a write access to a location of type T .

$= \{M \mid \exists T, (T \in AllReads(N) \text{ and } T \in DirectWrites(M)) \text{ or } (T \in AllWrites(N) \text{ and } (T \in DirectWrites(M) \text{ or } T \in DirectReads(M)))\};$

S is a method call that invokes method N .

If either $ConflictMethods(A)$ or $ConflictMethods(B)$ is an empty set, then the edge from A to B is not a delay edge. Else, it is a delay edge.

Figure 2.7: Algorithm for Simplified Delay Set Analysis

2.5.4 Complexity

The worst case time complexity of our simplified delay set analysis is bounded by $O(n)$, where n is the maximum number of access order graph nodes in the program. We test each intra-procedural program edge in the access order graph to check if it is a delay edge. Since Java is a structured language, and the maximum outdegree of each node due to control flow is a fixed constant, the number of edges we need to test is $O(n)$. A delay test between two nodes A and B searches for a pair of accesses X and Y that conflict with A and B respectively, and are possible end-points of a cycle. By itself, delay set analysis checks only to see if some conflicting access exists for both A and B, and it does not perform any further tests. Information for all accesses of a specific type can be summarized together. Since the type for A and B is known, it takes constant time to check if a conflicting access exists. Thus, the complexity is $O(n)$: $O(1)$ for each delay test, times $O(n)$ number of delay tests.

However, if synchronization analysis is used to refine the delay test, then all conflicting accesses may need to be determined. Our analysis summarizes access information for each method based on the types of accesses contained in the method. Instead of individual accesses X and Y, we search for methods that may contain conflicting accesses. Suppose there are m methods in the program. Then there are m^2 possible choices for a pair X and Y that may need to be tested for orders enforced by program synchronization. In this case, the complexity is $O(n * m^2 * s)$, where s is the time taken to determine ordering between two accesses based on synchronization information.

2.5.5 Type Resolution

We summarize method information based on the type of objects referenced by each access in the method. It is possible that the type of objects referenced by a particular access is concretely known only when the access actually executes. Our analysis conservatively assumes that such an access references objects that may be one of several types. To

determine a set of possible types for a reference, we consider all classes loaded in the system and check each one to see if it is compatible with the declared type of the reference (Section 4.3.2). If the type corresponding to a class is compatible with the reference type, then the reference may access an object of that type.

References to objects may be used to invoke a method call as well as to access shared memory. The type of the object accessed through the reference determines the method that is invoked by the call. As before, if we cannot concretely resolve the type of the object referenced, we conservatively assume that any of a set of possible types may be accessed, and thus any of a set of possible methods may be invoked by the call.

2.6 Optimizing for Object Constructors

Figure 2.8 shows the access order graph for an example program. Assume that the memory location X is initially accessible only by Thread 1. At point C, Thread 1 writes the reference to X into a shared memory location. Thread 2 reads this reference at point D, and then it can also access location X. The access order graph has three critical cycles: (B, C, D, E, B) , (A, B, E, F, A) , and (A, C, D, F, A) . Thus, all the program edges are delay edges to be enforced.

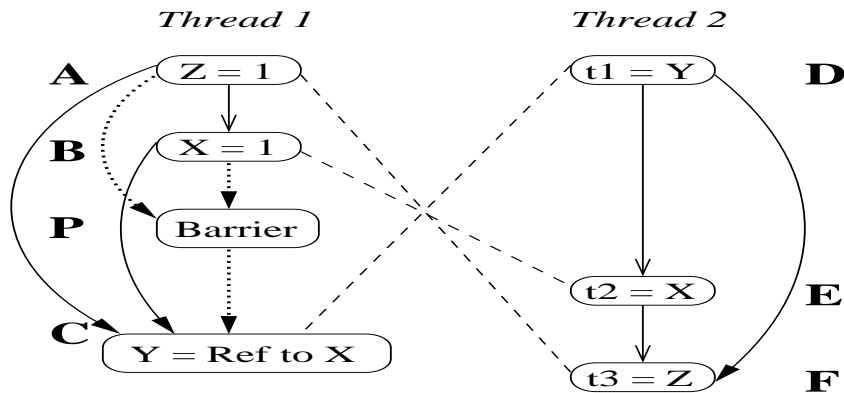


Figure 2.8: Delay Edges Related to the Point an Object Escapes

If B and E conflict, they must access the same location. B accesses location X , and the reference to X is not available outside of Thread 1 until C executes. Therefore, either C must execute before E , or there is no conflict edge between B and E . In the latter case, there is only one cycle in the program graph (A, C, D, F, A) , and the program edges in this cycle, from A to C and from D to F , need to be enforced.

Suppose B and E conflict, and C executes before E . Assume there is a memory barrier at point P in Thread 1 that ensures access B completes execution before C . The orders from B to C and from C to E transitively induce the order B to E on the conflict edge between B and E . As a result, the cycle (B, C, D, E, B) is no longer possible, and the delay edge from D to E need not be enforced.

The memory barrier at point P in Thread 1 also enforces the delay edge from A to C . This obviates the need to enforce the order from A to B . The orders A to C (due to the memory barrier) and C to E (due to the accessibility of location X) transitively result in the order A to E . This makes it impossible for an execution to have a shared memory access order of $B \rightarrow E \rightarrow F \rightarrow A$, where A occurs after E . Thus, the re-ordering of A and B is not visible to Thread 2 in any execution, and the delay edge from A to B need not be enforced.

In summary, if location X can be accessed only by Thread 1 until it makes the reference to X available to other threads at point C , and there is a memory barrier just before point C , we need not enforce the delay edges from A to B , and from D to E . This result can be used to reduce the number of delay edges.

It is common in Java programs for a new object to be created and initialized by one thread before it is made accessible to other threads. When a new object is created, memory is allocated for the object, and a special constructor method is invoked to initialize the object. If we determine that the reference to the newly created object (X in the previous example) is not shared by multiple threads at the point just before the constructor method returns (point P in the previous example), and we generate a

memory barrier at the return point of the constructor method, then we can apply the optimization illustrated by the example in Figure 2.8.

We do this in our implementation by using thread escape analysis (Section 4.3.1) to check, for each constructor method, if the reference to the object that is being constructed may escape at some point in the constructor method. We consider only those constructor methods for which the reference to the object being constructed does not escape in the constructor method. For each of these constructor methods, we insert a memory barrier instruction at the return points of the method. When performing delay set analysis, we can safely assume that delay edges do not involve any nodes within these constructor methods that access the object being constructed. Also, we disregard conflict edges directed into nodes within these constructor methods that access the object being constructed. Thus, the edge from D to E is determined not to be a delay edge. This optimization is useful in our implementation because our type-based alias analysis is conservative, and it is unable to eliminate conflict edges between accesses to different instances of objects of the same type. Figure 2.9 shows the effect this optimization has when compiling the benchmarks in Section 5.1.




Figure 2.9: Effect of Delay Set Analysis Optimization for Constructor Methods

The figure shows execution times for sequential consistency, normalized to the time taken by the base Jikes RVM implementation that uses weak consistency, for:

1. *NoOpt*: the case when our delay set analysis does not use the optimization for object constructors.
2. *WithOpt*: the case when our delay set analysis uses the optimization for object constructors.

We observe that the optimization improves performance for 5 out of 10 programs, and this improvement is 8.1% on average, ranging from 4% to 11.8%.

2.7 Related Work

Delay set analysis was first described by Shasha and Snir in [SS88]. Their algorithm is precise, but is shown to be NP-complete in the number of threads [KY96]. Krishnamurthy and Yelick [KY94, KY96] give a delay set analysis algorithm for SPMD programs. Their algorithm incorporates synchronization information, and uses the property that all threads execute the same code to achieve polynomial time complexity. Midkiff, Padua, and Cytron [MPC90] refine delay set analysis by considering array accesses and building a statement instance level flow graph. In this graph, edges between array accesses are labeled with a symbolic relation between the indices of the two accesses connected by the edge, and this information is used to reduce the set of shared memory access orders that need to be enforced. In [Mid95], Midkiff applies Kirchoff's Laws of Flow to the parallel program graph, and represents the path from a statement S back to itself as a system of equations. A linear system solver can then be used to check if the path may violate consistency requirements and the program edges along the path need to be enforced. This technique allows cycle detection in graphs considering all program accesses, including array accesses. In [CKY03], Chen, Krishnamurthy, and Yelick further present four polynomial time algorithms to perform delay set analysis for SPMD programs. One of

these algorithms use strongly connected components to achieve an analysis complexity of $O(n^2)$, where n is the number of shared memory accesses in the program. The other three algorithms target precision for programs that access array elements, and in the graph constructed for the analysis, edges between array accesses are labelled with some information that relates the indices of the two accesses. In [vP04], von Praun uses an object-based approach to simplify delay set analysis. In this approach, the analysis checks to see if there is an absolute order on all accesses to a memory location relative to all shared memory accesses, or a relative order between all accesses to two memory locations. The delay set analysis algorithm we present in this work has polynomial complexity and is applicable to MIMD programs.

Chapter 3

Synchronization Analysis

3.1 Problem Statement

In our synchronization analysis we aim to determine, for each shared memory access and each use of a synchronization primitive in the program code, constraints on the order in which they may be executed relative to one another. Then, by applying the semantics of the synchronization primitives, we can determine the order of execution of any two shared memory accesses.

This ordering information helps to restrict the conflict edges in the access order graph, and improve the precision of delay set analysis [KY96]. For example, the access order graph in Figure 3.1(a) has a cycle (A, B, E, F, A) , where A to B and E to F are delay edges. Thread 2 is spawned by Thread 1 at point C . Spawning of a thread implicitly synchronizes the creator thread and the new thread. So, point D which is the start of Thread 2 must execute after C . Java semantics ensure that all memory instructions in the program source that occur in the creator thread before the thread spawn point complete before the new thread starts. Thus, the values stored by A and B are visible before the new thread starts. Since E and F happen after D , and C executes before D , by transitivity E and F both happen after A and B . This enables us to direct the conflict

edges from A to F and from B to E , as shown in Figure 3.1(b). The cycle no longer exists, and no delay edges need to be enforced. Thus, synchronization helps reduce the number of orders that must be respected, which can improve the performance of program execution.

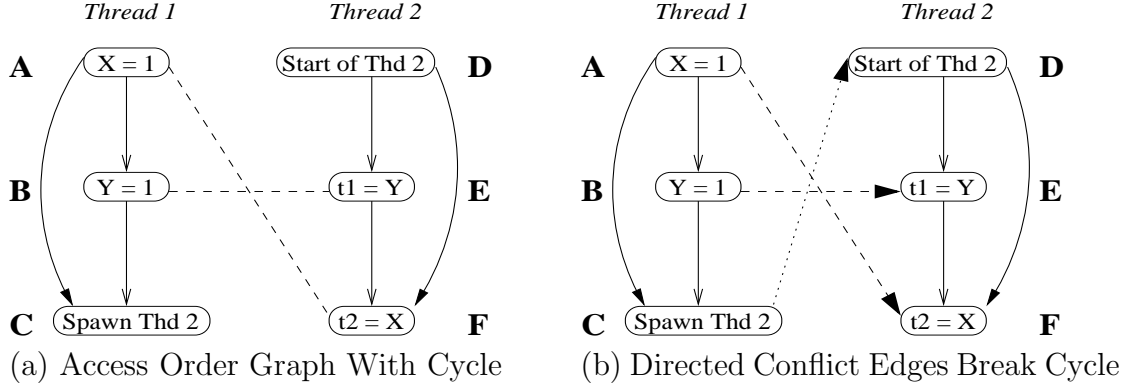


Figure 3.1: Use of Synchronization to Break Delay Cycles

Since we compile Java programs, we consider the following synchronization primitives that are part of the Java language:

- `thread start()` and `join()` calls, used to determine the program thread structure.
- `wait()`, `notify()`, and `notifyAll()` calls, used for event-based synchronization.
- `synchronized` blocks, used for lock-based synchronization.

When our compiler for sequential consistency generates machine code for these synchronization primitives, it ensures that a memory barrier is inserted before each `start()`, `notify()`, and `notifyAll()` call, after each `join()` and `wait()` call, and before and after each synchronized block.

In the next three sections, we describe our thread structure analysis, event-based synchronization analysis, and lock-based synchronization analysis respectively, and explain how we use the results of these analyses to determine delay edges. At the end of this chapter, we discuss some implementation issues for our analysis algorithms.

3.2 Thread Structure Analysis

Thread structure analysis determines orders enforced by thread spawning and thread termination. Java language semantics guarantee that when a thread is spawned via a thread `start()` call, all memory accesses of the creator thread that occur in the program source before the thread spawn point complete before the new thread starts. Also, if a thread `T` invokes a `join()` call to wait for another thread to terminate, then all memory accesses performed by the thread that terminates complete before `T` continues execution after the `join()`.

Figure 3.2 shows the execution timeline for an example program in which threads are spawned by a `start()` call `S`. The points `A`, `B`, and `C` represent static points in the program code. During an execution, a static instruction may be executed multiple times due to program control flow.

Definition 3.2.1 *For a static program statement S , an **instance** of S is some dynamic execution of S .*

Suppose we know from program control flow that:

- all instances of `A` must execute before any thread is spawned by `S`,
- all instances of `B` must execute after all instances of `S`, and before any `join()` point in the program that waits for all threads spawned at `S` to terminate, and
- all instances of `C` must execute after a `join()` point in the program that waits for all threads spawned at `S` to terminate.

Since shared memory accesses are not re-ordered across thread start and join points, we can infer the execution orders from `A` to `B`, from `B` to `C`, and from `A` to `C`.

We use dataflow analysis to determine orderings due to the thread structure of the program. For each pair (S, P) , where S is a thread `start()` call, and P is any point in

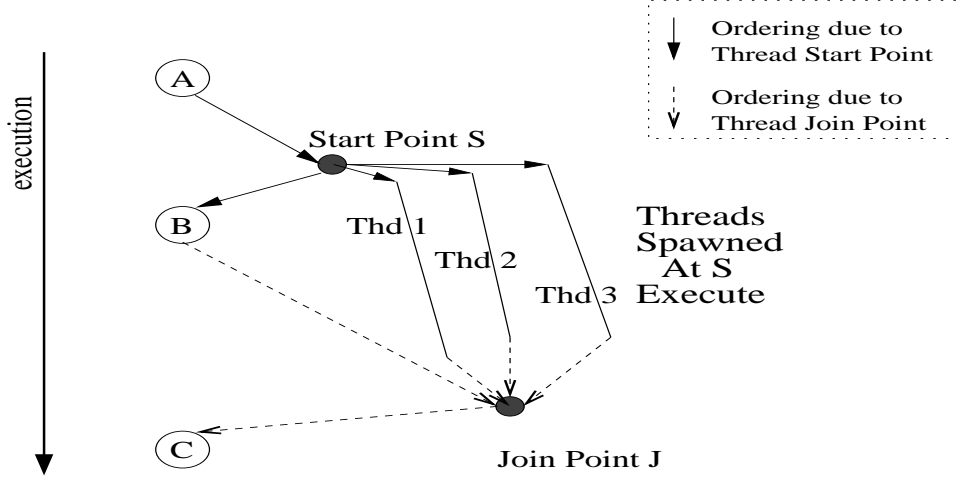


Figure 3.2: Orders Implied by Program Thread Structure

the code, we determine the relative order of execution of `P` and all threads started at `S`. Before we perform dataflow analysis, we need to identify all thread `start()` calls and thread `join()` calls in the source code, and match the `join()` calls to corresponding `start()` calls.

We present next, in Section 3.2.1, how we determine the `start()` calls that a `join()` matches. Then, in Section 3.2.2, we describe the algorithm to compute orders between program points and threads spawned by a `start()` call. In Section 3.2.3, we show how to apply the analysis results to determine ordering information that can be used in our delay set analysis.

3.2.1 Matching Thread Start and Join Calls

We say that a `join()` matches a `start()` if the `start()` and `join()` are both invoked on the same object. Pointer analysis is needed to determine if two objects referenced at different points in the code must be the same object. However, we do not perform pointer analysis in our implementation. In the absence of pointer analysis, two references are to the same object if they use the same program variable, this variable can be accessed by a single thread, and flow analysis can ensure that the variable is not re-assigned on

any path from the first reference to the second reference. We use thread escape analysis and a conservative intra-procedural flow analysis, and attempt to match a `join()` with a `start()` only if both occur in the same method. Note that a `join()` may match multiple `start()` calls that occur on different control flow paths reaching the `join()`.

If a `start()` is in a loop, then it must be invoked on a distinct object for each iteration of the loop ¹. So, a set of threads is associated with the `start()`. We say that a `join()` matches this `start()` if it is invoked on exactly the same set of objects as the `start()` is invoked on. A matching `join()` may be in the same loop as the `start()`, as in Figure 3.3(a), or in a different loop, as in Figure 3.3(b).

<pre> for (...) { x.start(); ... x.join(); } </pre> <p style="text-align: center;">(a)</p>	<pre> for (int i = LB; i < UB; i++) { thd[i].start(); } for (int i = LB; i < UB; i++) { thd[i].join(); } </pre> <p style="text-align: center;">(b)</p>
--	--

Figure 3.3: Matching `join()` for a `start()` in a Loop

In the case of Figure 3.3(b), if we know that the array `thd` can be accessed by a single thread, we can determine that the set of objects on which the `start()` and `join()` are invoked is exactly the same by observing that for both loops, the bounds and increments for the loop index variable `i` are same, and the array index expression used to access the array `thd` is also the same. In our implementation, we use simple pattern matching to detect this case.

When doing the analysis, we consider only those `join()` calls that are matched with some `start()`. It is safe to leave some `join()` unmatched because in that case, we consider fewer synchronizations, and this leads to a conservative, but correct, analysis.

¹Java semantics do not allow the thread corresponding to a specific object to be started multiple times.

3.2.2 Algorithm

We use an inter-procedural inter-thread analysis to compute, for each program point P , the sets $StartsAfter(P)$, $JoinsBefore(P)$, and $Concurrent(P)$.

Definition 3.2.2 *For a program point P , the set **StartsAfter** (P) is defined to be the set of all thread **start()** calls S such that some instance of S may execute after some instance of P .*

Definition 3.2.3 *For a program point P , the set **JoinsBefore** (P) is defined to be the set of all thread **start()** calls with a matching **join()** such that some instance of the **join()** may execute before some instance of P .*

Definition 3.2.4 *For a program point P , the set **Concurrent** (P) is defined to be the set of all thread **start()** calls S such that some instance of S starts a thread that may concurrently execute with some instance of P .*

In our analysis, we compute supersets that contain the sets defined above:

- For a thread **start()** S and any program point P , $S \in StartsAfter(P)$ if S does not dominate P .
- For a thread **start()** S and any program point P , $S \in JoinsBefore(P)$ if a thread **join()** that matches with S executes on some path from the program entry point to P .
- For a thread **start()** S and any program point P , $S \in Concurrent(P)$ if S executes on some path from the program entry point to P , and a **join()** that matches with S does not execute on that same path after S and before P .

Note that these are *may* sets, not *must* sets. It is conservative, but correct, to include any **start()** in any of these sets.

Initially, at the entry point of the application, say E , there is a single application thread and no program statements have executed. So $StartsAfter(E)$ is the set of all start calls in the program, and $JoinsBefore(E)$ and $Concurrent(E)$ are empty sets.

We define Gen and $Kill$ functions that determine how a program statement transforms each of the three sets we compute. For a statement S , these functions help compute the corresponding $StartsAfter()$, $JoinsBefore()$, and $Concurrent()$ sets, for successors of S in the control flow graph.

Definition 3.2.5 *For a program statement S , the function **Gen_Concurrent** (S) defines the set of **start()** calls that should be included in set $Concurrent(P)$ for any point P that is a successor of S in the control flow graph.*

Definition 3.2.6 *For a program statement S , the function **Kill_Concurrent** (S) defines the set of **start()** calls that should be removed from set $Concurrent(P)$ for any point P that is a successor of S in the control flow graph.*

Definition 3.2.7 *For a program statement S , the function **Gen_JoinsBefore** (S) defines the set of **start()** calls that should be included in set $JoinsBefore(P)$ for any point P that is a successor of S in the control flow graph.*

Definition 3.2.8 *For a program statement S , the function **Kill_StartsAfter** (S) defines the set of **start()** calls that should be removed from set $StartsAfter(P)$ for any point P that is a successor of S in the control flow graph.*

Consider the program control flow graph in Figure 3.4. Point E is the program entry point, point S is a thread **start()** call, and point J is a **join()** call that is matched with S . For point Q that occurs after the thread **start()** S , S should be removed from the set $StartsAfter(Q)$, and should be added to the set $Concurrent(Q)$. For point R that occurs after the thread **join()** J , S should be added to the set $JoinsBefore(R)$, and should be removed from the set $Concurrent(R)$. Thus, we get the definitions for Gen and $Kill$ as shown in Figure 3.5.

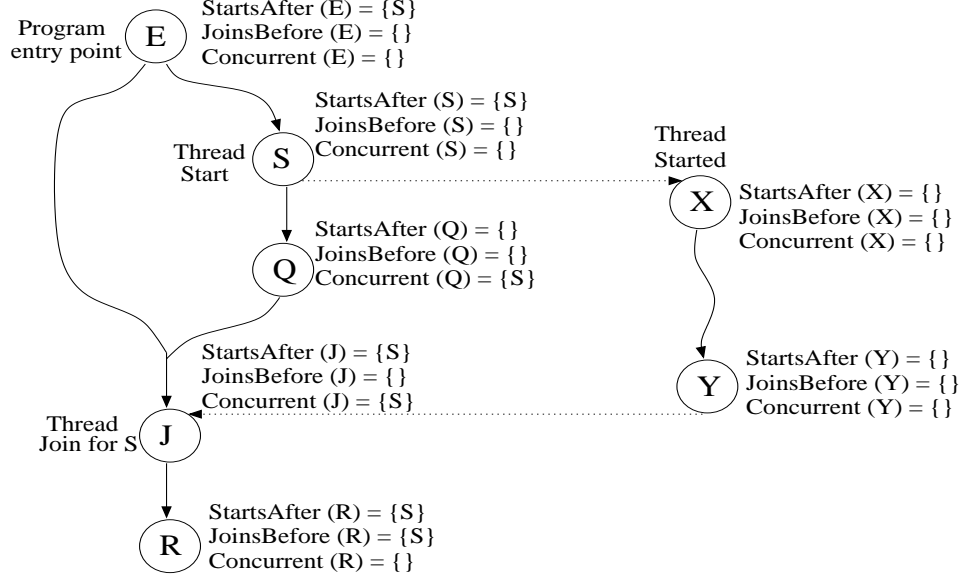


Figure 3.4: Effect of Thread Start and Join Calls

Point X in Figure 3.5 is the initial point in the new thread spawned by the `start()` at S . It illustrates the case when S is not contained in any of the sets $StartsAfter(X)$, $JoinsBefore(X)$, and $Concurrent(X)$. S does not occur after X since S must execute before the new thread is spawned. X executes as part of the thread spawned at S , so the thread spawned at S cannot finish before X . Also, assuming a single thread is spawned at S , X cannot be concurrent with itself.

In Figure 3.6, we give a simple algorithm that uses the *Gen* and *Kill* definitions to compute the sets $StartsAfter(P)$, $JoinsBefore(P)$, and $Concurrent(P)$ at each program point P .

We use dataflow analysis to propagate sets from the application entry point to all reachable points in the program. If a point is reachable from multiple other points, we propagate the union of the corresponding sets at all predecessor points in the method control flow graph. We also propagate information from each `start()` call to the entry point of the `run()` method for the thread type that is spawned by the `start()`. We iterate over the available code, propagating information until no sets change their values.

$StartsInMethod(M)$ is the set of all **start()** calls that may execute when method M is invoked, either directly in M , or indirectly in a method called in M .

$JoinsInMethod(M)$ is the set of all **start()** calls such that a matched **join()** must execute when method M is invoked, either directly in M , or indirectly in a method called in M .

$Kill_StartsAfter(x) = \{x\}$, x is a thread start() call
$= StartsInMethod(M)$, x is a call to method M
$= \phi$, otherwise.

$Gen_JoinsBefore(x) = \{y\}$, x is a matching join() for thread start() call y
$= JoinsInMethod(M)$, x is a call to method M
$= \phi$, otherwise.

$Gen_Concurrent(x) = \{x\}$, x is a thread start() call
$= StartsInMethod(M)$, x is a call to method M
$= \phi$, otherwise.

$Kill_Concurrent(x) = \{y\}$, x is a matching join() for thread start() call y
$= JoinsInMethod(M)$, x is a call to method M
$= \phi$, otherwise.

Figure 3.5: *Gen* and *Kill* Functions for Thread Structure Analysis

1. \forall method M , compute $StartsInMethod(M)$ and $JoinsInMethod(M)$ as defined in Figure 3.5.
2. For each program point P , initialize the sets $StartsAfter(P)$, $JoinsBefore(P)$, and $Concurrent(P)$ to be empty.
3. For the application entry point E , include all **start()** calls in $StartsAfter(E)$.
4. Repeat the following while there is a change in the sets for any program point:
 - \forall statement S in the program code
 - (a) $StartsAfter(S) \cup = \bigcup_{p \in Predecessor(S)} (StartsAfter(p) - Kill_StartsAfter(p))$
 - (b) $JoinsBefore(S) \cup = \bigcup_{p \in Predecessor(S)} (JoinsBefore(p) \cup Gen_JoinsBefore(p))$
 - (c) $Concurrent(S) \cup = \bigcup_{p \in Predecessor(S)} ((Concurrent(p) - Kill_Concurrent(p)) \cup Gen_Concurrent(p))$
 - (d) If S is a **start()**, and R is the entry point of the **run()** method for the thread spawned by S , then
 - (i) $StartsAfter(R) \cup = (StartsAfter(S) - \{S\})$
 - (ii) $JoinsBefore(R) \cup = JoinsBefore(S)$
 - (iii) $Concurrent(R) \cup = Concurrent(S)$

Figure 3.6: Algorithm to Compute Orders for Thread Starts and Joins

The analysis converges because there are a fixed number of `start()` calls in the code, and once a `start()` is added to a set due to propagation, it is never removed from it.

Predecessor Relation

The predecessor relation used in the analysis depends on the programming language memory model. For sequential consistency, it is defined as follows:

- The predecessor set of the first statement in a method includes the set of call sites that invoke the method.
- The predecessor set of the first statement of a basic block B includes the last statement of each basic block that is a predecessor of B in the method control flow graph.
- The predecessor set of all other statements includes the single statement that precedes it in its basic block.

However, in the general case, we need to construct a partial order graph for statements within each basic block using knowledge of the program edges that are implied by the programming language memory model.

Complexity

The worst case time complexity of our thread structure analysis is bounded by $O(n^2 * s)$, where n is the number of program points and s is the number of thread start calls in the program source. The analysis iterates until none of the sets being computed change. Each iteration does a pass over the program source and takes $O(n)$ time. The dataflow analysis is monotonic and a set can change only when an element is added to it. There are three sets computed for each program point n , and the s thread start calls in the program form the universe of elements that can be added to these sets. In the worst

case, a single element will be added to a single set in each iteration. Thus, the analysis needs at most $O(n * s)$ iterations, and the overall complexity is $O(n^2 * s)$.

3.2.3 Using Thread Structure Information in Delay Set Analysis

The results of thread structure analysis are used to refine our delay set analysis. In Section 3.2.3.1, we describe how these results can be applied to eliminate conflict edges. In Section 3.2.3.2, we describe how these results are used to determine an order between two accesses in the program. We illustrate these use cases with an example in Section 3.2.3.3.

3.2.3.1 Eliminating Conflict Edges

To help in determining if conflict edges can be eliminated, we determine for each thread type whether it satisfies the *single thread constraint*. We say that a thread type satisfies the single thread constraint if at most one thread of that type can execute at any given time. In Java, when a new thread is spawned, it begins execution by invoking a specific `run()` method that corresponds to the type of the thread spawned. A thread type `T` satisfies the single thread constraint if the `run()` method for `T` is not concurrent with any thread `start()` call that may directly or indirectly spawn a thread of type `T`. This information is readily available in the summary set $Concurrent(R)$, where R is any point that executes when the `run()` method corresponding to `T` is invoked. Section 3.5 describes summary sets and how they are computed.

The special thread type that distinguishes the main thread started at the program entry point will trivially satisfy the single thread constraint, since there are no thread `start()` calls corresponding to it.

We can eliminate conflict edges between two shared memory accesses that both execute as part of a single thread type, and where that thread type satisfies the single thread

constraint. This is because the single thread constraint implies that the two accesses are performed by the same thread, and in our delay set analysis, we search for conflict edges that occur between accesses performed in different threads.

3.2.3.2 Ordering Program Accesses

Definition 3.2.9 *For a thread type T , **Code** (T) is the set of all program statements such that an instance of the statement may be executed by a thread of type T in some program execution.*

Definition 3.2.10 *For a thread type T , **ThreadsSpawnedBy** (T) is the set of all thread types T' such that a thread of type T' may be spawned when a thread of type T executes. This includes threads spawned directly by a **start()** in **Code** (T), as well as threads spawned indirectly by a **start()** in any other thread that may be spawned as a result of the execution of a thread of type T .*

Consider two accesses X and Y . We cannot determine an order between these two accesses if:

- one of the accesses (say X) may be concurrent with a thread of type T , and
- the other access (say Y) is part of the code that may execute when a thread of type T is spawned, i.e. $Y \in \text{Code}(T)$ or $\exists T', T' \in \text{ThreadsSpawnedBy}(T)$ and $Y \in \text{Code}(T')$.

Otherwise, we may be able to order the two accesses if all instances of X or Y execute before some **start()**, or all of them execute after a **join()** that matches some **start()**.

If an instance of X in thread $T1$ and an instance of Y in another thread $T2$ are ordered by the thread structure of the program, then this order must be evident either from the ordering information for X and Y with respect to the **start()** that spawns $T1$, or from the ordering information for X and Y with respect to the **start()** that spawns

T2. To determine if there is an order between all instances of X and all instances of Y , we need only consider ordering information for X and Y with respect to **start()** calls S such that S may spawn a thread of type T , and either $X \in \text{Code}(T)$ or $Y \in \text{Code}(T)$. We can order $X \rightarrow Y$ (or $Y \rightarrow X$) if the *same* order is determined with respect to each of these **start()** calls S .

In Figure 3.7, we list all the possible orderings that may be inferred for two program accesses X and Y using their ordering information with respect to one particular **start()**, say S . The first row in the figure corresponds to the case when S is not contained in any of the sets $\text{StartsAfter}(X)$, $\text{JoinsBefore}(X)$ or $\text{Concurrent}(X)$. This implies that X is part of the code that may execute when a thread is spawned by S (refer to the discussion for the example in Figure 3.4). The third column in the figure corresponds to the case when S is contained in $\text{JoinsBefore}(Y)$, but not in $\text{StartsAfter}(Y)$ or $\text{Concurrent}(Y)$. Thus, the entry in the figure for the first row and third column determines the order $X \rightarrow Y$, since it corresponds to the case when X is part of the code executed by a thread spawned at S , and when all instances of Y execute after all threads spawned at S have finished execution. The other entries in the table in Figure 3.7 can be similarly interpreted.

We can infer some of the orders in Figure 3.7 only for conflict edges in our simplified delay set analysis. To infer these orders, we apply the extra check that is described next.

Refined Check to Order Conflict Edges

Consider a conflict edge between two accesses X and Y . In cases where all instances of one access (say X) occur before a **start()**, and instances of the other access (say Y) occur before the **start()** or concurrent with a thread spawned by the **start()**, we can apply an extra test to improve the precision of the orders determined. Let the type of the thread spawned by the **start()** be T . Then we check to see if all the following conditions hold:

A (Y) => StartsAfter (Y)			B (Y) => JoinsBefore (Y)				C (Y) => Concurrent (Y)			
<div>Access Y \ Access X</div>			A(Y) = {}	A(Y) = {}	A(Y) = {}	A(Y) = {}	A(Y) = {S}	A(Y) = {S}	A(Y) = {S}	A(Y) = {S}
			B(Y) = {}	B(Y) = {}	B(Y) = {S}	B(Y) = {S}	B(Y) = {}	B(Y) = {}	B(Y) = {S}	B(Y) = {S}
			C(Y) = {}	C(Y) = {S}	C(Y) = {}	C(Y) = {S}	C(Y) = {}	C(Y) = {S}	C(Y) = {}	C(Y) = {S}
A(X)={}	B(X)={}	C(X)={}			X -> Y		Y -> X			
A(X)={}	B(X)={}	C(X)={S}			X -> Y		Y -> X			
A(X)={}	B(X)={S}	C(X)={}	Y -> X	Y -> X		(Y -> X) ***see below	Y -> X	Y -> X		
A(X)={}	B(X)={S}	C(X)={S}			(X -> Y) ***see below		Y -> X			
A(X)={S}	B(X)={}	C(X)={}	X -> Y	X -> Y	X -> Y	X -> Y		(X -> Y) ***see below		
A(X)={S}	B(X)={}	C(X)={S}			X -> Y		(Y -> X) ***see below			
A(X)={S}	B(X)={S}	C(X)={}								
A(X)={S}	B(X)={S}	C(X)={S}								

These orders are contingent upon the refined check described in the text.

Figure 3.7: Orders Inferred From Thread Structure Analysis

1. all instances of X may execute as part of a single thread type, say T' .
2. all instances of Y may execute as part of at most two thread types, T and T' .
3. T' satisfies the single thread constraint.
4. The order $X \rightarrow R$ can be inferred, where R is the entry point of the `run()` method that corresponds to thread type T .

If all of the above conditions are satisfied, then $X \rightarrow Y$.

In the example code that follows, the access to x at S2 occurs before the `start()` at S4. The access to x at S6 in method `foo()` can occur at points S1, S5, and T1. Thus, S6 may occur before S4, concurrent with the thread started at S4, and as part of the thread started at S4. However, only the access S6 from point T1 is a candidate for a conflict edge between S2 and S6. This is because in our delay set analysis, we only consider conflict edges that exist between accesses in different threads, and S1 and S5 occur in the same thread as S2. Conditions 1-4 of the check above ensure that S1 and S5 are in the same thread as S2, and that T1 happens after S2. Thus, $S2 \rightarrow S6$.

```

main () {
    S1:  foo ();
    S2:  access to x;
    S3:  T t = new T ();
    S4:  t.start ();
    S5:  foo ();
}

class T extends Thread {
    void run () {
        T1:  foo ();
        ...;
    }
}

foo () {
    S6:  access to x;
}

```

A similar reasoning applies for the case when all instances of one access occur after a `join()` that matches a `start()` S , and instances of the other access occur concurrent with the threads spawned by S , or after the `join()`.

3.2.3.3 An Example

The example in Figure 3.8 shows the control flow graph of a program where the main thread spawns several threads at point D which is in a loop. The main thread and the spawned threads both make calls to method Z and execute point F in the code. The figure also shows the results of our thread structure analysis for each point in the graph. Note that D is propagated to the *Concurrent()* sets for all points in the loop, and then to point E and F . This captures the case when multiple threads spawned at a point are concurrent with one another.

We cannot infer any orders among points C , E , and F . This is because the *Concurrent()* sets at these points all include D , and E and F are part of the code executed by a thread spawned at D . Thus, the first condition to order conflict edges, described in Section 3.2.3.2, is not satisfied.

At point B , we observe that only the set *StartsAfter*(B) contains D , so we can deduce that B must execute before any thread is spawned at D . Also, at point E , only the set *Concurrent*(E) contains D , so we can deduce that E must execute concurrent with or as part of a thread spawned at D . Thus, $B \rightarrow E$ is an order imposed by thread start synchronization, and we can infer this using the table in Figure 3.7.

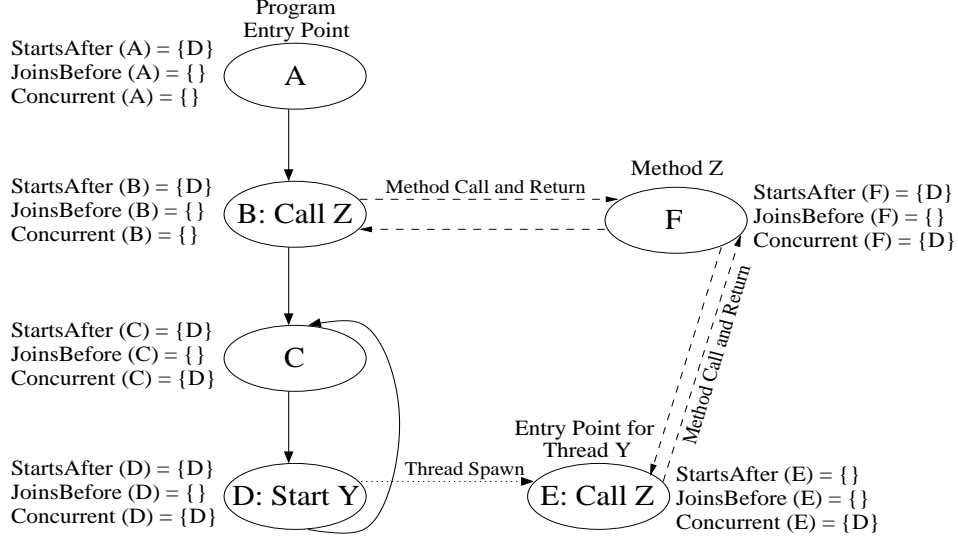


Figure 3.8: Example to Infer Orders From Thread Structure Synchronization

The points A , B , C , and D execute as part of the main thread that satisfies the single thread constraint (Section 3.2.3.1). We do not consider any conflict edges among them, since in our delay set analysis, we only consider conflict edges that occur between points that execute in different threads.

We can use the refined check described in Section 3.2.3.2 to determine the order $B \rightarrow F$. B only executes as part of the single main thread. F executes as part of two thread types: the single main thread, and the threads spawned at point D . E is the entry point of the threads spawned at D , and we previously inferred the order $B \rightarrow E$. Thus, all the conditions for the refined check are satisfied, and we infer $B \rightarrow F$.

3.3 Event-based Synchronization

In Java, event-based synchronization is provided by `wait()`, `notify()`, and `notifyAll()` calls. A thread that needs to wait for an event invokes the `wait()` call on the corresponding Java object. A thread that needs to signal a particular event invokes the `notify()` or `notifyAll()` call on the corresponding Java object. The `notify()` unblocks a single

thread that is waiting on that particular object, and the `notifyAll()` unblocks all threads that are waiting on that particular object.

Java language semantics prevent re-ordering of shared memory accesses across `wait()`, `notify()`, and `notifyAll()` calls. They guarantee that when a `notify()` or `notifyAll()` call executes, all shared memory accesses that occur in the notifying thread before the call complete before a thread is unblocked by the call. Also, shared memory accesses that occur in the waiting thread after the `wait()` call do not execute before the `notify()` or `notifyAll()` call that unblocks the waiting thread. This implies that if all instances of an access X in the program source execute before a `notify()` or `notifyAll()` call, and this call must unblock a particular `wait()`, and all instances of an access Y in the program source execute after the `wait()`, then $X \rightarrow Y$ is an order.

We use dataflow analysis to determine orderings due to event-based synchronization. For each pair (S, P) , where S is a `wait()`, `notify()`, or `notifyAll()` call, and P is any point in the code, we determine the relative order of execution of P and S .

3.3.1 Algorithm

The analysis first identifies all `wait()`, `notify()`, and `notifyAll()` calls in the program source.

Definition 3.3.1 *For a program point P , the set **SyncsAfter** (P) is defined to be the set of all `wait()`, `notify()`, and `notifyAll()` calls such that some instance of the call may execute after some instance of P .*

Definition 3.3.2 *For a program point P , the set **SyncsBefore** (P) is defined to be the set of all `wait()`, `notify()`, and `notifyAll()` calls such that some instance of the call may execute before some instance of P .*

For each program point P , we compute the sets $SyncsAfter(P)$ and $SyncsBefore(P)$ to be supersets that contain the sets defined above:

- For a `wait()`, `notify()`, or `notifyAll()` call S and any program point P , $S \in \text{SyncsAfter}(P)$ if S does not dominate P with respect to the entry point of a thread that executes both S and P .
- For a `wait()`, `notify()`, or `notifyAll()` call S and any program point P , $S \in \text{SyncsBefore}(P)$ if S executes on some path from the thread entry point to P , for a thread that executes both S and P .

Note that these are *may* sets, not *must* sets. It is conservative, but correct, to include any `wait()`, `notify()`, or `notifyAll()` call in any of these sets.

We determine the $\text{SyncsAfter}(P)$ sets at each program point P by computing for each thread type T in the program, the dominators for the flowgraph rooted at the entry point of T . Since we are interested only in `wait()`, `notify()`, and `notifyAll()` statements, we need to only consider these statements to include in the dominator sets computed. For a program point P and a `wait()`, `notify()`, or `notifyAll()` call S , if S is not a dominator of P in the dominator sets computed for some thread type T that may execute P , then we include S in the set $\text{SyncsAfter}(P)$.

We compute the $\text{SyncsBefore}(P)$ sets at each program point P using a dataflow analysis that is analogous to the one used in Section 3.2.2 to compute the $\text{JoinsBefore}()$ sets. The event-based synchronization analysis iterates over the code for each thread type in the program, one at a time.

$\text{MaySyncsInMethod}(M)$ is the set of all `wait()`, `notify()`, and `notifyAll()` calls that may execute when method M is invoked, either directly in M , or indirectly in a method called when M executes. We define the function $\text{Gen_SyncsBefore}(P)$ that determines, for any program statement P , how the execution of P transforms the sets we compute:

$$\begin{aligned} \text{Gen_SyncsBefore}(x) \\ = \{x\} \quad , x \text{ is a } \text{wait}(), \text{notify}(), \text{ or } \text{notifyAll}() \text{ call} \end{aligned}$$

$$\begin{aligned}
&= \text{MaySyncsInMethod}(M) \quad , x \text{ is a call to method } M \\
&= \phi \quad , \text{otherwise.}
\end{aligned}$$

Figure 3.9 outlines an algorithm that uses this definition to compute sets $\text{SyncsBefore}(P)$ at each program point P . The predecessor relation used in the algorithm is the same as that described in Section 3.2.2.

1. \forall method M , compute $\text{MaySyncsInMethod}(M)$.
2. \forall program point P , initialize the set $\text{SyncsBefore}(P)$ to be empty.
3. \forall thread type T in the program:
Repeat while there is a change in the sets for any program point $\in \text{Code}(T)$:
 \forall statement $S \in \text{Code}(T)$:
 $\text{SyncsBefore}(S) \cup =$
 $\bigcup_{p \in \text{Predecessor}(S)} (\text{SyncsBefore}(p) \cup \text{Gen_SyncsBefore}(p))$

Figure 3.9: Algorithm to Compute Orders for Event-based Synchronization

For a thread of type T that cannot execute any instance of a particular `wait()`, `notify()`, or `notifyAll()` call S , we conservatively include S in sets $\text{SyncsAfter}(P)$ and $\text{SyncsBefore}(P)$, at each point $P \in \text{Code}(T)$. Due to this, information from one thread type does not affect the results for program points in other thread types. Thus, our event-synchronization analysis need not propagate information at thread start sites to the entry point of the newly started thread type. This does not affect precision because thread structure analysis already accounts for orders implied by thread spawning and termination.

Complexity

The worst case time complexity of our event-based synchronization analysis is $O(t * n^2 * s)$ where t is the number of thread types in the program, n is the maximum number of

program points in the code for any thread type, and s is the number of `wait()`, `notify()`, and `notifyAll()` calls in the program. The analysis iterates until convergence over the code for each thread type individually. Each iteration does a pass over the code for a thread type and takes $O(n)$ time. The dataflow analysis is monotonic, and the sets $SyncsBefore(P)$ computed at each point P can change only when an element is added to them. There can be at most s elements in any of these sets. Thus, in the worst case when a single element is added to a single set in an iteration, we require $O(t * n * s)$ iterations, and the overall complexity is $O(t * n^2 * s)$.

3.3.2 Using Event-based Synchronization in Delay Set Analysis

For a program point P and a `wait()`, `notify()`, or `notifyAll()` call S , P occurs in the program only before S ($P \rightarrow S$) if $S \in SyncsAfter(P)$ and $S \notin SyncsBefore(P)$.

For a program point P and a `wait()`, `notify()`, or `notifyAll()` call S , P occurs in the program only after S ($S \rightarrow P$) if $S \notin SyncsAfter(P)$ and $S \in SyncsBefore(P)$.

For any two shared memory accesses X and Y , we can use event-based synchronization to infer the order $X \rightarrow Y$ if:

- X occurs in the program source only before a `notify()` or `notifyAll()`, say N ,
- Y occurs in the program source only after a `wait()`, say W , and
- N definitely is the synchronization operation that causes W to exit its wait.

N can cause W to exit its wait if both of them may occur concurrently, and are invoked on the same Java object. In general, there can be multiple points in the code where a `wait()` (or a `notify()` or `notifyAll()`) call are invoked on the same Java object. Also, we need alias analysis to determine if two references may in fact refer to the same Java object. Thus, during analysis we conservatively determine a set of `wait()` calls

that correspond to each `notify()` and `notifyAll()` call, and a set of `notify()` and `notifyAll()` calls that correspond to each `wait()` call.

Definition 3.3.3 For a `notify()` or `notifyAll()` call N , **WaitSet** (N) is the set of all `wait()` calls W such that N and W may be invoked on the same object, and thread structure analysis does not determine an order $N \rightarrow W$ or $W \rightarrow N$.

Definition 3.3.4 For a `wait()` call W , **NotifySet** (W) is the set of all `notify()` and `notifyAll()` calls N such that W and N may be invoked on the same object, and thread structure analysis does not determine an order $N \rightarrow W$ or $W \rightarrow N$.

For accesses X and Y , we can determine an order $X \rightarrow Y$ if either of the following is valid:

- there exists a `wait()` call W such that $W \rightarrow Y$, and for each call N in the set *NotifySet* (W), $X \rightarrow N$, or
- there exists a `notify()` or `notifyAll()` call N such that $X \rightarrow N$, and for each call W in *WaitSet* (N), $W \rightarrow Y$.

3.3.2.1 Limited Use in Java Programs

In practice, our event synchronization analysis is conservative when applied to Java programs. This is due to several reasons:

- The common usage for a `wait()` call is in a while loop, with some condition that determines if the loop will iterate and the `wait()` will be invoked. In these cases, it is possible that the loop performs zero iterations and the `wait()` is never invoked. So, there is no possibility of determining synchronization orders based on such conditional `wait()` calls.

- Java provides a timed version of the `wait()` call that is used more often in code than the untimed version. This timed `wait()` unblocks by default when a specified amount of time has elapsed. We cannot use it to determine orders due to synchronization.
- We use alias analysis to determine the sets $WaitSet(N)$ and $NotifySet(W)$ that are needed to apply the results of the analysis. The precision of the alias analysis impacts the utility of event-based synchronization analysis. Alias analysis is expensive to compute with high precision.

3.4 Lock-based Synchronization

Lock-based synchronization analysis uses properties of *synchronized blocks* in Java to eliminate some delay edges. A synchronized block in Java has a lock associated with it. This lock is acquired before the block begins execution, and is released when the block finishes execution. Synchronized blocks enforce the following properties:

1. They provide mutual exclusion for code blocks that are synchronized using the same lock.
2. They guarantee that all memory accesses in the source program prior to the end of the synchronized block happen before the lock for the synchronized block is released.
3. They guarantee that no memory access in the source program after the beginning of the synchronized block happens before the lock for the synchronized block has been acquired.

In our analysis, we compute a locking set at each program point that includes lock objects, corresponding to synchronized blocks, that have been acquired and not yet released when *any* instance of that program point executes. In the next section, we

describe our algorithm to compute these locking sets. Then, in Section 3.4.2, we show how these sets can be used to eliminate some conflict edges in our delay set analysis.

3.4.1 Algorithm

Figure 3.10 shows how to compute the sets of locks acquired due to synchronized blocks at each program point. For each statement in a synchronized block that uses the lock associated with an object L , we add L to the set of acquired locks for that statement. If the statement is a method call, then L is recursively propagated using set intersection to all statements reachable from the called method. The predecessor relation used in the algorithm is defined by the control flow graph of each method. The predecessors of the entry point of a method are all call sites for that method.

1. Initialize the sets of acquired locks as follows:
 - (a) For the `main()` method that is the application entry point: empty set.
 - (b) For each `run()` method that is the entry point of a thread: empty set.
 - (c) For all other methods: universe of all possible locks.
2. Repeat the following while there is a change in the set of acquired locks for any program point:
 - \forall statement S in the program code, with set of acquired locks $LS(S)$,
 $LS(S) \cap = \bigcap_{p \in \text{Predecessor}(S)} (\text{set of acquired locks for } p)$
 - If S is the first statement in a synchronized block that uses lock X ,
 add X to $LS(S)$.
 - If S is the first statement after a synchronized block that uses lock X ,
 remove X from $LS(S)$.
3. For each statement, add the “container” lock to its set of acquired locks if the statement only accesses fields of objects included in its set of acquired locks.

Figure 3.10: Algorithm to Compute Sets of Acquired Locks

Note that in our implementation, we conservatively initialize the sets of acquired locks at the entry point of a thread to be empty. We can improve the accuracy of locking analysis if we combine it with thread structure analysis. In this case, we would initialize

only the set of acquired locks at the program entry point to be empty, and initialize all other sets to be the universe of all possible locks. Then, in the algorithm in Figure 3.10, we would consider a thread `start()` point `S` to be the predecessor of entry points of threads that may be started at `S`.

Locks are regular Java objects. So we need pointer analysis to determine if a lock object is the same as one of the objects in a locking set. Since we do not have pointer analysis in the implementation, we take advantage of user-defined types to determine if two lock objects must be aliased. A lock object must be the same as another if both are of the same user-defined type, and we are guaranteed that a single instance of that type is instantiated. Thus, we consider a synchronized block only if its lock object is a write-once static field in Java.

We can circumvent the need for pointer analysis in the case when a synchronized block is used to protect accesses to fields of the same object whose associated lock is used by the synchronized block. We identify the lock in these cases by a special type that we call “**container**” lock. In the example code below, we include the “container” lock in the set of acquired locks for statements `X2` and `Y2`.

Thread X	Thread Y
<code>X1:synchronized (ObjectOne) {</code>	<code>Y1:synchronized (ObjectTwo) {</code>
<code>X2: ObjectOne.fld = ...;</code>	<code>Y2: ObjectTwo.fld = ...;</code>
<code>X3:}</code>	<code>Y3 }</code>

We use locking information in our delay set analysis to check if some conflict edges may be eliminated when searching for the end-points of a potential cycle. As shown in Section 3.4.2, if two accesses occur in synchronized blocks in the code, and the lock object used for both these blocks must be the same, then we can ignore a conflict edge between the two accesses in our simplified delay set analysis.

Consider accesses in synchronized blocks that access a field of the same object whose associated lock is used to synchronize the block. It is safe to discard any conflict edge between two such accesses in our delay set analysis. To see why this is the case, consider a conflict edge between accesses X2 and Y2 in the example above. There will be no conflict edge between X2 and Y2 because:

- either `ObjectOne` and `ObjectTwo` both refer to the same object, so both accesses are synchronized using the same lock, and we can eliminate conflict edges between them based on the discussion in Section 3.4.2, or
- `ObjectOne` and `ObjectTwo` refer to different objects, which implies that the two accesses are to different memory locations and therefore they do not conflict.

The special “container” lock is used to capture this case. We include this lock in the locking set corresponding to each access that refers to a field of object L, and the lock associated with L is used to synchronize the block that the access is contained in.

Complexity

The worst case time complexity of our locking analysis is bounded by $O(n^2 * s)$, where n is the number of program points and s is the number of locks used in synchronized blocks in the program. Each iteration does a pass over the program source and takes $O(n)$ time. The analysis iterates until there is no change in the locking sets being computed at each of the n program points. A locking set can only change monotonically. It can initially contain s elements, and it may have all of these elements removed from it. In the worst case, a single element will be removed from a single set in each iteration. Thus, the analysis needs at most $O(n * s)$ iterations, and the overall complexity is $O(n^2 * s)$.

3.4.2 Using Lock Information in Delay Set Analysis

In our approximate delay set analysis, when we search for initial conflict edges that form the end-points of a potential cycle, we can ignore conflict edges that occur between two accesses that execute as part of synchronized blocks with the same lock object.

Consider two conflicting memory accesses, say A and B, that execute as part of synchronized blocks with the same lock object. The mutual exclusion property of synchronized blocks implies that one access must finish before the other can start. Suppose that A happens before B in an execution of the program. For this case, Figure 3.11 shows the two possible ways in which a cycle may exist in the access order graph for delay set analysis, considering that only cycles with at most two consecutive conflict edges contribute to the delay set [SS88]. In the figure, thick lines denote program edges, thin lines denote conflict edges, and dashed lines denote some path that exists in the access order graph for delay set analysis.

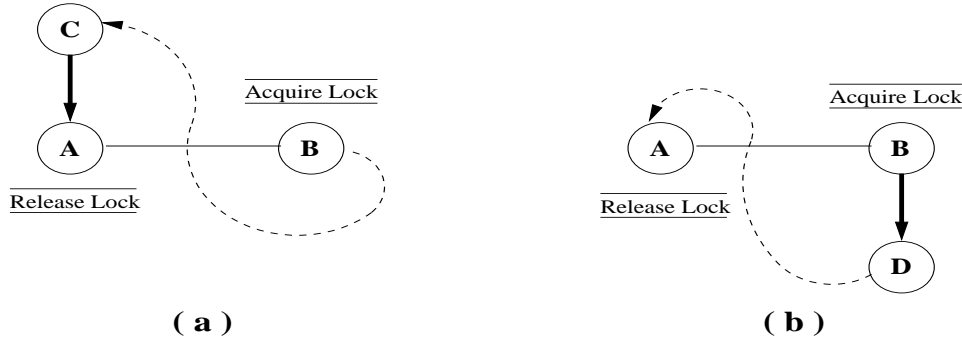


Figure 3.11: Possible Cycles for a Conflict Edge Between Synchronized Accesses

The cycle in Figure 3.11 (a) requires a program edge from some access C to A, and a path from B to C, i.e. the possibility that B happens before C. However, the second property listed above implies that since C is a memory access before the end of the synchronized block for A in the program source, C must happen before the lock for this synchronized block is released. From our assumption that A happens before B, we know B cannot happen before synchronization enforces ordering from the release of the lock after A to the acquire of the lock before B. So, it is not possible for B to happen before C.

Thus, when checking for a delay from C to A, we can safely ignore the cycle in Figure 3.11 (a), and the conflict edge between A and B can be eliminated.

The cycle in Figure 3.11 (b) requires a program edge from B to some access D, and a path from D to A. The last property listed above implies that since D is a memory access after the beginning of the synchronized block for B in the program source, it cannot happen before the lock for this synchronized block has been acquired. We know A happens before this lock is acquired, so it is not possible for D to happen before A. Thus, when checking for a delay from B to D, we can safely ignore the cycle in Figure 3.11 (b), and the conflict edge between A and B can be eliminated.

A similar reasoning applies to the case when B happens before A in a program execution. Thus, when checking if the program edge between two accesses X and Y is a delay edge, we can ignore conflict edges between X and another access if both of them have a common lock object in their locking sets. Likewise, we can ignore conflict edges between Y and another access if both of them have a common lock object in their locking sets.

3.5 Implementation

We efficiently implement dataflow analysis using a topological traversal of the strongly connected components of each method control flow graph. Also, we use a bitset implementation for our thread structure analysis algorithm. We have not implemented event-based synchronization analysis since it does not impact the benchmark programs we use. After the analysis, we preserve only the sets corresponding to method entry points. The sets for each program point can be computed intra-procedurally on demand.

Even though our synchronization analysis is inter-procedural and flow-sensitive, it is efficient in practice (Section 5.3). This is because the time it takes for the analysis to converge depends on the number of synchronization constructs used in the program. Typically, the number of points in the code where threads are spawned, or locks are

acquired or released, is small compared to the total number of program points. Also, for locking analysis, we need not perform a flow-sensitive analysis of methods that do not contain any synchronized blocks. For these methods, the locking information at each point in the method will be the same as the locking information at the method entry point. Likewise, for thread structure analysis, we need not perform a flow-sensitive analysis of methods that do not start or join any threads, either directly or indirectly through another call invoked when the method executes.

We use the results of synchronization analysis to perform delay set analysis. When searching for conflict edges, delay set analysis aggregates shared memory accesses within each method, instead of considering each individual access separately (Section 2.5.3). Also, nodes in the access order graph that correspond to method calls are treated as a collection of accesses performed either directly in the method called, or indirectly through a call invoked when the method executes. To efficiently handle method calls in delay checks, we compute for each method a summary of the synchronization information:

- over all read accesses of a given type that may execute when the method is called,
- over all write accesses of a given type that may execute when the method is called,
- and
- over all possible accesses that may execute when the method is called.

Thus, we store at most $2 * t + 1$ summaries for each method, where t is the number of types in the program.

For thread structure analysis, we can compute this summary by performing a union over the corresponding *StartsAfter*(P), *JoinsBefore*(P), and *Concurrent*(P) sets for each point P that is included in the summary information. Similarly, for event-based synchronization analysis, we can compute this summary by performing a union over the corresponding *SyncsAfter*(P) and *SyncsBefore*(P) sets for each point P that is included in the summary information. For locking analysis, we can compute this sum-

mary by intersecting the locking sets for each point P that is included in the summary information.

3.5.1 Exceptions

Exceptions in Java affect the flow of control in a program. Many instructions in Java are defined such that they either complete normally, or they throw an exception to indicate that something out of the ordinary happened. The program execution may proceed along different control flow paths depending on whether an instruction execution throws an exception or not. Programmers can define *exception handlers* that specify the code to be executed when a particular exception is thrown. Also, programmers can define new types of exceptions² and explicitly incorporate exception control flow in their programs due to these new exceptions.

Figure 3.12 shows an example where the execution of instructions I1 and I2 may throw some exception. If I1 executes normally, the next instruction to execute is I2. However, if an exception is thrown when I1 executes, the program continues execution at instruction I4 in the first exception handler block B5. Likewise, if an exception is thrown when I2 executes, the program continues execution at the second exception handler block B6.

Traditionally, a control flow graph is constructed so that the branching flow of control due to each instruction that may throw an exception terminates the basic block that the instruction is contained in. A majority of instructions in Java codes may throw exceptions. Due to this, there may be a large number of basic blocks, and this may increase the time for dataflow analysis in the compiler. Also, some optimization opportunities may be missed due to few instructions within each individual basic block. To alleviate these problems, the Jikes Research Virtual Machine compiler builds a *factored control flow graph* [CGHS99]. In a factored control flow graph, a basic block is not terminated

²Note that exceptions are encapsulated as Java objects, and each exception has a specific type associated with it.

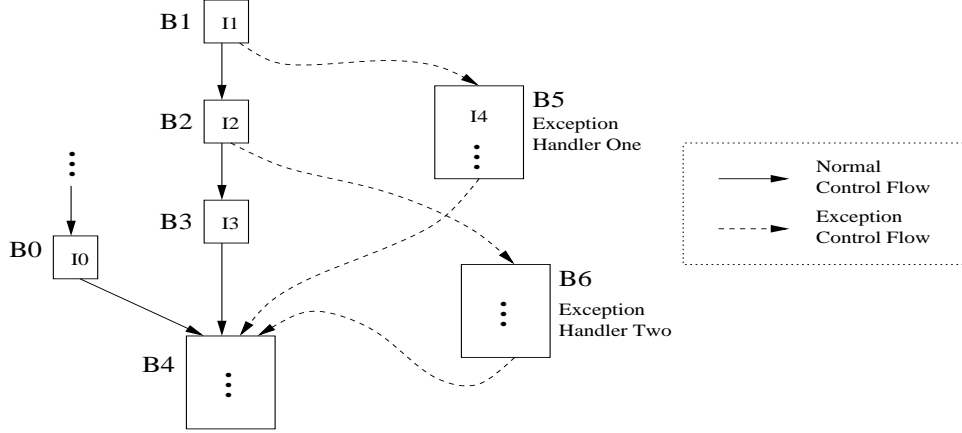


Figure 3.12: Flow Graph With Normal Control Flow and Exception Control Flow

at an instruction if exceptions are the only reason for a branch in control flow after the instruction executes. An edge exists between two basic blocks in a factored control flow graph if an exception due to any instruction in the first block may transfer flow of control to the second block. Figure 3.13 shows the factored control flow graph for the example in Figure 3.12.

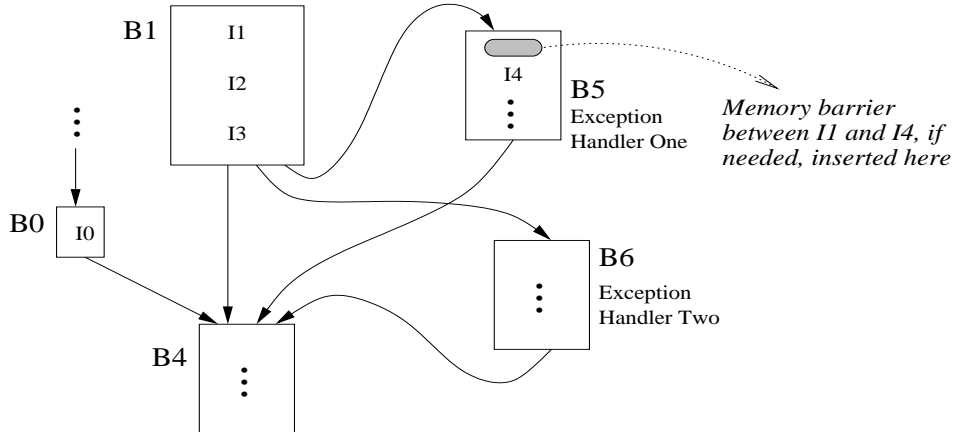


Figure 3.13: Factored Control Flow Graph in the Jikes RVM

For our synchronization analysis, we use dataflow analysis that traverses the control flow graph in a forward direction. The information propagated to exception handlers may be different for a traditional control flow graph and a factored control flow graph. In Figure 3.13, the information propagated to B5 and B6 is the information at instruction I3, instead of the information at I1 and I2 respectively. Only synchronization instruc-

tions change the sets that are propagated within a basic block. Thus, we can ensure correct analysis using a factored control flow graph if each synchronization instruction (`start()`, `join()`, `wait()`, `notify()`, and `notifyAll()` calls, and locks and unlocks corresponding to synchronized blocks) is the only instruction in its basic block. Since programs typically contain few synchronization instructions relative to the total number of program instructions, this change does not significantly impact the performance of program analysis.

Our delay set analysis considers each edge between an instruction and its successors in the control flow graph to be a program edge (Section 2.5.2). Successors of an instruction *I* contained in a basic block *B* include all instructions in basic blocks that are successors of *B*, as well as instructions within *B* that follow *I* in sequential order. In the example of Figure 3.13, delay edges that we test (among the instructions labelled in the figure) include edges from *I1* to *I2*, *I1* to *I3*, *I1* to *I4*, *I2* to *I3*, *I2* to *I4*, and *I3* to *I4*. Comparing this with Figure 3.12, we conservatively test the edges from *I2* to *I4*, and *I3* to *I4*. However, all successor edges determined from the flow graph in Figure 3.12 are also determined from the flow graph in Figure 3.13.

All paths in the original graph are contained in the factored control flow graph. This means that if our simplified delay set analysis determines an edge to be a delay edge in the original control flow graph, it will also determine that edge to be a delay edge in the corresponding factored control flow graph.

The design of a factored control flow graph ensures successors of an instruction include successors due to exception control flow. Thus, our delay set analysis correctly handles exception semantics in Java.

3.5.2 Space Complexity

As previously described, we store the results of synchronization analysis in terms of summaries for each method in the program. We have at most $2 * t + 1$ summaries for

each method, where t is the number of object types in the program. Each instance of synchronization information is proportional to the total number of synchronization constructs in the program: 3 bits for each thread `start()` call, 2 bits for each `wait()`, `notify()`, and `notifyAll()`, and 1 bit for each write-once static lock object associated with a synchronized block. Suppose the program contains s of these synchronization constructs. Then the space required to store the results of the analysis is $O(m * t * s)$, where m is the number of methods in the program.

During delay set analysis, we check program edges within each method to determine if they are delay edges. For this, we perform an intra-procedural flow-sensitive analysis to compute the synchronization information at each point in the method. This analysis considers one method at a time, and requires space proportional to $i * s$ where i is the maximum number of program points within any single method in the program.

Thus, the total space requirement for our analysis is $O((m * t * s) + (i * s))$.

3.6 Related Work

Synchronization enforces orders between different events in a program execution. In [NM90], Netzer and Miller prove that in general it is NP-hard to determine if two events in a program may be ordered, and it is co-NP-hard to determine if two events must be ordered. Past work has considered synchronization analysis in the context of race detection, deadlock detection, and synchronization removal. Early work in [YT88, HM91] describes how to perform an exhaustive search of all possible states that may occur in a parallel execution. This information is represented as *concurrency history graphs*, and takes exponential time to compute. Other research [EGP89, KY96, RM94] has focused on constructing a *task graph* for control flow amongst parallel threads of execution, and computing dominance relations on this graph to determine orders. This takes polynomial time, but is not as precise as the dataflow analysis methods described in [CS89, DS91], which also take polynomial time.

Callahan and Subhlok [CS89] give dataflow equations to compute synchronization orders, but their method cannot be generalized to handle recursion. Duesterwald and Soffa [DS91] use dataflow analysis to compute synchronization orders for a program that may use recursion, and whose code comprises of multiple methods. Their approach determines when one event must happen before and/or after another event. It cannot determine when two events may/must happen concurrently, and it is unable to differentiate between contexts specific to different call sites for the same method. Masticola and Ryder [MR93] build on previous work and describe a framework to determine events that may happen in parallel by successively refining the analysis results using various techniques. Recent work [NAC99] applies dataflow analysis to compute events that may happen in parallel for programs that use synchronization constructs defined in the Java language, but it does not attempt to compute any happens-before or happens-after orders. The synchronization analysis we present in this work also uses dataflow analysis. It computes happens-before, happens-after, and happens-in-parallel relations. It is more precise than previous methods based on dataflow analysis, and can distinguish information at different call sites for the same method.

There is substantial research on dynamic data race detection that attempts to determine orders based on a particular program execution. This work is related to our synchronization analysis, but the most advanced techniques [NG92, CLL⁺02] take advantage of dynamic properties of a particular execution. These techniques cannot be applied to a software implementation of memory consistency that must determine the delay set to honor the memory consistency model for *any* execution of the program.

In [KY96], Krishnamurthy and Yelick perform delay set analysis and incorporate synchronization information to improve the precision of their results. Their work aims to optimize the communication backend when compiling Titanium programs [YSP⁺98]. They target SPMD programs that do not include synchronization due to thread spawning or termination. Also, their use of locking synchronization to refine the access order graph differs from our approach.

Chapter 4

Compiler Framework

4.1 Java Terminology

Our analysis algorithms target the Java programming language. We explain some terms that refer to semantics of Java, and are used in the description of our algorithms.

Java programs manipulate data values. Each value in Java has a specific *type* associated with it. There are two basic types in Java: *primitive* and *reference*. A primitive type denotes an integer, floating point, or boolean value, as defined in the Java Language Specification[GJS96]. The value of a reference type denotes a collection of memory locations. A reference type may be any array, class, or interface declared in the code.

An *object* is an aggregation of primitive and/or reference values in memory that are collectively identified by a single reference value. A *field* is an identifier for a specific member of an object. A *class* or *interface* defines an aggregation by specifying fields, the types of values stored in fields, and methods to manipulate values in a program. An object is an *instance* of a class C, and the reference that identifies the object is said to be of type C. We say *user-defined types* are classes or interfaces that are defined in the user program.

A *memory access* is any instruction that reads a value from, or writes a value to, a location in memory. A single memory access can only read or write one primitive value, or one reference value. A memory access has an associated reference value that determines what object is accessed, and an associated field that specifies what member within the object is accessed.

A *variable* is a container for a primitive or reference value. It is declared to be of a specific type. Java is *strongly typed*, which means that a container can only be assigned values that are of a type compatible with the type of the container. Section 4.3.2 defines when one type is compatible with another.

4.2 The Jikes Research Virtual Machine and Compiler

We implemented the algorithms described in Chapters 2 and 3 in the Jikes Research Virtual Machine (Jikes RVM) from IBM [AAB⁺00]. The Jikes RVM is an open-source virtual machine that executes bytecodes contained in Java classes. It uses a dynamic just-in-time compiler to generate machine code for each Java method before it is executed. Except for some basic functionality, the source code for the Jikes RVM is written in Java. This enables the compiler to seamlessly integrate the virtual machine code and the application code when performing certain optimizations, such as code inlining.

The Jikes RVM has the ability to perform *online adaptive feedback-directed optimization* [AFG⁺00]. The virtual machine includes a monitor that profiles code executions to determine hot methods that take up a significant percentage of execution time. The monitor provides this information to the optimizing compiler. The compiler can then optimize hot methods in the code to improve performance. Mechanisms used to implement this adaptive optimization are method invalidation and re-compilation. Multiple instances of machine code for the same Java method may exist at the same time, and

outdated versions of code can be marked as invalid to indicate that they must no longer be used. We use this facility in our compiler system to make our analysis effective when dealing with partial program evaluation (Section 4.3.6).

Since compilation occurs dynamically in the Jikes RVM, compile time counts as part of an application's execution time. So it is essential that our analysis be fast. However, there are significant advantages of using dynamic compilation and a runtime system:

- We can dynamically detect when an application spawns a new thread during execution. Since our analysis is only applied to multithreaded programs, we need not analyze code until a thread is spawned. Therefore we can execute single-threaded applications with no overhead.
- Suppose a particular execution instance of an application never loads some class nor invokes some method present in the code for the application. A dynamic compiler can exclude such classes and methods when performing analysis, but a static compiler, in general, cannot. Thus, the analysis performed by a dynamic compiler can be more precise.




Figure 4.1: Performance of the Jikes RVM and Commercial JVMs

Figure 4.1 shows the performance of the Jikes RVM with respect to two commercial Java Virtual Machines: the Sun JDK 1.4.2 and the IBM JDK2-1.4.2. All measurements were obtained on a Dell PowerEdge 6600 SMP with 4 Intel hyperthreaded 1.5GHz Xeon processors running Linux. For each of the benchmark programs described in Section 5.1, the figure shows two ratios: the Sun JDK execution time over the Jikes RVM execution time, and the IBM JDK execution time over the Jikes RVM execution time¹. We observe that on average the performance of the Jikes RVM is comparable with that of commercial Java virtual machines. Thus, we believe our results are comparable to those that would have been obtained using one of the commercial Java Virtual Machines.

4.3 Our Memory-model Aware Compiler

In Section 1.2.2, Figure 1.4, we gave an overview of the components in our compiler system. We implement our system using the Jikes RVM, and extend its optimizing compiler with our analyses and transformations. Currently the programming language memory model is fixed to be sequential consistency, but in the long term this is expected to be a programmable feature and the system will automatically adapt to different memory consistency models.

The delay set analysis and synchronization analysis components of Figure 1.4 were discussed in Chapter 2 and 3 respectively. In this section, we discuss the thread escape analysis, alias analysis, and memory barrier insertion components of our compiler implementation. We describe the code re-ordering optimizations in the Jikes RVM that need to be augmented for sequential consistency. We also describe how our algorithms adapt to perform partial program analysis in the presence of dynamic classloading.

¹SPEC JBB 2000 [Cor00] is a server benchmark and its performance is measured in terms of throughput, i.e. number of transactions completed per second. Thus, the ratios for SPEC JBB 2000 are computed as the Jikes RVM throughput over the Sun JDK throughput, and the Jikes RVM throughput over the IBM JDK throughput.

4.3.1 Thread Escape Analysis

Definition 4.3.1 *A reference to a memory location **escapes** if that location is accessible by multiple threads in the program.*

Thread escape analysis identifies shared memory accesses in the program that may use references that escape. In our system, we use a flow-insensitive, partially context-sensitive, iterative algorithm for thread escape analysis. It differs from previous algorithms in that it considers individual fields of Java `Thread` objects, and checks if they do indeed escape. Also, it makes use of user-defined types to speed up convergence of the iterative analysis.

We tested our implementation of sequential consistency using different escape analysis algorithms to determine the effect escape analysis has on delay set analysis (Section 5.2.2). The different escape analysis algorithms varied in the accuracy of their analysis results. Not surprisingly, we found that delay set analysis is very sensitive to the precision of escape analysis. For some of the benchmark programs described in Section 5.1, we observed slowdowns in execution time that were an order of magnitude greater when using an escape analysis that is less precise than the one described in Section 4.3.1.2.

4.3.1.1 Algorithm

In Java, there are three causes for a reference to escape a thread:

1. A reference to a static field may be accessed by multiple threads in the program, causing it to escape.
2. When a new thread T1 is created by a thread T2, the reference to T1 may be accessed by T1 itself as well as its creator T2. So, the reference to T1 escapes.
3. If a reference R1 escapes, then any reference R2 read from or written into a memory location accessed through reference R1 also escapes.

When a type may escape, all accesses to all objects that are instances of that type may escape. Similarly, when a field may escape, all accesses to that field may escape.

When a new thread object T1 is created, its creator thread may use the reference to T1 to access fields of T1. If these fields are of a reference type, the values they contain escape because they are accessible both by the creator thread, and by the execution of T1. The second and third conditions enumerated above together identify that these references escape. However, the second condition is conservative because it causes all references contained in fields of T1 to escape. In reality, only some fields of T1 will be accessed by the creator. That the reference to T1 is available to its creator thread will cause only those references to escape that are read from or written into fields of T1 by its creator thread.

We use this in our analysis to refine the second condition enumerated above. We do not assume that the reference to a new thread T1 escapes by default. Instead, each reference that is a field of T1 that may be accessed in the creator thread escapes. To determine these fields, we need to analyze all code that may be executed by the creator thread. In our implementation, we use a conservative test that only examines the code for two methods: the method that contains the thread `start()` call for T1, say M, and the constructor for T1, say C. This test determines if the reference to T1 is available for use only within these two methods:

1. T1 is newly allocated in M, and
2. the reference to T1 is not stored to memory in M or C, and
3. the reference to T1 is not passed as a parameter for a method called from M or C, and
4. the reference to T1 is not returned from M.

If so, then only fields of T1 accessed in M or C escape. Otherwise, our analysis conservatively assumes that the thread type for T1 escapes, and that *all* fields of T1 also escape.

Thus, the root causes for a reference to escape in our analysis are:

1. Any reference to a static field escapes.
2. For a new thread T1 that is spawned by a thread T2, the thread type that corresponds to T1 escapes, unless our conservative intra-procedural analysis determines that the reference to T1 is not available for use outside the method that spawns the thread T1, and the constructor for T1.
3. For a thread type T1 that does *not* escape according to the above condition, a field F of T1 escapes if:
 - F is of reference type, and
 - a statement S, in the method that spawns T1 or in the constructor method for T1, accesses a value in F, except when:
 - (a) the value accessed by S is `null`, or
 - (b) S is a write access that stores a reference to an object newly allocated in the constructor for T1, and this reference is not found to be escaping.

There are four categories of statements in our intermediate representation that cause a reference to escape because of another escaping reference: load or store to an object field in memory, read or write access to an array element, statements that use register operands to perform a unary or binary operation, and method calls. We describe each of these below. In our discussion, F identifies a field type, and R, R_0, R_1, \dots identify temporary registers in the intermediate code.

Loads and Stores of an Object Field A load of an object field, $R_2 = R_1.F$, loads into R_2 the value from field F in the object referenced by R_1 . A store of an object field,

$R_1.F = R_2$ stores the value from R_2 into the field F in the object referenced by R_1 . In both cases, escape properties are affected only if the value loaded or stored is of reference type. If R_1 escapes, then everything accessible through it also escapes, and so R_2 escapes. If F is a field of a thread type that escapes (the second root cause enumerated above), then we conservatively assume that both R_1 and R_2 escape. If F is an escaping field type (the third root cause enumerated above), then the value read from or written into F may be accessible by multiple threads, and so R_2 escapes.

Array Element Accesses A read access of an array element, $R_2 = R_1[i]$, reads into R_2 the value that is at index i of the array referenced by R_1 . A write access of an array element, $R_1[i] = R_2$, stores the value from R_2 into the element at index i in the array referenced by R_1 . We treat array accesses very conservatively since we do not use subscript analysis to determine when two array accesses may refer to the same location in memory. If an element of an array escapes, we assume the entire array escapes. Thus, if R_2 escapes, then R_1 escapes. Also, if R_1 escapes, then each element in the array also becomes accessible by multiple threads, and so R_2 escapes.

Register-to-register Operations A unary operation, $R_1 \text{ op } = R_2$, performs some operation on the value in R_2 , and stores the result in R_1 . If R_2 escapes, we conservatively assume R_1 also escapes. Note that a move instruction also counts as a unary operation that simply transfers the value from R_2 to R_1 . A binary operation, $R_1 = R_2 \text{ op } R_3$, performs some operation on the values in R_2 and R_3 and stores the result in R_1 . If R_2 or R_3 escape, we conservatively assume R_1 also escapes.

Method Calls For each method M we compute two sets during the iterative analysis: $EscapeParamsOnEntry(M)$ and $EscapeParamsOnExit(M)$.

Definition 4.3.2 For a method M , $i \in \text{EscapeParamsOnEntry}(M)$ if some call to M in the program may be invoked with an escaping reference as the argument for the i th parameter.

Definition 4.3.3 For a method M , $i \in \mathbf{EscapeParamsOnExit}(M)$ if execution of M may cause the reference contained in the i th parameter to escape.

Consider a method call, $R_0 = \text{foo}(R_1 \dots R_n)$, where foo is the method invoked. foo takes n parameters whose values are passed in registers R_1 through R_n . It returns a value that is stored in register R_0 . This return value can be treated as an out parameter that is available only on return from the method. Thus, we consider it to be the zero'th parameter when computing the set $\text{EscapeParamsOnExit}(M)$.

For parameter i of method M , $0 \leq i \leq n$, if $i \in \text{EscapeParamsOnExit}(M)$, then R_i escapes. Conversely, for $1 \leq i \leq n$ if R_i escapes, then we add i to the set $\text{EscapeParamsOnEntry}(M)$.

To compute $\text{EscapeParamsOnExit}(M)$ for a method M , we determine all registers in M that escape when initially no parameter passed into the method escapes. If register R_i corresponding to parameter i is found to be escaping, then we know the actions of the method cause the reference passed in parameter i to escape. So we add i to $\text{EscapeParamsOnExit}(M)$. It is also possible that a parameter that does not initially escape on method entry later escapes during method execution only because of an escaping reference that is passed in through some other parameter. Due to this, we individually consider the effect of each parameter that may be escaping on method entry. In each case, we determine what other parameters may escape due to actions of the method, and add them to the set $\text{EscapeParamsOnExit}(M)$. We only need to consider the effects of parameters that escape on method entry one at a time, since the sets of escaping references determined in each case are independent of one another.

Figure 4.2 illustrates how our analysis is partially context-sensitive. Method M takes two parameters u and v , and the actions of M are such that it causes v to escape only when the reference passed in u escapes. Method M is called from three points in the program:

1. Method A calls M with neither parameter initially escaping.

<pre> A () { /* X, Y do not escape */ call M (X, Y); /* Y escapes */ } </pre>	<pre> B () { /* X escapes, Y does not*/ call M (X, Y); /* X, Y escape */ } </pre>	<pre> C () { /* X, Y escape */ call M (X, Y); /* X, Y escape */ } </pre>
---	---	--


```

void M (u, v) {
    /* M causes v to escape only
       if u escapes on entry */
}

```

```

EscapeParamsOnEntry (M) = {1, 2}
EscapeParamsOnExit  (M) = {2}

```

Figure 4.2: Different Calling Contexts of Methods for Escape Analysis

2. Method *B* calls *M* with *u* initially escaping but *v* not escaping.
3. Method *C* calls *M* with both parameters initially escaping.

The figure also shows sets $EscapeParamsOnEntry(M)$ and $EscapeParamsOnExit(M)$ that our analysis computes. Based on this, we correctly determine that *Y* escapes after the call to *M* in *B* returns, and that *X* does not escape after the call to *M* in *A* returns. However, we conservatively determine that *Y* escapes after the call to *M* in *A* returns.

Figure 4.3 gives an algorithm to perform our iterative thread escape analysis. It uses the rules in Figure 4.4 that describe how to process statements in a method during analysis.

At the end of the iterative analysis, we save sets $EscapeParamsOnEntry(M)$ and $EscapeParamsOnExit(M)$, and discard escape information for registers in the intermediate code. This information can later be computed intra-procedurally on demand, using the saved sets.

4.3.1.2 Optimization for Fast Analysis

We take advantage of user-defined types in Java to speed up the convergence of our iterative analysis. We conservatively assume that if an access to an object field escapes, then all accesses to the same field also escape. To incorporate this change, we need only

- I Initialize the following sets to be empty:

 - 1) *EscapingThreads*: set of all thread types found to be escaping.
 - 2) *EscapingThreadFields*: set of all fields that are part of thread objects and that are found to be escaping.

II For each method M , initialize the sets $EscapeParamsOnEntry(M)$ and $EscapeParamsOnExit(M)$ to be empty.

III Repeat the following until there is no change in the escape property of any parameter, thread type, or field:

For each method M in the program with parameters contained in registers $R_1 \dots R_n$, and the return value contained in register R_0 :

 - 1) Process M to determine escaping references as described in Fig 4.4.
 - 2) $\forall i, 0 \leq i \leq n$, if R_i escapes \Rightarrow add i to $EscapeParamsOnExit(M)$.
 - 3) $\forall i \in EscapeParamsOnEntry(M)$:
 - a) Set R_i to be escaping.
 - b) Process M to determine escaping references as in Figure 4.4.
 - c) $\forall j, 0 \leq j \leq n$,
if $i \neq j$ and R_j escapes \Rightarrow add j to $EscapeParamsOnExit(M)$.
 - d) Reset all registers in M to be not escaping.

Figure 4.3: Escape Analysis Algorithm

Repeat the following until there is no change in the escape property of any register, parameter, thread type, or field:

\forall statement S in method M :

- 1) If S is an access to a static field F : $R = F$ or $F = R$,
 R is of reference type $\Rightarrow R$ escapes.
- 2) If S is a thread start call invoked on object R of thread type T : $R.start()$,
 R escapes \Rightarrow add T to set *EscapingThreads*.
 R accessed in a method other than M or the constructor method for R
 $\Rightarrow R$ escapes, and add T to set *EscapingThreads*.
- 3) If S accesses a field in a thread of reference type: $R_1.F = R_2$ or $R_2 = R_1.F$,
Thread type of $R_1 \in \text{EscapingThreads}$
 \Rightarrow add F to set *EscapingThreadFields*.
 S is a read in constructor for R_1 or in method that starts thread R_1
 \Rightarrow add F to set *EscapingThreadFields*.
 S is a write in constructor for R_1 or in method that starts thread R_1 , and
 R_2 is not null, and R_2 is not a thread-local object newly allocated in the
constructor for R_1
 \Rightarrow add F to set *EscapingThreadFields*.
- 4) If S is an access to field F of object type T : $R_1.F = R_2$ or $R_2 = R_1.F$,
 R_2 is a reference and $T \in \text{EscapingThreads} \Rightarrow R_1$ and R_2 escape.
 R_2 is a reference and $F \in \text{EscapingThreadFields} \Rightarrow R_2$ escapes.
 R_2 is a reference and R_1 escapes $\Rightarrow R_2$ escapes.
- 5) If S is an access to an element of array R_1 : $R_1[] = R_2$ or $R_2 = R_1[]$,
 R_1 escapes $\Rightarrow R_2$ escapes.
 R_2 escapes $\Rightarrow R_1$ escapes.
- 6) If S is a unary or binary operation: $R_1 \text{ op } = R_2$ or $R_1 = R_2 \text{ op } R_3$,
 R_2 escapes or R_3 escapes $\Rightarrow R_1$ escapes.
- 7) If S is a method call: $R_0 = \text{foo}(R_1 \dots R_n)$,
For each i , $0 \leq i \leq n$, $i \in \text{EscapeParamsOnExit}(\text{foo}) \Rightarrow R_i$ escapes.
For each i , $1 \leq i \leq n$, R_i escapes \Rightarrow add i to *EscapeParamsOnEntry(foo)*.

Figure 4.4: Effect of Statements in Escape Analysis Algorithm

change the rules to process statements that access object fields. The revised rules are shown in Figure 4.5.

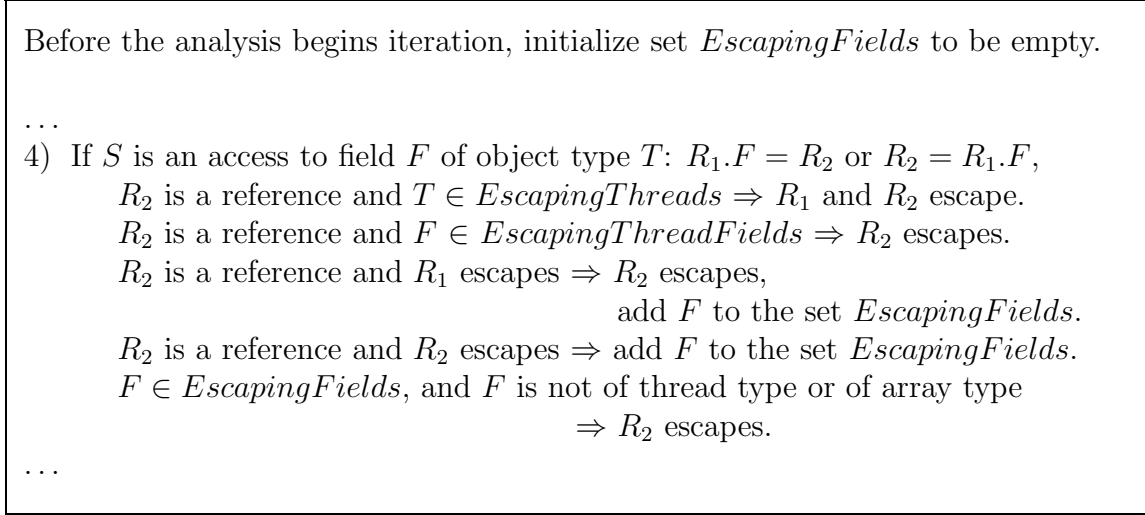


Figure 4.5: Optimizing Escape Analysis for Fast Convergence

During the analysis, a set *EscapingFields* is maintained. A field $F \in \text{EscapingFields}$ if there is a statement in the code that loads an escaping reference value from field F of some object, or stores an escaping reference value into field F of some object. The last rule in Figure 4.5 is not required for correctness; it serves only to accelerate the convergence of the analysis. It conservatively assumes that the register R_2 involved in an object field access escapes if the access is to a field $F \in \text{EscapingFields}$. Note that we do not apply this conservative rule to fields that are of array type or thread type. This is because our analysis is very conservative for array element accesses, and accesses to objects of thread type are one of the main causes for references to initially escape.

Applying this optimization speeds up escape analysis for the largest two benchmarks described in Section 5.1. Escape analysis for the SPEC JBB 2000 benchmark [Cor00] takes 60% less time, and for *mtrt* in the SPEC JVM 98 benchmark suite [Cor98] takes 35% less time.

4.3.2 Alias Analysis

Given two references, alias analysis determines if they may be aliased, i.e. if they may refer to the same memory location. This information is used to test for conflict edges during delay set analysis. A conflict edge can only exist between two accesses that may be aliased. Precision of alias analysis affects the number of conflict edges found, and so it affects the accuracy of delay set analysis. Ideally, we want to perform pointer analysis and trace when specific accesses in the program may refer to the same memory location. However, precise pointer analysis is expensive. Since compilation occurs during runtime in our dynamic compiler, we choose not to use pointer analysis in our implementation. Instead, we take advantage of strong typing in Java, and use a type-based alias analysis.

Definition 4.3.4 *We say type A is **compatible** with type B in Java if:*

1. *A is the same type as B , or*
2. *A is a subclass of B , or*
3. *A is a class that implements interface B .*

Each variable in the program is declared to be of a specific type. A reference variable can only refer to objects that are of a type compatible with the declared type of the reference. For two references of type A and B , if A is compatible with B , or B is compatible with A , then our alias analysis assumes they may refer to the same memory location, and so they may be aliased.

If the class corresponding to some type has not yet been loaded, then the Java virtual machine has no knowledge of its superclass, subclasses, or interfaces implemented by it. In the absence of such information, our alias analysis optimistically assumes that the unknown type is not compatible with any other type in the application². Later, if a new class is loaded, we need to re-analyze the code, as described in Section 4.3.6.

²The exception is that all classes are by default subclasses of the `Object` type in Java.

For accesses to array elements, we need array dependence analysis to process subscript expressions, and determine if two accesses may refer to elements at the same array index. In our implementation, we do not perform array dependence analysis. Instead, we conservatively assume that all accesses to elements of a given array may access the same element.

4.3.3 Inhibiting Code Transformations

Subset correctness is the correctness criterion for our memory-model aware compiler. It requires that the set of outcomes possible from execution of a transformed program be a subset of outcomes possible from the original program execution. Transformations that preserve orders determined by delay set analysis (and normal intra-thread dependences) preserve subset correctness [LPM99].

To preserve orders determined by delay set analysis, we need to prevent transformations in the compiler that re-order or eliminate accesses related by delay edges. Specifically, for each delay edge from access A to B , the order of A and B must not be changed, and neither A nor B must be eliminated.

In the Jikes RVM optimizing compiler, there are four optimization phases that may re-order or eliminate shared memory accesses:

1. loop invariant code motion,
2. loop unrolling,
3. common subexpression elimination, and
4. redundant load elimination.

Figure 4.6 shows an example of a loop invariant code motion transformation that is valid in a single-threaded application. However, for multithreaded execution, this optimization cannot be applied when there is another thread that writes the value 1

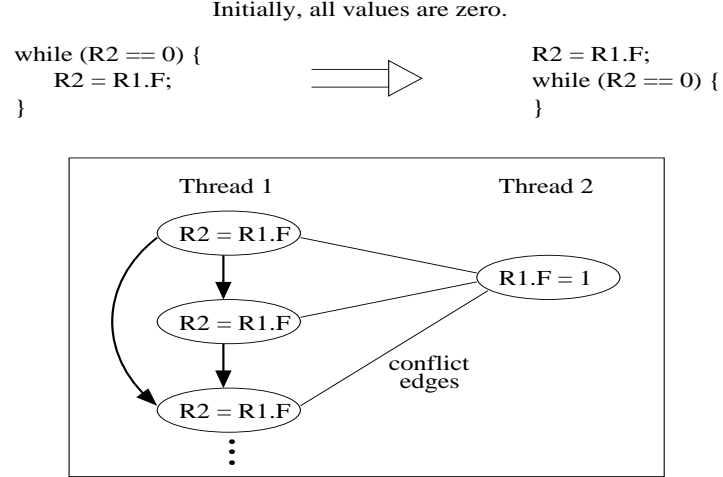


Figure 4.6: Loop Invariant Code Motion Example

into the memory location referred by $R1.F$. In this case, a possible outcome for the transformed code is that the loop repeats forever. However, this is not possible in the original code as it must eventually observe the value written by the other thread. The access order graph for sequential consistency has a delay edge from access $R1.F$ in one loop iteration to access $R1.F$ in subsequent loop iterations. The transformation amounts to eliminating all accesses to $R1.F$ except the one in the first loop iteration. Thus, the transformation does not preserve delay edges intact, and the transformed program violates subset correctness.

Figure 4.7 shows an example of loop unrolling that is valid in a single-threaded application when $R2$ is a loop constant, and the number of loop iterations can be symbolically computed. However, in a multithreaded context, $R2$ is not a constant if concurrent threads write to the memory location referred by $R3.F$. In this case, the number of loop iterations executed by the transformed code may differ, thus violating subset correctness. The access order graph for sequential consistency has delay edges between access $R3.F$ in one iteration and access $R3.F$ in subsequent loop iterations in the original code. The transformation either results in some of the accesses to $R3.F$ being eliminated or some

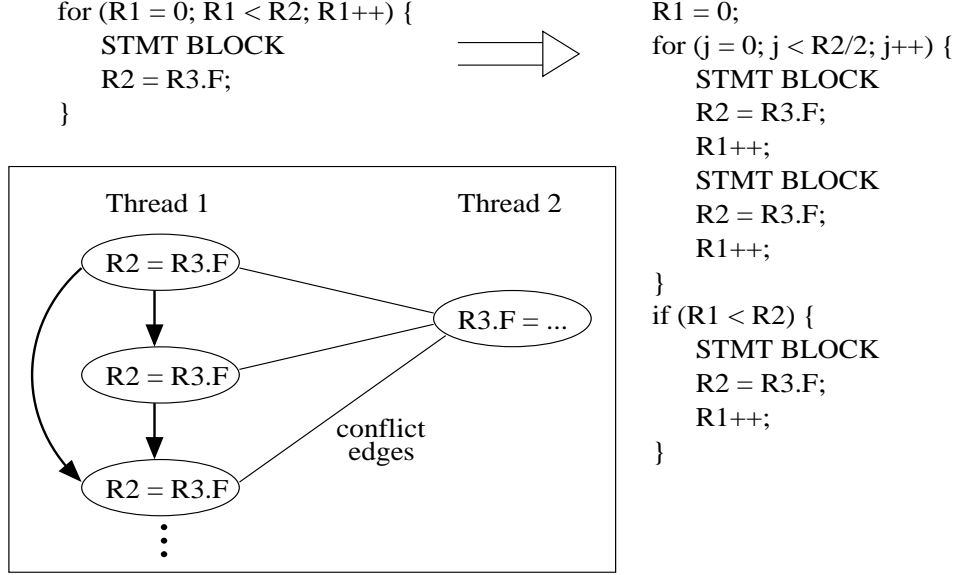


Figure 4.7: Loop Unrolling Example

new accesses to $R3.F$ being introduced³. The transformed code does not preserve delay edges intact, and the transformation is invalid.




Figure 4.8: Common Subexpression and Load Elimination Example

Figure 4.8 shows an example of common subexpression elimination and load elimination that is valid in a single-threaded application. However, this transformation cannot be

³Note that transformations typically do not introduce extra shared memory accesses when optimizing code. However, in general, it is not correct to insert extra shared memory accesses.

applied when another thread writes to the memory locations accessed by $R1.F$ and $R4.F$, and sequential consistency is assumed. The access order graph in the figure illustrates such an execution. For the original code, if $R5 = 2$, then the value for $R6$ must be 1 and $R7$ must be 3. However, in the transformed code, it is possible to observe $R5 = 2$ and $R7 = 2$, which is not a possible behavior in the original code. Thus, subset correctness does not hold. The transformation eliminates the second load of $R1.F$ and this violates the delay edges shown in the access order graph.

In our implementation, we observe no significant impact on performance due to inhibiting code re-ordering and code elimination transformations. This is despite the fact that we use a conservative test to determine when to inhibit transformations. To check if accesses A and B may be re-ordered, we ensure that neither A nor B is thread-escaping. Likewise, to check if an access A may be eliminated, we ensure A is not thread-escaping. We do not test for a delay edge from A to B to preclude re-ordering them, and we do not search for some delay edge that involves A to determine if A can be eliminated. This is because when a method is compiled, we test for delay edges after optimizations have been applied and we are ready to generate machine code for it. We could test for delay edges earlier in the compilation process and store the results. Then, as the method code is transformed, we will need to either update or re-compute the set of delay edges found. However, by using the conservative test based only on results of thread escape analysis, we save the overhead of doing this without sacrificing performance in practice.

4.3.4 Memory Barrier Insertion

Memory barrier insertion occurs when machine code is being generated for a method. It ensures that the access orders required by the programming language memory model are not violated by the processor executing the machine code.

On the IBM Power3 architecture, memory barrier insertion inserts *sync* instructions between access pairs whose order needs to be enforced. A *sync* instruction ensures that all

data memory accesses⁴ issued by the processor before the *sync* instruction are complete and available to all processors before any data memory access after the *sync* is issued. On the Intel Xeon architecture, there are three types of memory barrier instructions that are used to enforce access orders:

1. An *mfence* ensures all memory accesses that occur in code order before the *mfence* are complete before any access after it is issued.
2. An *lfence* ensures all load accesses that occur in code order before the *lfence* are complete before any load access after it is issued.
3. An *sfence* ensures all store accesses that occur in code order before the *sfence* are complete before any store access after it is issued.

Note that *lfences* and *sfences* cost less than *mfences* in terms of execution time.

Our system uses the local-optimized version of memory barrier insertion described in [FLM03]. After code optimization and transformation, the compiler determines the sequence of machine instructions to be generated for a method being compiled. Memory barrier insertion searches this sequence and constructs the set of all memory access instructions and method call instructions. For each pair of instructions (I, J) such that I and J are in this set, and I is a predecessor of J in the method control flow graph, it queries delay set analysis to determine if a delay edge exists from I to J . For each delay edge (I, J) , a memory barrier instruction is inserted on every path from I to J in the control flow graph: if I and J are in different basic blocks of the control flow graph, the memory barrier is inserted at the beginning of the basic block for J , else the memory barrier is inserted just before J .

Memory barrier insertion is optimized to coalesce consecutive memory barrier instructions in the code into a single instruction. In the example control flow graph in

⁴To order accesses to program instructions, the IBM Power3 architecture provides a separate *isync* instruction.

Figure 4.9, if there is a delay edge from A to C, and from B to C, then a single memory barrier before C enforces both these delay edges.




Figure 4.9: Example for Memory Barrier Insertion

Also, on the Intel Xeon system, the analysis is optimized to use cost-effective *lfences* and *sfences* instead of *mfences* wherever possible. Note that there is further scope for optimizing memory barrier insertion that is not currently part of our system implementation. In particular, we can take into account loop structure and branch probabilities in the control flow graph when deciding where to insert a memory barrier instruction.

4.3.5 Specialization of the Jikes RVM Code

The Jikes RVM is mostly written in Java, and the compiler can integrate the virtual machine (VM) code with the application code. Thus, some methods may execute both as part of the VM and the application. Ideally, for an implementation of sequential consistency, we want to analyze all methods to determine delay edges, including methods shared by the VM itself. Note that this would not be an issue if the Jikes RVM was written in C/C++, as are commercial Java virtual machines. If our analysis includes VM methods, it results in high overhead:

- If VM methods are included in the analysis, their number tends to be much greater than the number of application methods. As a result, the analysis time increases, and the majority of it is spent analyzing VM code. Moreover, this is done each time an application executes.

- The VM code is large and complex, and it uses several concurrent thread types, as well as static variables. As a result, our analysis is not precise for VM code, and this conservativeness infiltrates the application code as well. Thus, overall accuracy and performance are affected.
- The VM and the application may share some library methods. The application may pass an escaping reference to the shared method while the VM does not. The method will need memory barriers in the context of the application, but not in the context of the VM. Thus, an application can potentially slow down the execution of the entire VM.

In our implementation, we want to ensure sequential consistency for *all* code executed by an application without having to analyze VM code each time an application runs. We identify VM methods that may execute as part of application code. To do this, we use manual inspection of the VM code that is guided by execution traces from our benchmark programs. Methods identified include:

- methods to clone an object,
- methods to efficiently copy arrays,
- methods to query the execution state of a thread, and
- methods in the implementation of the `java.lang.Throwable` class that access the state of an exception⁵.

For all such methods identified, we provide two specialized versions of the method in the system: one version has memory barriers conservatively inserted for each shared memory access in it, and the other version does not have any memory barriers inserted. The VM always uses the version with no memory barriers⁶. For the application, the compiler

⁵All exceptions in Java are derived from the type `java.lang.Throwable`.

⁶We assume the VM is well-synchronized.

decides which version to invoke at each call site depending on whether escaping references are passed as parameters to the method.

Thus, the number of memory barriers inserted is similar to the number that would be inserted if the VM was written in C/C++.

4.3.6 Dynamic Class Loading

Java allows dynamic classloading, so the entire code for a program may not be available when it begins execution. During our inter-procedural analysis, we may need to consider the effect of some methods whose code is not accessible when the analysis is performed. In such cases, we can be conservative and ensure that the analysis results are valid irrespective of what the method does. However, this will cause delay set analysis and synchronization analysis to be ineffective, since they will have to assume that the unknown method reads and writes all user-defined types that are accessible from it, and it starts threads of all types accessible from it. Instead, our analysis optimistically assumes that the unknown method does not spawn any new threads, and that it contains no shared memory accesses that conflict with an access in another method. Later, when code for the method becomes available, these properties must be verified to determine if the analysis results are valid. If an assumption is found to be wrong, we must perform the analysis again, invalidate previously compiled methods that may now need extra memory barriers inserted, and recompile all these methods to generate correct code.

4.3.6.1 Incremental Analysis

We modify the runtime system of the Jikes RVM to detect the first time the application spawns a new thread. Until a thread is spawned, we assume the application is single-threaded, and delay set analysis trivially finds no delay edges. When the runtime system is first invoked by the application to spawn a new thread, it passes control to the optimizing compiler. At that point, synchronization analysis is performed for the

first time, and the check for delay edges is no longer trivial. The analysis is performed inter-procedurally over code for all methods possibly executed by the application and available at the point of the analysis. We do not repeat the inter-procedural analysis if the just-in-time compiler later compiles a method whose code was available during the inter-procedural analysis phase.

Due to classloading, code for a method that was previously unavailable may become available. When the just-in-time compiler attempts to compile this method, our analysis must perform an inter-procedural phase to account for the effect of the new method. Instead of repeating the entire analysis from scratch, we design it to build upon previous results.

Since each inter-procedural phase is potentially expensive, we want to minimize the number of times it occurs. For this, we want to perform each analysis phase over all application code that is available. Therefore, we use aggressive classloading for class types seen in the application code that is already available to the analysis. We attempt to load such classes before they are actually needed by the execution, so we can analyze them as part of the current inter-procedural analysis phase. Java semantics allow this as long as no exception is thrown if the class being aggressively loaded does not exist, and if the class initializer is not executed till the point when the class is actually used.

Call Graph

For our analysis, we maintain an explicit call graph. A call site may invoke one of several possible methods, depending on the type of the object on which the call is invoked (Section 2.5.5). When a new method triggers our inter-procedural analysis, it is possible that a new class has been loaded and is used in the application. If a class C is newly loaded, we determine the set of all types T such that C is compatible with T (Section 4.3.2). For each call to a method (say M) invoked through a reference of type T , we check if class C defines a method (say N) that is a candidate for that call. If so, we update our call graph to include this new method N in the set of possible methods that the call to M resolves to.

Also, we set the synchronization information at the entry point of method N ⁷ to be the conservative summary of the synchronization information at the entry points of method M and method N . Section 3.5 describes how synchronization information is summarized over multiple points.

Thread Structure Analysis

We check the code for any new methods that have become available and look for thread `start()` calls and matching thread `join()` calls. If no such calls are found, the incremental analysis traverses the call graph rooted at each new method to propagate thread structure information across all methods called directly or indirectly from the new method. Note that we only need to iterate over individual statements of the new methods, since we can use the call graph along with the summarized synchronization information stored for each method that was previously analyzed, to reconstruct its inter-procedural effect.

If any thread `start()` call is identified in the new methods, then the iterative analysis outlined in Figure 3.6 is performed with respect to the newly identified `start()` calls. In this case, we need to iterate over individual statements of all the new methods, as well as any method that may invoke a newly identified `start()` or `join()` indirectly through another method that it calls. With reference to Figure 3.6, we iterate over statements of a method M if either *StartsInMethod*(M) or *JoinsInMethod*(M) contains one of the newly identified `start()` calls. For all other previously analyzed methods, we only need to propagate information from that method to all other methods that it may call. We maintain an explicit call graph that can be used for this purpose.

Locking Analysis

Incremental locking analysis traverses the call graph rooted at each new method to propagate locking information across all methods called directly or indirectly from the new method. With reference to Figure 3.10, we perform step 2 of the algorithm, iterating over individual statements only in the new methods and using the call graph to propa-

⁷Recall from Section 3.5 that we preserve the synchronization information for the entry point of each method.

gate information across entry points of methods previously analyzed. Also, step 3 of the algorithm in Figure 3.10 is performed for each new method.

Delay Set Analysis

For each new method whose code has become available since the last inter-procedural analysis phase, we compute synchronization information summaries over read and write accesses of a specific type in the method, and over all possible accesses that result from the method execution, as described in Section 3.5. The summary information over all possible accesses in a method M takes into account accesses in methods that may be called directly or indirectly by M . So, we need to re-compute this summary over all accesses for any method that may directly or indirectly call one of the new methods.

For a newly loaded class C , we determine the set of all types T such that C is compatible with T (Section 4.3.2). Then, for each method in the application, and each type T in the set, we update information on types accessed by the method (Section 3.5) as follows:

- If the method may read accesses of type C as well as type T , then we conservatively incorporate the summarized synchronization information for T into C .
- If the method may read accesses of type T but not of type C , we add an entry in the summarized synchronization information for read accesses to type C , and this entry has the same summary information as that for type T .
- If the method may write accesses of type C as well as type T , then we conservatively incorporate the summarized synchronization information for T into C .
- If the method may write accesses of type T but not of type C , we add an entry in the summarized synchronization information for write accesses to type C , and this entry has the same summary information as that for type T .

We choose to store all possible types in the method summary information instead of dynamically checking for compatible types each time the summary information is used.

This makes it easy to check for methods that must be invalidated because of a change in the analysis results (Section 4.3.6.2).

4.3.6.2 Method Invalidation

If the entire program source is not available the first time we perform our inter-procedural analysis, we make optimistic assumptions about parts of the code that are not available. Later, when more code becomes available, we may determine that some of our optimistic assumptions are not valid. At this point, we re-analyze the program, determine methods that may now need more memory barriers inserted in them, and re-compile these methods. We need to ensure that the program execution henceforth uses the new versions of re-compiled methods, so we invalidate the old code that was previously generated for these methods.

We consider the analysis results for a method to have changed in the last inter-procedural analysis phase if any of the following is valid:

1. The analysis determines a new type that some shared memory access in the method may refer to. This may happen either because the code for a method is available for the first time, or due to type resolution that takes into account newly loaded classes.
2. The analysis causes some *StartsAfter()*, *JoinsBefore()*, or *Concurrent()* set in any of the synchronization information summaries for the method to include a new element not previously contained in the set.
3. The analysis causes some locking set in any of the synchronization information summaries for the method to exclude some element previously contained in the set.
4. The analysis determines that the method may execute as part of additional thread types, or that the single thread type that executes the method no longer satisfies the single thread constraint (Section 3.2.3.1).

Figure 4.10 shows how we compute the set of methods that need to be invalidated. Recall that our delay set analysis checks for a delay edge between two accesses A and B in the same method by searching for conflict edges that involve A or B (Section 2.3). So when the analysis information for a method M changes, it may affect delay edges within all methods N such that a conflict edge exists between some access in M and some access in N .

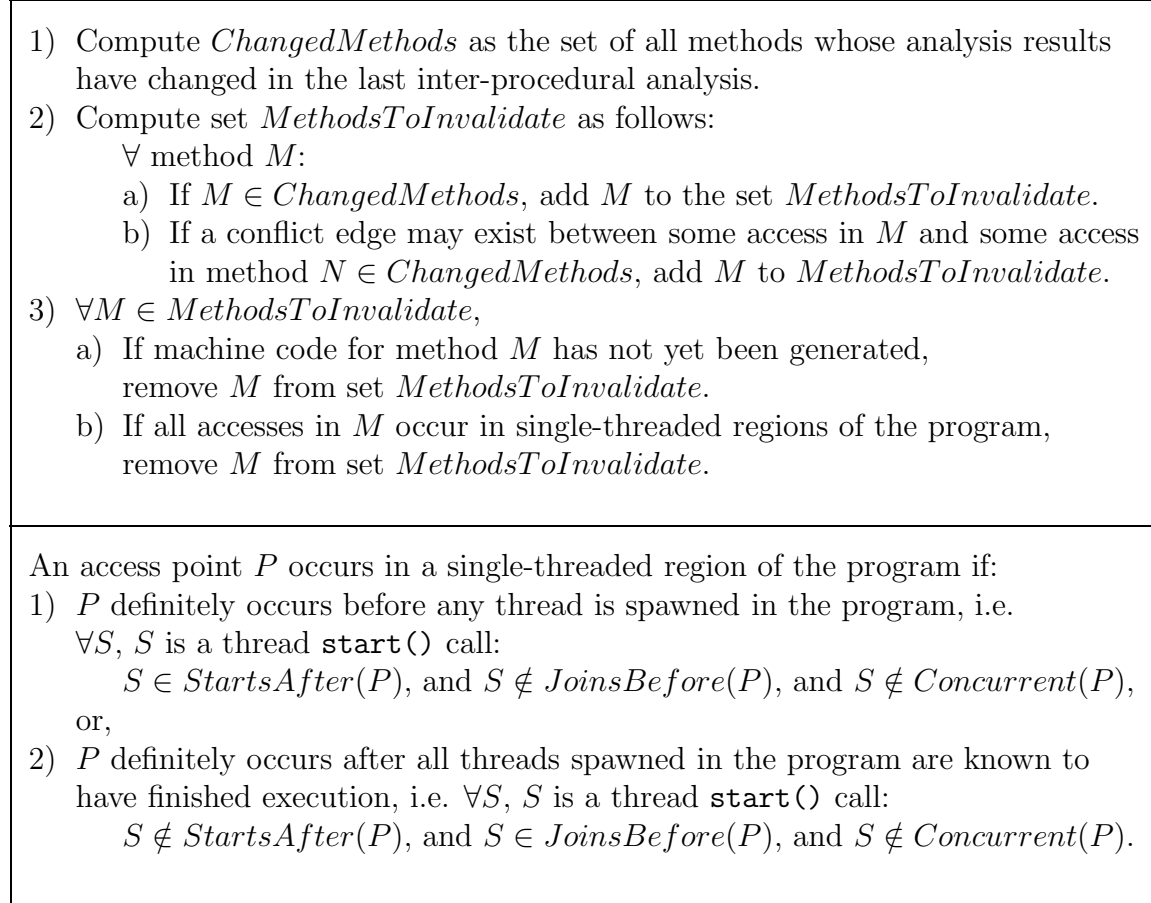


Figure 4.10: Determining the Set of Methods to Invalidate

We use the method invalidation mechanism provided in the Jikes RVM to invalidate the code generated for these methods. This ensures that any subsequent calls to the method will use newly generated code. However, this mechanism does not handle the case when a method call has already been invoked by a thread and is currently active on its execution stack. In this case, we would need to interrupt the thread execution

and modify its stack. For each stack frame that corresponds to a method that has been invalidated and re-compiled, we would need to re-write the frame to conform to the newly generated code. Also, before the thread is allowed to resume execution, we would need to ensure that all previous memory accesses issued by it are complete. Moreover, we would need to do this stack re-writing before any method that was analyzed for the first time in the last inter-procedural analysis begins its execution. In our implementation, we do not handle stack re-writing to invalidate methods that are currently executing. This is not an issue for the experiments we perform, as no methods on the stack need to be invalidated during execution of our benchmark programs.

4.4 Related Work

In [AG96], Adve and Gharachorloo review the different memory consistency models proposed in literature and implemented in commercial processors.

It has been widely conjectured that a system that supports sequential consistency is bound to degrade performance so much as to make it an impractical option. Research has been done to gauge the performance impact of implementing sequential consistency in hardware [GFV99, RPA97]. Experiments with some benchmark programs have shown that speculative techniques can be used to build sequentially consistent architectures with a performance degradation of about 20% on average [Hil98].

In our work, we determine the performance impact of implementing sequential consistency in software. We are aware of two publications that quantitatively compare the performance of a relaxed consistency model and sequential consistency implemented in software. In [LAY03], Liblit, Aiken, and Yelick develop a type system that identifies shared data accesses between multiple threads in SPMD Titanium programs [YSP⁺98]. They insert a memory barrier at each shared data access identified by the type inference system to show the performance difference between sequential consistency and the Titanium memory model (a relaxed memory model). However, it is difficult to gauge

the impact of their technique. This is because the specific benchmarks that they use show no significant performance benefit of using a relaxed consistency model over sequential consistency on the system used in their experiments. In [vP04], von Praun uses an object-based approach to help the compiler analyse what constraints to enforce for sequential consistency. Our approach is different because it uses inter-procedural analysis and considers control flow in the program, and this makes it more accurate. For the common benchmarks used, the work in [vP04] gives 98.6% performance degradation, whereas the techniques we use give a degradation of 11.5%. This comparison shows that compiler techniques can have a significant impact on performance.

It is encouraging to note that all these techniques are complementary, and a combination of them can potentially yield a generally applicable, low-overhead, sequentially consistent system.

Shen, Arvind, and Rudolph have proposed the Commit-Reconcile & Fences (CRF) notation and algebra for manipulating consistency models and cache protocols [SAR99]. This may be useful to specify the memory model in our memory-model aware compiler. Pugh has described problems with the semantics of the original memory model for the Java programming language[Pug99], and has led efforts to revise this memory model[Jav03].

Thread escape analysis identifies memory locations that may be accessed by multiple threads of execution. These are the memory locations that are of interest when performing synchronization and delay set analysis. Several researchers have studied thread escape analysis [VR01, Ruf00, BH99, CGS⁺99, Bla98]. In [vPG03], Praun and Gross use *object use graphs* that augment the prior escape analyses with the *happened-before* relation for object accesses to improve the precision of escape analysis. Recently, there has been research on using specialized programming language features, such as keywords, user annotations, or a restrictive type system to limit shared memory accesses [YSP⁺98, FF00, LAY03, BLR02].

Chapter 5

Experimental Evaluation

In our experiments, we determine the impact that our compiler techniques have on performance of benchmark programs compiled for sequential consistency. We measure the performance penalty incurred when sequential consistency is used as the programming language memory model, instead of the default consistency model implemented in the Jikes RVM. The Jikes RVM model is weak consistency, that allows shared memory accesses to be freely re-ordered except when:

- the re-ordering violates single-threaded data dependences, or
- the re-ordering is across an explicit synchronization point in the code. This includes thread structure, event-based, and lock-based synchronization described in Chapter 3. This does not include accesses to Java volatile variables.

5.1 Benchmark Programs

We use the following benchmark programs in our experiments to gauge the utility of synchronization analysis and delay set analysis:

1. *HashMap*: A program based on [Lea99b], designed to test the performance of `ConcurrentHashMap` from the implementation of JSR 166 [Lea]. This implementation allows reads and writes to access the hashmap concurrently, and uses Java volatile variables to synchronize these accesses. Each thread in the program continuously looks for an object in the hashmap. If the object is found, it deletes the object with 0.05 probability, else it adds the object with 0.1 probability. The program tries to mimic typical usage patterns where reads dominate all accesses.
2. *GeneticAlgo*: A parallel genetic algorithm based on the sequential version found in [Har98]. The implementation allows the selection, crossing-over, and mutation phases to proceed concurrently. It uses the `WaitFreeQueue` class from Doug Lea's concurrency utilities package, which is a precursor to the JSR 166 [Lea] implementation. The `WaitFreeQueue` class implementation includes volatile variables for synchronization.
3. *BoundedBuf*: A producer-consumer application that uses the `BoundedBuffer` class from Doug Lea's concurrency utilities package, which is a precursor to the JSR 166 [Lea] implementation. Threads mutate from being producers to consumers, and back again, depending on whether the buffer is full or empty.
4. *Sieve*: An implementation of sieve of Erasthones to generate the first "n" prime numbers. This is an example program in [Har98]. It is modified to use a bounded number of threads instead of "n" threads, to prevent out of memory errors due to excessive thread spawning.
5. *DiskSched*: A disk scheduler that uses an elevator algorithm to schedule concurrent read and write requests to different disk locations. This is based on an example program in Doug Lea's book [Lea99a]. The program includes busy-wait synchronization implemented using Java volatile variables.
6. *Raytrace*: A ray tracing application that is part of the Java Grande Forum multi-threaded benchmarks suite [jgf].

7. *Montecarlo*: A Monte Carlo simulation that is part of the Java Grande Forum multithreaded benchmarks suite [jgf].
8. *MolDyn*: A molecular dynamics application that is part of the Java Grande Forum multithreaded benchmarks suite [jgf].
9. *SPECmtrt*: A ray tracing application from the SPECJVM98 benchmark suite [Cor98]. This program is the only multithreaded application in the suite, and it uses dynamic class loading.
10. *SPECjbb00*: A server side application that is a standard SPEC benchmark [Cor00]. It focuses on the middle layer in business servers that is responsible for implementing business logic and manipulating objects. This program uses dynamic class loading.

The benchmarks we use include many standard multithreaded Java benchmarks, as well as applications that use concurrent implementations of data structures such as hashmaps and queues. These data structures are widely used in general applications, and the concurrent implementations are to be incorporated in the Java standard libraries. However, there are programs with characteristics that are not reflected in our set of benchmarks:

- Applications that are trivial to analyze: These include programs that have no data shared between threads, and a very simple escape analysis can detect this (e.g. the n-Queens problem, and some game codes). These also include programs that are computationally intensive (e.g. integration and interpolation functions). The computationally intensive programs perform just as well using a naive implementation for sequential consistency without any of our analysis, and using an implementation for relaxed consistency.
- Applications that need powerful alias analysis or shape analysis: None of the multithreaded Java benchmarks that we collected require powerful alias analysis or

shape analysis. For our analysis requirements on these benchmarks, a simple type-based alias analysis is sufficient. However, in general, there exist programs that use a recursive data structure such as a tree, and different threads work on data in distinct parts of the data structure (e.g. some programs in the Olden C++ benchmark suite have this property). Our type-based alias analysis cannot detect that individual portions of the data structure are not shared between threads, and so our speculation is that we would have significant performance degradation on these programs when sequential consistency is used.

- Applications that primarily need array dependence analysis: Many numerical programs analyze large amounts of data stored in arrays. These programs use data partitioning parallelism, and each thread accesses its own region in a shared array. In such cases, array dependence analysis is needed to determine that individual array elements do not escape ¹. The programs in Section 2 of the Java Grande Forum multithreaded benchmark suite [jgf] exhibit this property. We have previously experimented with a simple dependence analysis, and we obtained perfect results for 3 out of 5 Java Grande Forum Section 2 benchmarks. The two benchmarks that showed performance degradation either used subscripts of subscripts, or different array elements to synchronize different iterations of a loop. Our system lacks array dependence analysis. Since this analysis is expensive in general, there is a need to develop a sophisticated technique for it that takes advantage of the adaptive system. Such a technique would initially perform a simple analysis over all code, and then progressively refine the results for hot methods that take up a significant amount of execution time.

¹In our set of benchmark programs, *MolDyn* includes some arrays that are partitioned across threads, and we observe performance degradation for *MolDyn* due to the lack of array dependence analysis in our system.

Benchmark	Bytecodes	Methods	Thread Types	IP Analyses	Invalidates
Hashmap	24,989	108	1	8	0
GeneticAlgo	30,147	230	6	20	25
BoundedBuf	12,050	79	1	4	0
Sieve	10,811	976	2	3	0
DiskSched	21,186	1,220	2	14	18
Raytrace	33,198	113	1	16	0
Montecarlo	63,452	100	1	26	0
MolDyn	26,913	51	1	9	0
SPECmtrt	290,260	51	2	25	67
SPECjbb00	521,021	98	1	357	626

Table 5.1: Benchmark Characteristics

For each of the benchmarks we use, Table 5.1 shows the number of bytecodes that need to be analyzed, the number of methods in the call graph used for analysis, the number of thread types in the application, the number of inter-procedural analysis phases that are performed (Section 4.3.6), and the number of methods that need to be invalidated (Section 4.3.6.2).

5.2 Performance Impact

For our experiments, we used the memory-model aware compiler infrastructure described in Chapter 4. For the benchmark programs in Section 5.1, we obtained execution times for the following configurations:

1. *Base_RelaxedConsistency*: This is the default Jikes RVM execution without any of our changes. It uses weak consistency, and freely performs re-ordering optimizations like loop invariant code motion, load elimination, and dead store elimination.

2. *Escape_SequentialConsistency*: This is an implementation of sequential consistency that uses only the results of our escape analysis to determine the memory barriers that need to be inserted. It assumes there is a delay edge between a pair of shared memory accesses if both accesses may refer to escaping objects.
3. *Delay_SequentialConsistency*: This is an implementation of sequential consistency that uses escape analysis, type-based alias analysis, synchronization analysis, and delay set analysis to determine the memory barriers that need to be inserted.

We ran our experiments using a FastAdaptiveSemiSpace configuration for the Jikes RVM version 2.3.1. In this configuration, adaptive optimization is enabled, a copying garbage collector is used, and much of the Java code used to implement the virtual machine is pre-compiled and optimized. We set options so that each method is compiled using the optimizing compiler modified by us.

We obtain performance numbers for two architectures with different memory consistency models:

- Intel Xeon: We use a Dell PowerEdge 6600 SMP with 4 Intel hyperthreaded 1.5GHz Xeon processors (each having 1MB cache), and 6GB of system memory, running the Linux operating system. This platform provides a stronger hardware memory model than the Power3-based one.
- IBM Power3: We use an IBM SP 9076-550 with 8 375MHz processors, and 8GB of system memory, running the AIX operating system. This platform provides a more relaxed hardware memory model than the Xeon-based one.

5.2.1 Execution Times

For each benchmark program listed in Section 5.1, we obtain performance data by averaging over five executions of the program.

Intel Xeon Platform

Figure 5.1 and Table 5.2 show the execution times on the Intel Xeon platform for each of the three configurations, normalized to the *Base_RelaxedConsistency* configuration. We observe that our implementation of sequential consistency shows a slowdown of 10% on average over the default weak consistency model implemented by the Jikes RVM. Delay set analysis and synchronization analysis have a high impact on performance, since the average slowdown is 26.5 *times* when these two analyses are not used.




Figure 5.1: Slowdowns for Sequential Consistency on Intel Xeon

Benchmark	Slowdowns with respect to Base_RelaxedConsistency	
	Escape_SequentialConsistency	Delay_SequentialConsistency
Hashmap	7.76	1.01
GeneticAlgo	18.56	1.00
BoundedBuf	2.49	1.02
Sieve	3.54	1.01
DiskSched	13.34	1.50
Raytrace	138.06	1.01
Montecarlo	0.99	0.98
MolDyn	77.25	1.24
SPECmtrt	1.35	1.21
SPECjbb00	1.30	1.06

Table 5.2: Slowdowns for the Intel Xeon Platform

For 7 out of 10 benchmarks, the slowdowns were 6% or less. The three benchmarks that did not perform as well are:

- *DiskSched*: This application uses busy-wait synchronization that is implemented with the help of a volatile variable. A thread, say T1, tests the value of a volatile variable in each iteration of its busy-wait loop, and exits the loop only when the variable contains a specific value. This value is written by another thread in the program, say T2. Synchronization through the volatile variable ensures shared memory accesses that occur in T2 before the write, are completed before accesses that occur in T1 after the busy-wait loop. This ordering is not captured by our synchronization analysis, and delay set analysis detects delay edges in the code for T1 and T2 that do not have to be enforced for sequential consistency. This results in superfluous memory barriers being inserted, and a loss in performance.
- *MolDyn*: This application uses several shared arrays that are accessed by all threads. However, some of these arrays are such that each thread accesses one specific element of the array. Our analysis does not capture this. In previous

work [WFPS02], we developed a technique to perform array dependence analysis in Java. Note that in Java, array dependence analysis is complicated by the fact that all arrays are one-dimensional objects. Multi-dimensional arrays are considered to be arrays of array objects. Individual arrays in each dimension need not be distinct and they may have a differing number of elements. The analysis in [WFPS02] can handle the semantics of Java arrays, but it is too expensive to apply in the context of our just-in-time compiler system. In the absence of array dependence analysis, accesses to some arrays in the application are considered to be escaping, even though in reality each individual element of the array is accessed by a single thread. Thus, our analysis conservatively finds delay edges between these array accesses and inserts memory barriers to enforce them. This slows down the application in the case of sequential consistency.

- *SPECmtrt*: Threads in this application all read data from the same input file. As a result of this shared file reference, our conservative escape analysis determines several fields in the Java standard I/O library classes to be escaping. Delay set analysis detects delay edges between accesses to these fields, and memory barriers are inserted to enforce them. This contributes to some of the performance slowdown observed. However, as described in Section 5.3.1, a significant amount of the slowdown for *SPECmtrt* is due to adaptive re-compilation for optimization that overlaps with the application execution.

Power3 Platform

Figure 5.2 and Table 5.3 show the execution times for the Power3 platform for each of the three configurations, normalized to the *Base_RelaxedConsistency* configuration. We observe that our implementation of sequential consistency shows a slowdown of 26% on average over the default relaxed memory model implemented by the Jikes RVM. Again, delay set analysis and synchronization analysis have a significant impact on performance, and the average slowdown is 8.21 times when these two analyses are not used.




Figure 5.2: Slowdowns for Sequential Consistency on IBM Power3

Benchmark	Slowdowns with respect to Base_RelaxedConsistency	
	Escape_SequentialConsistency	Delay_SequentialConsistency
Hashmap	4.10	1.01
GeneticAlgo	11.25	1.06
BoundedBuf	1.27	1.03
Sieve	2.43	1.06
DiskSched	3.20	3.00
Raytrace	31.20	1.05
Montecarlo	1.67	1.00
MolDyn	23.93	1.25
SPECmtrt	1.15	1.06
SPECjbb00	1.94	1.07

Table 5.3: Slowdowns for the IBM Power3 Platform

For 8 out of 10 benchmarks, the slowdowns were 7% or less. The 2 benchmarks that did not perform as well are *DiskSched* and *MolDyn*, that did not perform well on the Intel Xeon platform either.

5.2.2 Effect of Escape Analysis

Escape analysis determines nodes in the access order graph for delay set analysis. The accuracy of escape analysis affects the accuracy of delay set analysis. Figure 5.3 shows the performance impact when our memory-model aware compiler uses the escape analysis implemented in the Jikes RVM distribution. The Jikes RVM escape analysis is a very conservative, simple, flow-insensitive analysis. It initially assumes that no variables escape. Then, each time a method is analyzed, a reference variable is considered to be escaping if:

1. it contains a reference used in a load or store operation, or
2. it is the destination of a load from a memory location, or
3. it is the source of a store to a memory location, or
4. it is the operand of a statement that returns a value from the method, or
5. it is the operand of a reference move statement.

Just like Figure 5.1, Figure 5.3 shows execution times on an Intel Xeon platform for each of the three configurations we test. The difference between the two is that we obtain the data in Figure 5.3 by replacing the escape analysis described in Section 4.3.1.1 with the default Jikes RVM escape analysis. In this case, delay set analysis is not very effective. We observe an average slowdown of 32.2 times without delay set analysis, and an average slowdown of 22.9 times with delay set analysis. In contrast, when our escape analysis is used, the performance improves from an average slowdown of 26.5 times without delay set




Figure 5.3: Slowdowns for Sequential Consistency Using the Jikes RVM Escape Analysis

analysis, to an average slowdown of 10% with it. Thus, the effect of our delay set analysis on execution time performance is highly sensitive to the accuracy of escape analysis.

5.2.3 Accuracy of Delay Set Analysis

Benchmark	Number of Delay Edges		
	Tested For	Need Enforcing	% Removed
Hashmap	16,089	60	99.6%
GeneticAlgo	9,857	367	96.3%
BoundedBuf	8,139	884	89.1%
Sieve	151,699	1,866	98.8%
DiskSched	4,041	58	98.6%
Raytrace	17,418	46	99.7%
Montecarlo	7,156	153	97.9%
MolDyn	13,370	6,924	48.2%
SPECmtrt	174,263	28,434	83.7%
SPECjbb00	1,415,865	115,580	91.8%

Table 5.4: Delay Edge Counts for the Intel Xeon Platform

For each benchmark program we use in our experiments, Table 5.4 shows statistics for performance of delay set analysis on the Intel Xeon platform. In our system, delay set analysis is invoked individually for each program edge to determine if it is a delay edge. The first column in Table 5.4 shows the number of times that the test for a delay edge is invoked during compilation. The second column shows the number of times this test finds a delay edge that needs to be enforced for correct execution under sequential consistency. Note that we consider only those delay tests that are non-trivial, i.e. tests performed after the application has spawned a thread. The last column shows the percentage of delay tests for which the analysis is able to determine the absence of a delay edge. This is an indication of the accuracy of delay set analysis: the higher the percentage, the more accurate the analysis. On the Intel Xeon platform, our delay set analysis can eliminate on average 90% of the delay edges that the system would otherwise enforce.

Benchmark	Number of Delay Edges		
	Tested For	Need Enforcing	% Removed
Hashmap	17,265	347	98.0%
GeneticAlgo	17,664	750	95.8%
BoundedBuf	3,890	824	78.8%
Sieve	5,256	1,121	78.7%
DiskSched	3,296	202	93.9%
Raytrace	78,786	47	99.9%
Montecarlo	16,568	314	98.1%
MolDyn	13,938	6,583	52.8%
SPECmtrt	150,501	19,861	86.8%
SPECjbb00	1,162,591	88,309	92.4%

Table 5.5: Delay Edge Counts for the Power3 Platform

Table 5.5 shows statistics for performance of delay set analysis on the Power3 platform (analogous to Table 5.4 for the Intel Xeon platform). In this case, we eliminate on average 87.5% of the delay edges that the system would otherwise enforce.

Note that execution time performance is not proportional to the number of delay edges found. This is because one memory barrier instruction may enforce multiple delay edges [FLM03]. The number of memory barriers inserted depends on the number of delay edges found, the specific pairs of accesses in the code that these delay edges are between, as well as the placement of barriers in the code that is determined by the memory barrier insertion algorithm.

5.2.4 Memory Barrier Instruction Counts

Benchmark	Static Memory Barriers			Dynamic Memory Barriers		
	Escape	Delay	% of Escape	Escape	Delay	% of Escape
Hashmap	1,459	6	0.4%	857,291 K	62	0.0%
GeneticAlgo	1,427	56	3.9%	898,186 K	4,066 K	0.5%
BoundedBuf	1,265	20	1.6%	3,918,543 K	263,951 K	6.7%
Sieve	9,860	394	4.0%	1,955,615 K	4,538 K	0.2%
DiskSched	1,217	9	0.7%	528,132 K	50,129 K	9.5%
Raytrace	1,235	11	0.9%	3,562,060 K	6 K	0.0%
Montecarlo	1,556	18	1.2%	5,925 K	40 K	0.7%
MolDyn	1,096	78	7.1%	1,710,646 K	26,057 K	1.5%
SPECmtrt	4,175	970	23.2%	9,354 K	3,700 K	39.6%
SPECjbb00	19,613	2,344	12.0%	1,245,037 K	238,103 K	19.1%

Table 5.6: Memory Fences Inserted and Executed on the Intel Xeon Platform

Table 5.6 shows memory barrier instruction counts for each of the benchmark programs we test on the Intel Xeon platform. For all these applications, no *lfences* or *sfences* are inserted (Section 4.3.4). Thus, all counts are for the number of *mfence* instructions. The table shows two sets of data:

1. Static Memory Barriers: the number of memory barrier instructions inserted by our compiler to enforce the set of delay edges determined.
2. Dynamic Memory Barriers: the number of memory barrier instructions dynamically executed when an application runs.

The first column in each set is the number of memory barrier instructions for the *Escape_SequentialConsistency* configuration that uses only escape analysis to enforce sequential consistency. The second column in each set is the number of memory barrier instructions for the *Delay_SequentialConsistency* configuration that uses all our

analyses to enforce sequential consistency. The last column in each set is the number of memory barrier instructions for the *Delay_SequentialConsistency* configuration expressed as a percentage of the number of memory barrier instructions for the *Escape_SequentialConsistency* configuration. The number of fences inserted and executed is much higher when our delay set analysis is not used.

Benchmark	Static Memory Barriers			Dynamic Memory Barriers		
	Escape	Delay	% of Escape	Escape	Delay	% of Escape
Hashmap	1,217	23	1.9%	446,031 K	129	0.0%
GeneticAlgo	1,414	76	5.4%	446,131 K	2,514 K	0.6%
BoundedBuf	1,134	25	2.2%	142,456 K	13,373 K	9.4%
Sieve	1,091	94	8.6%	215,313 K	4,964 K	2.3%
DiskSched	1,068	12	1.1%	13,696 K	15 K	0.1%
Raytrace	1,755	4	0.2%	12,609 K	1 K	0.0%
Montecarlo	1,461	56	3.8%	73,948 K	106 K	0.1%
MolDyn	960	79	8.2%	290,997 K	12,998 K	4.5%
SPECmtrt	3,735	914	24.5%	28,992 K	8,339 K	28.8%
SPECjbb00	17,039	2,657	15.6%	1,337,196 K	173,385 K	13.0%

Table 5.7: Sync Instructions Inserted and Executed on the Power3 Platform

For the Power3 platform, Table 5.7 shows memory barrier instruction (i.e. sync) counts for each of the benchmark programs we test (analogous to Table 5.6 for the Intel Xeon platform).

5.2.5 Optimizations and Memory Usage

Benchmark	Optimizations		Performance Cost With	Memory Overhead
	Inhibited	Performance Cost	All Optimizations Inhibited	
Hashmap	509 (75%)	0.0%	3.1%	44 M (62.3%)
GeneticAlgo	578 (63%)	0.0%	0.0%	13 M (12.1%)
BoundedBuf	491 (78%)	0.0%	0.0%	48 M (32.6%)
Sieve	803 (64%)	0.0%	0.0%	18 M (26.4%)
DiskSched	525 (75%)	0.0%	0.0%	71 M (97.2%)
Raytrace	864 (57%)	0.05%	0.05%	36 M (37.5%)
Montecarlo	718 (45%)	0.0%	4.0%	41 M (14.4%)
MolDyn	782 (66%)	0.0%	0.0%	36 M (46.7%)
SPECmtrt	1,141 (41%)	0.0%	0.0%	254 M (164.9%)
SPECjbb00	7,304 (59%)	0.1%	4.7%	517 M (106.9%)

Table 5.8: Optimizations Inhibited and Memory Overhead

In this section, we consider only code re-ordering and code elimination optimizations that depend on the memory consistency model. The first column of Table 5.8 shows the number of instances of these optimizations that are inhibited in our compiler to honor sequential consistency (Section 4.3.3). The majority of the inhibited optimization instances correspond to common subexpression elimination. We count the number of optimization instances performed for two cases: when a benchmark is executed with all optimizations enabled, and when a benchmark is executed with optimizations inhibited based on sequential consistency requirements. The difference of these two counts is the absolute number of optimizations inhibited, shown in the first column of the table. This column also shows the optimizations inhibited as a percentage of the total number of optimizations performed when all of them are enabled.

The second column in Table 5.8 shows the percentage slowdown in execution time when using sequential consistency over a relaxed model, that is due to inhibiting optimizations. We observe that even though a significant percentage of code re-ordering

Thread 1	Thread 2
<pre> for (i=0; i<10⁸; i++) Y.field += X.field; print (X.field, Y.field); </pre>	<pre> X.field = 10; Y.field = 3; </pre>

Figure 5.4: Example for Loop Invariant Code Motion

optimizations are inhibited, there is a negligible effect on execution time performance. The third column in Table 5.8 shows the percentage slowdown in execution time when using sequential consistency, that is due to inhibiting *all* code re-ordering and code elimination optimizations in the Jikes RVM. We find that these optimizations are not very effective for our set of benchmark programs, and only 3 benchmarks show any significant performance degradation when these optimizations are disabled: *Hashmap*, *Montecarlo*, and *SPECjbb00*. Thus, for our benchmarks, inhibiting optimizations for sequential consistency does not impact performance.

The last column in Table 5.8 shows the memory overhead of our compiler analysis, i.e. the heap space required by the analysis during program execution. It also gives the amount of this extra space as a percentage of the maximum heap space used by a program when it executes without our analysis.

To further study the effect of optimizations in our system, we use the microbenchmarks illustrated in Figure 5.4 and Figure 5.5. In these examples, X and Y refer to shared memory locations that are accessed by both threads in the example programs.

For the example in Figure 5.4, a traditional compiler may attempt to perform loop invariant code motion, and move the reference to `X.field` outside the `for` loop in Thread 1. In our sequentially consistent system, we inhibit this optimization if the compiler attempts to perform it (Section 4.3.3). When bytecode is generated by *javac*² for the Java

²We use the Sun *javac* compiler to translate Java source code to bytecodes. From Sun Java 2 SDK onwards, *javac* does not perform common optimizations such as loop invariant code motion, loop unrolling, algebraic simplification, or strength reduction [Hag01].

Thread 1	Thread 2
<pre> for (i=0; i<10⁸; i++) { T1 = X.field; T2 = T1 + 10; T3 = Y.field; T4 = X.field; T5 = T4 + 10; } print (T2, T3, T5); </pre>	<pre> X.field = 7; Y.field = 3; </pre>

Figure 5.5: Example for Load Elimination

source code corresponding to this example, loop invariant code motion is not performed. Also, we inspect the final machine code generated by the default Jikes RVM system, and determine that it does not perform loop invariant code motion for this example. Thus, in this case, inhibiting optimizations for sequential consistency has no effect, as the optimization is not performed at all in the default Jikes RVM system that we use to build our compiler.

For the example in Figure 5.5, a traditional compiler may attempt to eliminate the load of `X.field` into `T4`, since `X.field` has previously been loaded into `T1`. Also, the compiler may attempt to re-use the common subexpression computed in `T2` for uses of `T5`. In our sequentially consistent system, we inhibit these optimizations if the compiler attempts to perform them (Section 4.3.3). When bytecode is generated by *javac* for the Java source code corresponding to this example, neither the load elimination for `T4`, nor the common subexpression elimination for `T5` is performed. When we inspect the final machine code generated by the default Jikes RVM system, we determine that it eliminates the load into `T4`, and uses the value loaded in `T2` instead. However, it does not perform common subexpression elimination for `T5`.

For this second microbenchmark, we obtained the following execution times:

1. *Base_RelaxedConsistency* configuration: 1.35 seconds.

2. *Base_RelaxedConsistency* configuration with all load eliminations disabled: 1.35 seconds.
3. *Delay_SequentialConsistency* configuration: 66.7 seconds.
4. *Delay_SequentialConsistency* configuration with no load eliminations disabled: 56.7 seconds.

From these results, we observe that inhibiting optimizations may impact performance. The slowdown for sequential consistency over the default Jikes RVM memory model is 49 times. However, if load elimination is not disabled and only memory barriers are inserted, the slowdown is 42 times. The example also illustrates the importance of reducing the number of memory barriers inserted in hot methods, since this can lead to large slowdowns otherwise.

5.2.6 Effect of Synchronization Analysis

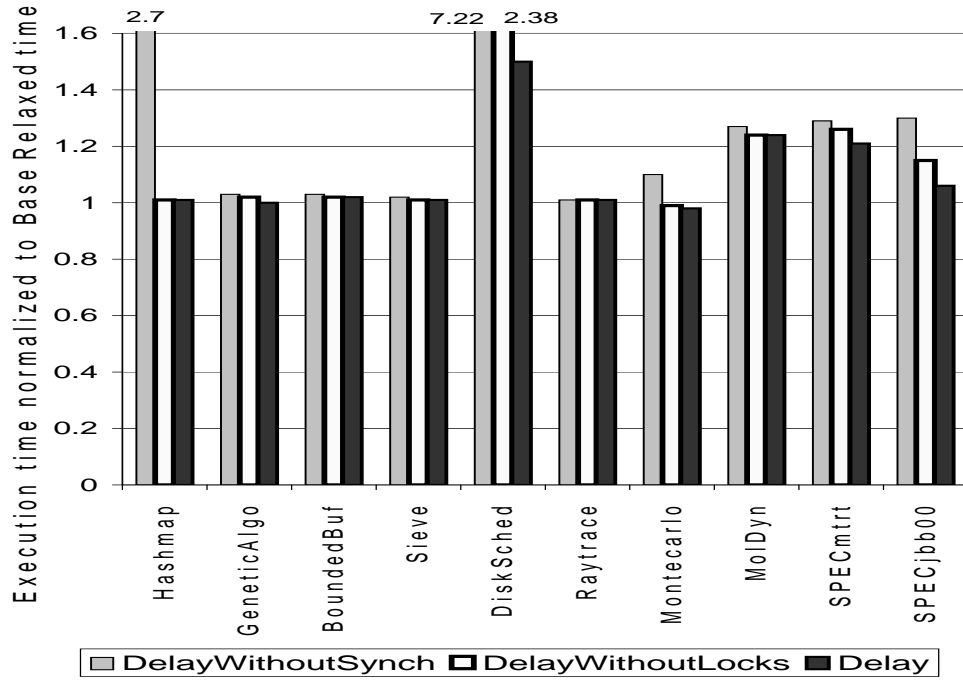


Figure 5.6: Effect of Thread Structure Analysis and Locking Synchronization

Figure 5.6 shows the effect of synchronization analysis on execution time performance. The figure shows the execution times, normalized to the *Base_RelaxedConsistency* configuration, for the following three cases:

1. *DelayWithoutSynch*: This is an implementation of sequential consistency that uses escape analysis, type-based alias analysis, and delay set analysis to determine the memory barriers that need to be inserted. It does not use synchronization analysis.
2. *DelayWithoutLocks*: This is an implementation of sequential consistency that uses escape analysis, type-based alias analysis, thread structure analysis, and delay set analysis to determine the memory barriers that need to be inserted. It does not use locking information based on synchronized blocks in Java.
3. *Delay*: This is the same as the *Delay_SequentialConsistency* configuration described earlier, that uses escape analysis, type-based alias analysis, synchronization analysis, and delay set analysis to determine the memory barriers that need to be inserted.

The effect of synchronization analysis depends on the particular characteristics of a specific program. For our benchmarks, synchronization analysis has a significant impact on the performance of 5 out of 10 programs. Performance of *Hashmap* improves by 169% over the *Base_RelaxedConsistency* configuration, *DiskSched* by 572%, *Montecarlo* by 10%, *SPECmtrt* by 8%, and *SPECjbb00* by 24%. The locking analysis component of synchronization analysis improves the performance of *DiskSched* by 88%, *SPECmtrt* by 5%, and *SPECjbb00* by 9%.

5.3 Analysis Times

The performance numbers in Figure 5.1 do not include the compile time for an application, except for re-compilation due to adaptive optimization. On the Intel Xeon platform,

the compile time is significantly large for three benchmarks *MolDyn*, *SPECmtrt*, and *SPECjbb00*. *SPECjbb00* is a server application that runs for long periods of time in practice. We expect the compilation overhead to be amortized over this long running time. A similar argument applies to *MolDyn* and *SPECmtrt* if they are used to process large data sets, and take a long time to finish execution. Note that a significant portion of the compilation overhead for these benchmarks is due to adaptive optimization of methods (Section 5.3.1). On the Power3 platform, compile time is significantly large enough to impact performance numbers for most of our benchmarks. In this case, the applications will have to run for longer periods of time if the compilation cost is to be amortized.

Benchmark	Delay	Synch	Escape	Barrier	TotalSC	TotalRC
Hashmap	0.1 (7.7)	0.6 (46.2)	1.2 (92.3)	0.9 (69.2)	4.7 (361.5)	1.3
GeneticAlgo	0.1 (5.6)	0.8 (44.4)	1.0 (55.6)	0.7 (38.9)	4.8 (266.7)	1.8
BoundedBuf	0.02 (1.2)	0.3 (17.6)	0.5 (29.4)	0.7 (41.2)	2.9 (170.6)	1.7
Sieve	0.6 (10.5)	0.2 (3.5)	0.3 (5.3)	2.1 (36.8)	8.9 (156.1)	5.7
DiskSched	0.1 (8.3)	0.5 (41.7)	1.9 (158.3)	0.7 (58.3)	4.9 (408.3)	1.2
Raytrace	0.1 (3.8)	0.7 (26.9)	1.5 (57.7)	1.4 (53.8)	6.6 (253.8)	2.6
Montecarlo	0.1 (4.8)	1.1 (52.4)	2.0 (95.2)	0.8 (38.1)	7.2 (342.9)	2.1
MolDyn	0.7 (38.9)	0.6 (33.3)	1.4 (77.8)	14.2 (788.9)	22.6 (1255.6)	1.8
SPECmtrt	5.5 (83.3)	9.1 (137.9)	57.9 (877.3)	3.2 (48.5)	89.1 (1350.0)	6.6
SPECjbb00	60.1 (140.4)	33.0 (77.1)	340.1 (794.6)	125.1 (292.3)	716.5 (1674.1)	42.8

Table 5.9: Analysis Times in Seconds for the Intel Xeon Platform

Table 5.9 shows the total analysis times (including all re-compilations) in seconds for each benchmark executing on the Intel Xeon platform. The first five columns also show in brackets the time for their respective analysis components expressed as a percentage of the total Jikes RVM base compile time, shown in the column titled *TotalRC*.

- *Delay*: This column shows the time taken by delay set analysis.

- *Synch*: This column shows the time taken by synchronization analysis.
- *Escape*: This column shows the time taken by escape analysis.
- *Barrier*: This column shows the time taken by memory barrier insertion.
- *TotalSC*: This column shows the total compilation time in the *Delay_SequentialConsistency* configuration when all analyses are enabled.
- *TotalRC*: This column shows the total compilation time in the *Base_RelaxedConsistency* configuration when none of our analysis is performed.

We observe that delay set analysis and synchronization analysis take at most 1.1 seconds each for 8 out of 10 programs. Also, for 7 out of 10 programs, the difference between the total compile time for our implementation of sequential consistency, and the total compile time for the default Jikes RVM system is at most 5 seconds. The 3 programs that have higher overheads are *MolDyn*, *SPECmtrt*, and *SPECjbb00*. Most of the overhead for *MolDyn* comes from memory barrier insertion. *SPECmtrt* and *SPECjbb00* both make use of dynamic class loading, which leads to many non-trivial inter-procedural analysis phases and re-compilations (Section 4.3.6). These contribute to increase the analysis time.

Benchmark	Delay	Synch	Escape	Barrier	TotalSC	TotalRC
Hashmap	0.2 (4.8)	2.0 (47.6)	7.9 (188.1)	2.3 (54.8)	18.2 (433.3)	4.2
GeneticAlgo	0.3 (6.3)	2.5 (52.1)	5.7 (118.8)	2.3 (47.9)	16.6 (345.8)	4.8
BoundedBuf	0.03 (0.7)	0.9 (22.0)	1.9 (46.3)	1.9 (46.3)	9.1 (222.0)	4.1
Sieve	0.1 (2.9)	0.8 (23.5)	1.1 (32.4)	1.9 (55.9)	7.9 (232.4)	3.4
DiskSched	0.1 (3.1)	1.7 (53.1)	6.1 (190.6)	1.9 (59.4)	13.9 (434.4)	3.2
Raytrace	0.3 (3.8)	2.3 (28.8)	7.5 (93.8)	4.9 (61.3)	24.3 (303.8)	8.0
Montecarlo	0.3 (3.6)	4.7 (60.0)	9.7 (115.5)	2.7 (32.1)	25.7 (306.0)	8.4
MolDyn	0.4 (8.3)	1.8 (37.5)	5.0 (104.2)	2.6 (54.2)	17.5 (364.6)	4.8
SPECmtrt	12.0 (82.2)	19.4 (132.9)	93.1 (637.7)	11.7 (80.1)	179.7 (1230.8)	14.6
SPECjbb00	45.0 (94.3)	31.1 (65.2)	361.2 (757.2)	118.5 (248.4)	716.7 (1502.5)	47.7

Table 5.10: Analysis Times in Seconds for the Power3 Platform

Table 5.10 shows the total analysis times (including all re-compilations) in seconds for each benchmark executing on the Power3 platform (analogous to Table 5.9). For most benchmarks, the absolute compile times are higher than the times on the Intel Xeon platform. The base compile time for the *Base-RelaxedConsistency* configuration is also higher in general.

5.3.1 Effect of Adaptive Re-compilation for Optimization

We base our compiler on the adaptive configuration of the Jikes RVM that profiles an application as it executes, and attempts to dynamically optimize methods that take a significant percentage of the execution time. The re-compilation of methods for optimization purposes can be switched off using a command-line option. We experimented with this option to determine the effect of adaptive re-compilations on analysis times and execution time performance.

Benchmark	Base	Delay
BoundedBuf	1	3
Sieve	1	6
DiskSched	1	6
MolDyn	1	5
SPECmtrt	4	23
SPECjbb00	6	23

Table 5.11: Number of Methods Re-compiled for Adaptive Optimization

Benchmark	Delay	Synch	Escape	Barrier	TotalSC	TotalRC
BoundedBuf	0.02 (2.0)	0.2 (20.0)	0.5 (50.0)	0.5 (50.0)	2.4 (240.0)	1.0
Sieve	0.02 (2.0)	0.2 (20.0)	0.4 (40.0)	0.6 (60.0)	2.3 (230.0)	1.0
DiskSched	0.1 (9.1)	0.5 (45.5)	1.6 (145.5)	0.6 (54.5)	3.9 (354.5)	1.1
MolDyn	0.2 (12.5)	0.6 (37.5)	1.7 (106.3)	0.9 (56.3)	6.6 (412.5)	1.6
SPECmtrt	3.6 (69.2)	9.1 (175.0)	39.4 (757.7)	2.7 (51.9)	65.0 (1250.0)	5.2
SPECjbb00	9.9 (34.9)	24.6 (86.6)	63.2 (222.5)	62.0 (218.3)	212.3 (747.5)	28.4

Table 5.12: Analysis Times in Seconds With No Re-compilation for Optimization

Table 5.11 shows the number of methods that are re-compiled for adaptive optimization in the *Base_RelaxedConsistency* configuration (Base) and the *Delay_SequentialConsistency* configuration (Delay). Table 5.12 shows the analysis times analogous to Table 5.9 (Intel Xeon platform) for the case when re-compilations for adaptive optimization are disabled. The table includes only those benchmarks that show a difference in compile times for this case. We observe a significant reduction in the compile time for the 3 programs that have the highest compile time overheads. The total compile time for *MolDyn* reduced by 70.8%, for *SPECmtrt* by 27%, and for *SPECjbb00* by 70.4%. Moreover, the performance of *SPECmtrt* and *SPECjbb00* improved in this case. *SPECmtrt* shows a slowdown of 8.8%, and *SPECjbb00* shows a slowdown of 2% over the *Base_RelaxedConsistency* configuration. However, 3 programs show a performance

degradation when re-compilations are disabled. For the *Delay_SequentialConsistency* configuration, *MolDyn*, *Montecarlo*, and *Raytrace* respectively perform 10%, 15%, and 20% worse when re-compilations are turned off. Thus, adaptive re-compilation is important in general, and disabling it is not a solution. We believe it is worth exploring how to fine-tune the use of adaptive re-compilation in the context of a compiler like ours, that performs expensive inter-procedural analysis.

Chapter 6

Conclusion

6.1 Contributions

This thesis makes the following contributions:

- It describes a fast and effective algorithm for delay set analysis that is applicable to MIMD programs. The algorithm has polynomial time complexity in the number of shared memory accesses. Program synchronization information is used to improve the accuracy of the results.
- It describes a dataflow analysis algorithm to determine orders enforced by explicit synchronization in the program. The algorithm computes ordering information that includes happens-before, happens-after, and happens-in-parallel relations, and represents the results in a compact form. The algorithm places no restraints on recursion or dynamic thread creation, and is more precise than previous methods based on dataflow analysis.
- It describes our memory-model aware compiler and the range of analysis techniques needed to account for the effect of inter-thread memory accesses.

- It evaluates performance of our proposed algorithms in the framework of a Java just-in-time compiler. It reports data to compare the performance of two programming language memory models: sequential consistency and weak consistency.

The synchronization and delay set analysis algorithms are designed to be fast. Experience with the implementation indicates that the analysis times required are within 15% of the total compile time in most cases. On the Intel Xeon platform, the absolute analysis times are in the order of a few seconds for all applications except *SPECjbb00*. Thus, these algorithms may be implemented in a just-in-time compiler where the compile time contributes to the application execution time. Several factors contribute to make our analysis fast:

- We use a conservative and simple algorithm for delay set analysis. This algorithm is effective because for the most part applications communicate through shared memory in a disciplined manner using specific sharing patterns.
- Our synchronization analysis is efficient since only orders involving a synchronization construct are determined. Applications typically include relatively few synchronization constructs.
- We summarize the analysis results for each method. This makes the inter-procedural analysis fast since the effect of a method call can be easily incorporated when testing for shared memory access orders to enforce.
- We take advantage of user-defined types in the programming language to determine if two shared memory locations may be the same. In most cases, user-defined types naturally group together locations with similar access properties. This enables the system to perform delay set analysis with a high level of accuracy without having to use an expensive alias analysis phase.

Experiments on a set of benchmark programs show that our techniques are effective in reducing the number of delays that need to be enforced when implementing sequential

consistency in software. This results in low performance overheads for a sequentially consistent system over a system that uses weak consistency. In our experiments, we observe a slowdown of 10% on average for our benchmarks executing on an architecture based on the Intel Xeon processor, and 26% on average for an architecture based on the IBM Power3. Moreover, 7 out of 10 applications on the Intel Xeon platform, and 8 out of 10 applications on the Power3 platform, show a performance loss of 7% or less.

This is important because sequential consistency is the memory consistency model that enforces the most constraints on an implementation. We show the possibility of achieving reasonable performance with sequential consistency. This opens up a wide space for the design of programming language memory models since performance need no longer be a prohibitive factor. We can consider using models that are more stringent than the relaxed consistency models popular today. Use of such strong models may help make parallel programming more efficient.

6.2 Open Problems

In this work, we use the analysis techniques described to reduce the number of fences that need to be inserted when implementing memory consistency in software. These analysis results can also be applied to optimize programs by:

- Removing redundant synchronization: We can use our analyses to determine delay edges without taking into account some synchronization constructs in the program. Then we can match these delay edges with orders enforced by the synchronizations constructs, and determine if some of the synchronizations are unnecessary.
- Improving thread scheduling: We can apply our analysis results to infer knowledge about code sections that may execute independently of each other. This knowledge can be used to determine when or where to schedule threads so that the overall execution is efficient.

- Applying code specialization: Our analysis results retain information about the source of concurrency in terms of the `start()` calls that spawn threads that may be concurrent with a program point. This information can be used by a code specialization phase to optimize code sections that are common to multiple thread types, and may be executed by a thread of any of these types. Code specialization generates multiple versions of a piece of code, each optimized for use in different contexts.

The analysis can also be used to aid program development and debugging by:

- Statically detecting possible races and deadlocks in the program code.
- Reducing the overhead of dynamic data race detection by eliminating the need to track some shared memory accesses in the program.
- Efficiently inserting fences to automatically enforce orders on a subset of shared memory accesses specified by the programmer. This can reduce the complexity of reasoning about different permutations of shared memory accesses when developing a program.

Bibliography

- [AAB⁺00] B. Alpern, C. R. Attanasio, J. J. Barton, M. G. Burke, P.Cheng, J.-D. Choi, A. Cocchi, S. J. Fink, D. Grove, M. Hind, S. F. Hummel, D. Lieber, V. Litvinov, M. F. Mergen, T. Ngo, J. R. Russell, V. Sarkar, M. J. Serrano, J. C. Shepherd, S. E. Smith, V. C. Sreedhar, H. Srinivasan, and J. Whaley. The Jalapeño Virtual Machine. *IBM System Journal*, 39(1), February 2000.
- [AFG⁺00] Matthew Arnold, Stephen Fink, David Grove, Michael Hind, and Peter F. Sweeney. Adaptive Optimization in the Jalapeño JVM. *Third ACM Workshop on Feedback-Directed and Dynamic Optimization*, December 2000.
- [AG96] Sarita V. Adve and Kourosh Gharachorloo. Shared Memory Consistency Models: A Tutorial. *IEEE Computer*, pages 66–76, December 1996.
- [AH90] Sarita V. Adve and Mark D. Hill. Weak Ordering - A New Definition. In *Proceedings of The 17th Annual International Symposium on Computer Architecture (ISCA)*, pages 2–14, May 1990.
- [AK02] Randy Allen and Ken Kennedy. *Optimizing Compilers for Modern Architectures*. Morgan Kaufmann Publishers, 2002.
- [BH99] Jeff Bogda and Urs Holzle. Removing Unnecessary Synchronization in Java. In *Proceedings of the 14th ACM SIGPLAN Conference on Object-oriented Programming, Systems, Languages, and Applications*, pages 35–46. ACM Press, 1999.

- [BHJ⁺03] Konstantin Berlin, Jun Huan, Mary Jacob, Garima Kochhar, Jan Prins, Bill Pugh, P. Sadayappan, Jaime Spacco, and Chau-Wen Tseng. Evaluating the Impact of Programming Language Features on the Performance of Parallel Applications on Cluster Architectures. In *16th Annual Workshop on Languages and Compilers for Parallel Computing*, October 2003.
- [BK89] Vasanth Balasundaram and Ken Kennedy. Compile-time Detection of Race Conditions in a Parallel Program. In *Proceedings of the 3rd international conference on Supercomputing*, pages 175–185. ACM Press, 1989.
- [Bla98] Bruno Blanchet. Escape Analysis: Correctness Proof, Implementation and Experimental Results. In *Proceedings of the 25th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 25–37. ACM Press, 1998.
- [BLR02] Chandrasekhar Boyapati, Robert Lee, and Martin Rinard. Ownership Types for Safe Programming: Preventing Data Races and Deadlocks. In *Object-oriented Programming, Systems, Languages, and Applications (OOPSLA)*, November 2002.
- [CADG⁺93] David E. Culler, Andrea C. Arpaci-Dusseau, Seth Copen Goldstein, Arvind Krishnamurthy, Steven Lumetta, Thorsten von Eicken, and Katherine A. Yelick. Parallel programming in Split-C. In *Supercomputing*, pages 262–273, 1993.
- [CGHS99] Jong-Deok Choi, David Grove, Michael Hind, and Vivek Sarkar. Efficient and Precise Modeling of Exceptions for the Analysis of Java Programs. *ACM SIGPLAN-SIGSOFT Workshop on Program Analysis for Software Tools and Engineering (PASTE)*, September 1999.

- [CGS⁺99] Jong-Deok Choi, Manish Gupta, Mauricio Serrano, Vugranam C. Sreedhar, and Samuel P. Midkiff. Escape Analysis for Java. In *Proceedings of the ACM SIGPLAN 1999 Conference on Object-oriented Programming Systems, Languages, and Applications (OOPSLA)*, pages 1–19, November 1999.
- [CKS90] David Callahan, Ken Kennedy, and Jaspal Subhlok. Analysis of Event Synchronization in a Parallel Programming Tool. In *Proceedings of the Second ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, pages 21–30, Seattle WA, 1990.
- [CKY03] Wei-Yu Chen, Arvind Krishnamurthy, and Katherine Yelick. Polynomial-Time Algorithms for Enforcing Sequential Consistency in SPMD Programs with Arrays. In *Sixteenth Annual Workshop on Languages and Compilers for Parallel Computing*, October 2003.
- [CL97] M. Cierniak and W. Li. Just-in-time Optimization for High-performance Java Programs. *Concurrency: Practice and Experience*, 9(11):1063–73, November 1997.
- [CLL⁺02] J.-D. Choi, K. Lee, A. Loginov, R. O’Callahan, V. Sarkar, and M. Sridharan. Efficient and Precise Datarace Detection for Multithreaded Object-oriented Programs. In *Proceedings of the ACM SIGPLAN 2002 Conference on Programming Language Design and Implementation (PLDI)*, pages 258–269, June 2002.
- [Cor98] The Standard Performance Evaluation Corporation. SPEC JVM Client98 Suite. 1998. URL: <http://www.specbench.org/jvm98/jvm98>.
- [Cor00] The Standard Performance Evaluation Corporation. SPEC JBB 2000 Benchmark. 2000. URL: <http://www.specbench.org/jbb2000>.

- [CS89] D. Callahan and J. Subhlok. Static Analysis of Low-level Synchronization. In *Proceedings of the ACM SIGPLAN and SIGOPS Workshop on Parallel and Distributed Debugging*, pages 100–111, January 1989.
- [DS91] Evelyn Duesterwald and Mary Lou Soffa. Concurrency Analysis in the Presence of Procedures Using a Data-flow Framework. In *Proceedings of the Symposium on Testing, Analysis, and Verification*, pages 36–48. ACM Press, 1991.
- [DSB88] M. Dubois, C. Scheurich, and F.A. Briggs. Synchronization, Coherence, and Event Ordering in Multiprocessors. 21(2):9–21, 1988.
- [EGP89] Perry A. Emrath, Sanjoy Ghosh, and David A. Padua. Event Synchronization Analysis for Debugging Parallel Programs. In *Proceedings of Supercomputing '89*, pages 580–588, 1989.
- [FF00] Cormac Flanagan and Stephen N. Freund. Type-based race detection for Java. *ACM SIGPLAN Notices*, 35(5):219–232, 2000.
- [FLM03] Xing Fang, Jaejin Lee, and Samuel P. Midkiff. Automatic Fence Insertion for Shared Memory Processing. In *2003 ACM International Conference on Supercomputing*, June 2003.
- [FLR98] Matteo Frigo, Charles E. Leiserson, and Keith H. Randall. The Implementation of the Cilk-5 Multithreaded Language. In *SIGPLAN Conference on Programming Language Design and Implementation*, pages 212–223, 1998.
- [GFV99] K. Gniady, B. Falsafi, and T. Vijaykumar. Is SC + ILP = RC? In *Proc. of the 26th Annual International Symposium on Computer Architecture (ISCA '99)*, 1999.
- [GJS96] James Gosling, Bill Joy, and Guy Steele. *The Java Language Specification, Second Edition*. Addison-Wesley, 1996.

- [GLL⁺90] Kourosh Gharachorloo, Daniel Lenoski, James Laudon, Phillip Gibbons, Anoop Gupta, and John Hennessy. Memory Consistency and Event Ordering in Scalable Shared-Memory Multiprocessors. In *Proceedings of The 17th Annual International Symposium on Computer Architecture (ISCA)*, pages 15–26, May 1990.
- [GS93] D. Grunwald and H. Srinivasan. Data Flow Equations for Explicitly Parallel Programs. In *Proceedings of the Fourth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, San Diego, CA, 1993.
- [Hag01] Peter Hagggar. Understanding Bytecode Makes You a Better Programmer. July 2001. URL: www-106.ibm.com/developerworks/ibm/library/it-hagggar_bytecode.
- [Har98] Stephen Hartley. *Concurrent Programming: the Java Programming Language*. Oxford University Press, 1998.
- [Hil98] Mark D. Hill. Multiprocessors Should Support Simple Memory-Consistency Models. *IEEE Computer*, August 1998.
- [HM91] David P. Helmbold and Charles E. McDowell. Computing Reachable States of Parallel Programs. *Proceedings of the ACM/ONR Workshop on Parallel and Distributed Debugging, published in ACM SIGPLAN Notices*, 26(12):76–84, 1991.
- [IA3] IA-32 Intel Architecture Software Developer’s Manual, Volume 2 and 3. URL: developer.intel.com/design/pentium4/manuals/index_new.htm.
- [Jav03] JavaMemoryModel. Java Memory Model Mailing List. In *Archive at* http://www.cs.umd.edu/~pugh/java/memoryModel/arc_hive/, 2003.
- [jgf] The Java Grande Forum Multi-threaded Benchmarks. URL: <http://www.epcc.ed.ac.uk/javagrande/threads/contents.html>.

- [KJCS99] Sanjeev Kumar, Dongming Jiang, Rohit Chandra, and Jaswinder Pal Singh. Evaluating Synchronization on Shared Address Space Multiprocessors: Methodology and Performance. In *Proceedings of the 1999 ACM SIGMETRICS International Conference on Measurement and Modeling of Computer Systems*, pages 23–34. ACM Press, 1999.
- [KSV96] Jens Knoop, Bernhard Steffen, and Jürgen Vollmer. Parallelism for Free: Efficient and Optimal Bitvector Analyses for Parallel Programs. *ACM Transactions on Programming Languages and Systems*, 18(3):268–299, May 1996.
- [KY94] Arvind Krishnamurthy and Katherine Yelick. Optimizing Parallel SPMD Programs. In *Seventh Annual Workshop on Languages and Compilers for Parallel Computing*, August 1994.
- [KY96] Arvind Krishnamurthy and Katherine Yelick. Analyses and Optimizations for Shared Address Space Programs. *Journal of Parallel and Distributed Computing*, 38:139–144, 1996.
- [Lam79] Leslie Lamport. How to Make a Multiprocessor Computer That Correctly Executes Multiprocess Programs. *IEEE Transactions on Computers*, C-28(9):690–691, September 1979.
- [LAY03] Ben Liblit, Alex Aiken, and Katherine Yelick. Type Systems for Distributed Data Sharing. In *Proceedings of the Tenth International Static Analysis Symposium*, 2003.
- [Lea] Doug Lea. Java Specification Request (JSR) 166: Concurrency Utilities. In <http://gee.cs.oswego.edu/dl/concurrency-interest/index.html> .
- [Lea99a] Doug Lea. *Concurrent Programming in Java*. Addison Wesley, 1999. URL: <http://gee.cs.oswego.edu/dl/cpj>.
- [Lea99b] Doug Lea. JavaMemoryModel: recap: concurrent reads. December 1999. URL: www.cs.umd.edu/~pugh/java/memoryModel/archive/0358.html.

- [LMP97] Jaejin Lee, Samuel P. Midkiff, and David A. Padua. Concurrent Static Single Assignment Form and Constant Propagation for Explicitly Parallel Programs. In Z. Li, P.-C. Yew, S. Chatterjee, C.-H. Huang, P. Sadayappan, and D. Sehr, editors, *Proceedings of The 10th International Workshop on Languages and Compilers for Parallel Computing*, number 1366 in Lecture Notes in Computer Science, pages 114–130. Springer, August 1997.
- [LP00] Jaejin Lee and David A. Padua. Hiding Relaxed Memory Consistency with Compilers. In *Proceedings of the IEEE International Conference on Parallel Architectures and Compilation Techniques (PACT)*, pages 111–122, October 2000.
- [LPM99] Jaejin Lee, David A. Padua, and Samuel P. Midkiff. Basic Compiler Algorithms for Parallel Programs. In *Proceedings of the 7th ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 1–12, May 1999.
- [MC93] John M. Mellor-Crummey. Compile-Time Support for Efficient Data Race Detection in Shared-Memory Parallel Programs. In *Workshop on Parallel and Distributed Debugging*, pages 129–139, 1993.
- [Mid95] S. Midkiff. Dependence Analysis in Parallel Loops with $i+/-k$ Subscripts. In *1995 Workshop on Languages and Compilers for Parallel Computing*, 1995. available as Springer Lecture Notes in Computer Science Vol. N. 1033.
- [MP87] Samuel P. Midkiff and David A. Padua. Compiler Algorithms for Synchronization. *IEEE Transactions on Computers*, C-36(12):1485–1495, December 1987.
- [MP90] S. Midkiff and D. Padua. Issues in the Compile-Time Optimization of Parallel Programs. In *Proceedings of the 1990 International Conference on Parallel Processing, Vol. II*, pages 105–113, August 1990.

- [MP01] J. Manson and W. Pugh. Core Semantics of Multithreaded Java. In *Proceedings of the ACM SIGPLAN 2001 ISCOPE/Java Grande Conference*, pages 29–38, 2001.
- [MPC90] Samuel P. Midkiff, David A. Padua, and Ron Cytron. Compiling Programs with User Parallelism. In *Languages and Compilers for Parallel Computing*, pages 402–422, 1990.
- [MR93] Stephen P. Masticola and Barbara G. Ryder. Non-concurrency Analysis. In *Proceedings of the Fourth ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming*, pages 129–138. ACM Press, 1993.
- [NAC99] Gleb Naumovich, George S. Avruninand, and Lori A. Clarke. An Efficient Algorithm for Computing MHP Information for Concurrent Java Programs. In *Proceedings of Seventh European Software Engineering Conference and Seventh ACM SIGSOFT Symposium on the Foundations of Software Engineering*, Sept. 1999.
- [NG92] Robert H. B. Netzer and Sanjoy Ghosh. Efficient Race Condition Detection for Shared-Memory Programs with Post/Wait Synchronization. In *Proceedings of the 1992 International Conference on Parallel Processing*, volume II, Software, pages II:242–246, Boca Raton, Florida, 1992. CRC Press.
- [NM90] Robert H. B. Netzer and Barton P. Miller. On the Complexity of Event Ordering for Shared-Memory Parallel Program Executions. In *Proceedings of 1990 International Conference on Parallel Processing*, pages II.93–II.97, University Park PA, 1990.
- [NM91] Robert H. B. Netzer and Barton P. Miller. Improving the Accuracy of Data Race Detection. *Proceedings of the ACM SIGPLAN Symposium on Principles and Practice of Parallel Programming PPOPP, published in ACM SIGPLAN NOTICES*, 26(7):133–144, 1991.

- [NM92] Robert H. B. Netzer and Barton P. Miller. What are Race Conditions? Some Issues and Formalizations. *ACM Letters on Programming Languages and Systems*, 1(1), March 1992.
- [PPCa] Book II: PowerPC Virtual Environment Architecture. URL: www-106.ibm.com/developerworks/eserver/articles/archguide.html.
- [PPCb] PowerPC Microprocessor Family: Programming Environments Manual. URL: www-3.ibm.com/chips/techlib/techlib.nsf/products/PowerPC_970_and_970FX_Microprocessors.
- [Pug99] William Pugh. Fixing the Java Memory Model. In *Proceedings of the ACM 1999 Java Grande Conference*, June 1999.
- [Rin01] Martin Rinard. Analysis of Multithreaded Programs. *Lecture Notes in Computer Science*, 2126, 2001.
- [RL98] Martin C. Rinard and Monica S. Lam. The Design, Implementation, and Evaluation of Jade. *ACM Transactions on Programming Languages and Systems*, 20(3):483–545, 1 May 1998.
- [RM94] J. Ramanujam and A. Mathew. Analysis of Event Synchronization in Parallel Programs. In *Languages and Compilers for Parallel Computing*, pages 300–315, 1994.
- [RPA97] Parthasarathy Ranganathan, Vijay S. Pai, and Sarita V. Adve. Using Speculative Retirement and Larger Instruction Windows to Narrow the Performance Gap between Memory Consistency Models. In *Proceedings of The 9th ACM Symposium on Parallel Algorithms and Architectures (SPAA)*, pages 199–210, June 1997.
- [Ruf00] Erik Ruf. Effective synchronization Removal for Java. In *Conference on Programming Languages, Design, and Implementation (PLDI)*, 2000.

- [SAR99] Xiaowei Shen, Arvind, and Larry Rudolph. Commit-Reconcile & Fences (CRF): A New Memory Model for Architects and Compiler Writers. In *Proceedings of The 26th Annual International Symposium on Computer Architecture (ISCA)*, pages 150–161, May 1999.
- [SBN⁺97] Stefan Savage, Michael Burrows, Greg Nelson, Patrick Sobalvarro, and Thomas Anderson. Eraser: A Dynamic Data Race Detector for Multi-threaded Programs. *ACM Transactions on Computer Systems*, 15(4):391–411, 1997.
- [Sch89] Edmond Schonberg. On-the-fly Detection of Access Anomalies. In *Proceedings of the ACM SIGPLAN '89 Conference on Programming Language Design and Implementation*, volume 24, pages 285–297, Portland, OR, June 1989.
- [Sin96] Pradeep K. Sinha. *Distributed Operating Systems, Concepts and Design*. IEEE Press, 1996.
- [SL94] Daniel J. Scales and Monica S. Lam. The Design and Evaluation of a Shared Object System for Distributed Memory Machines. In *Operating Systems Design and Implementation*, pages 101–114, 1994.
- [SS88] Dennis Shasha and Marc Snir. Efficient and Correct Execution of Parallel Programs that Share Memory. *ACM Transactions on Programming Languages and Systems*, 10(2):282–312, April 1988.
- [Ste90] Guy L. Steele, Jr. Making Asynchronous Parallelism Safe for the World. In *Proceedings of the 17th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 218–231. ACM Press, 1990.

- [SWF⁺02] Z. Sura, C.-L. Wong, X. Fang, J. Lee, S.P. Midkiff, and D. Padua. Automatic Implementation of Programming Language Consistency Models. In *15th Annual Workshop on Languages and Compilers for Parallel Computing*, July 2002.
- [Tay83] Richard N. Taylor. A General-purpose Algorithm for Analyzing Concurrent Programs. *Commun. ACM*, 26(5):361–376, 1983.
- [vP04] Christoph von Praun. Efficient Computation of Communicator Variables for Programs with Unstructured Parallelism. In *Seventeenth Annual Workshop on Languages and Compilers for Parallel Computing*, September 2004.
- [vPG03] Christoph von Praun and Thomas R. Gross. Static Conflict Analysis for Multi-threaded Object-oriented Programs. In *Proceedings of the ACM SIGPLAN 2003 Conference on Programming Language Design and Implementation (PLDI)*, June 2003.
- [VR01] Frederic Vivien and Martin Rinard. Incremental Pointer and Escape Analysis. In *Proceedings of the ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI)*, 2001.
- [WFPS02] Peng Wu, Paul Feautrier, David Padua, and Zehra Sura. Instance-wise Points-to Analysis for Loop-based Dependence Testing. In *ACM International Conference on Supercomputing*, June 2002.
- [WR99] John Whaley and Martin Rinard. Compositional Pointer and Escape Analysis for Java Programs. In *Proceedings of the ACM SIGPLAN 1999 Conference on Object-oriented Programming Systems, Languages, and Applications (OOPSLA)*, pages 187–206, November 1999.
- [WSF⁺02] C.-L. Wong, Z. Sura, X. Fang, S.P. Midkiff, J. Lee, and D. Padua. The Pensieve Project: A Compiler Infrastructure for Memory Models. *International Symposium on Parallel Architectures, Algorithms, and Networks*, May 2002.

- [YSP⁺98] K. Yelick, L. Semenzato, G. Pike, C. Miyamoto, B. Liblit, A. Krishnamurthy, P. Hilfinger, S. Graham, D. Gay, P. Col ella, and A. Aiken. Titanium: A High-Performance Java Dialect. In *ACM 1998 Workshop on Java for High-Performance Network Computing*. ACM SIGPLAN, 1998. URL: <http://www.cs.ucsb.edu/conferences/java98>.
- [YT88] M. Young and R. M. Taylor. Combining Static Concurrency Analysis with Symbolic Execution. *IEEE Transactions on Software Engineering*, 14(10):1499–1511, 1988.

Vita

Zehra Sura was born on April 12, 1976 in Dubai, U.A.E. She earned a Bachelor of Engineering degree in computer science from Visvesvaraya College of Engineering, Nagpur, India in 1998. Immediately after that, she began her graduate study in computer science at the University of Illinois, Urbana-Champaign (UIUC). In Fall 1998, she worked as a research assistant in the Parallel Programming Laboratory, UIUC under the guidance of Professor Laxmikant Kale. In Spring 1999, she was a teaching assistant for an introductory course in parallel programming in the department of computer science, UIUC. Since the summer of 1999, she has worked as a research assistant under the guidance of Professor David Padua. She did an internship at Intel KAI Lab, Champaign, Illinois in the summer of 2001, and at IBM T.J. Watson Research Center, Yorktown Heights, New York in the summer of 2003. In 2004, she completed her Doctor of Philosophy degree in computer science. After graduation, she will join the IBM T.J. Watson Research Center in Yorktown Heights, New York. Her research interests include analysis and transformation of programs for parallel processing, and the use of runtime techniques to improve compiler optimizations.